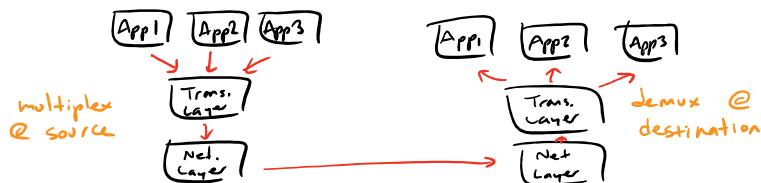


Chapter 3 - Transport Layer

Wednesday, February 10, 2016 9:24 AM

- Between application layer & network layer
- provides logical communication between processes
- can provide reliable communication even when underlying network protocol is unreliable (IP)
 - network layer provides communication between hosts
- convert application-layer messages to and from transport-layer segments
 - ↳ direct messages sent from host-to-host to the right processes using multiplexing & demultiplexing → passing through sockets

MULTIPLEXING / DEMULTIPLEXING



source port #	dest port #
other header fields	↳ takes data from sockets & adds header information to chunks to create segments to pass to network layer
application data	

port #'s are 16-bits (0-65535), 0-1023 are restricted for well-known application protocols

UDP

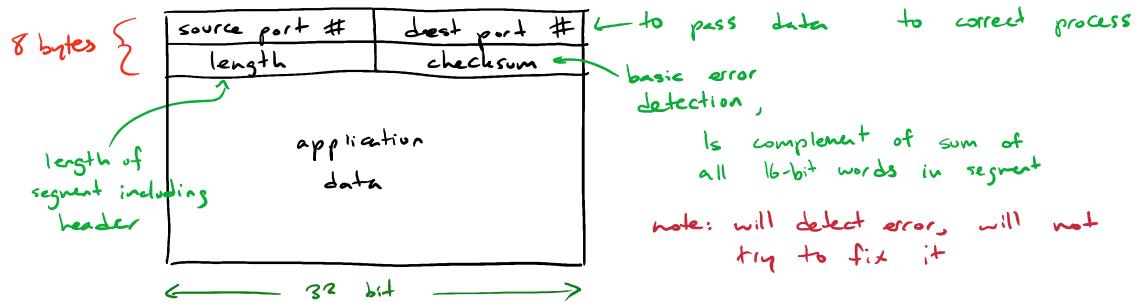
- minimal, connectionless transport protocol
 - ↳ no handshaking, each UDP segment handled independently
- basically talking directly with IP except for some basic integrity checking & muxing/demuxing
- no guarantee of reliable delivery
 - ↳ must be built into the application itself

BENEFITS

- precise application-level control of what is sent & when
- no connection establishment delay → why DNS uses UDP
- no state/buffer information needed

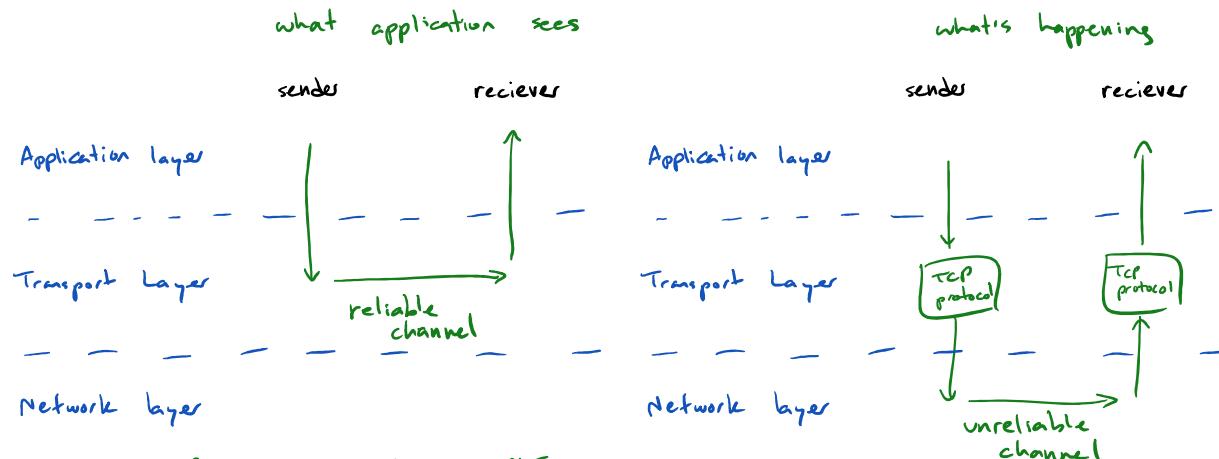
- ↳ can handle many active clients
- ↳ TCP needs states for reliable data
- small packet overhead
- no congestion control → good for the app, bad for network

SEGMENT HEADER



RELIABLE DATA TRANSFER

→ no data is corrupted or lost and all are delivered in order



PRINCIPLES OF RELIABLE DATA TRANSFER

- retransmission → bad/missing packets retransmitted
- error detection → checksum
- receiver feedback → acknowledgements (positive & negative)

IP

ARQ → reliable transfer protocols based on retransmission

Automatic
Repeat
request

STOP & WAIT

→ sends packet, waits for acknowledgement, sends another, if receives NAK, resends

TIMOUT → if doesn't receive ACK, will resend after timeout

ACK & NAK should also have checksum for reliability

↳ if garbled ACK or NAK, resend packet to be sure

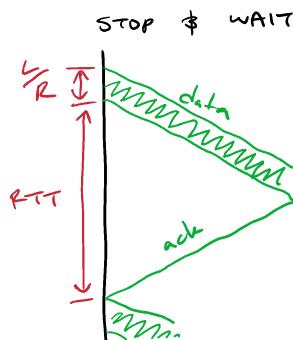
↳ causes duplicate packets, so ACKs & NAKs

must also have sequence numbers associated with packet

PIPELINING

→ sending a certain amount of packets one-after-another
before the previous have been acknowledged

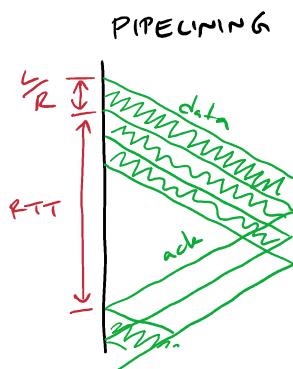
→ causes buffering



utilization

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R}$$

fraction of time sender is busy doing work
over entire time process for one sent packet



$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R}$$

↳ pipelining with 3
triples our utilization
factor

requirements:

- must increase range of sequence numbers
- sender/receiver have to buffer multiple packets
- more complicated loss/ack techniques

↳ Go-Back N, selective repeat

GO-BACK-N (GBN)

→ sender sends N packets w/o waiting for an ACK
"window size"

→ receiver will send ACK for all packets received in order

↳ "cumulative acknowledgement" → has received all packets up to # including #

→ if packet arrives out-of-order, it will be discarded & ACK for most recently received packet in order will be resent

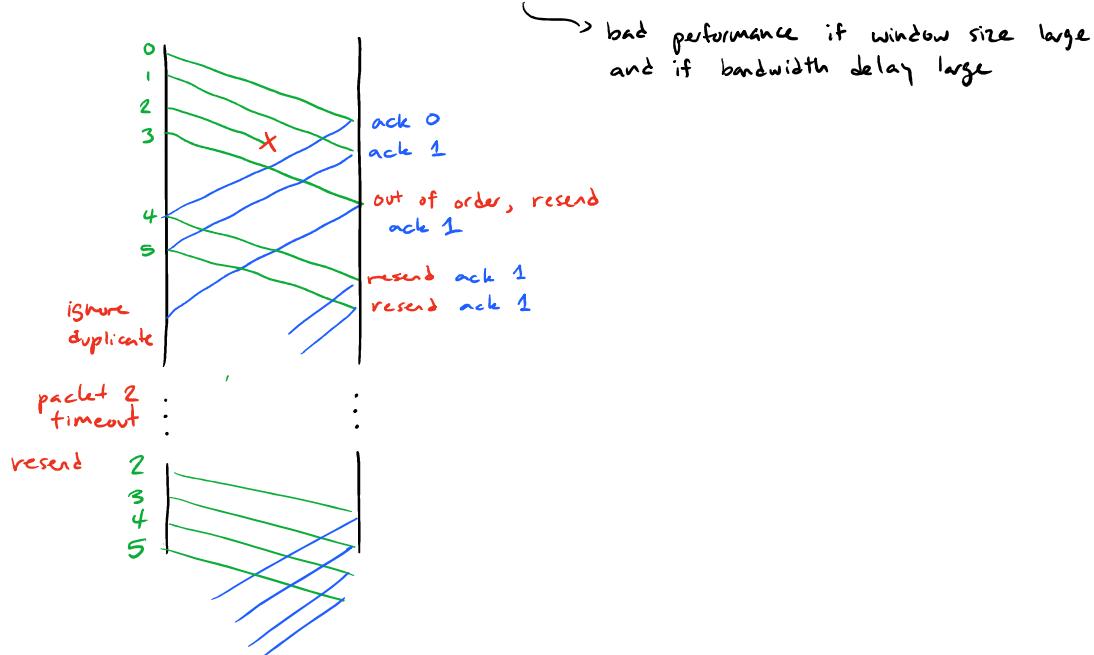
→ still has a timer like stop & wait, but upon timeout will resend unacknowledged packet and resend packets above unacknowledged

PROS

- receiver doesn't buffer any out-of-order packets
- only need to use one timer, for oldest unacknowledged packet

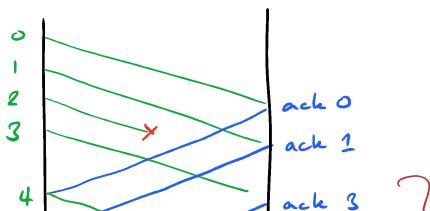
CONS

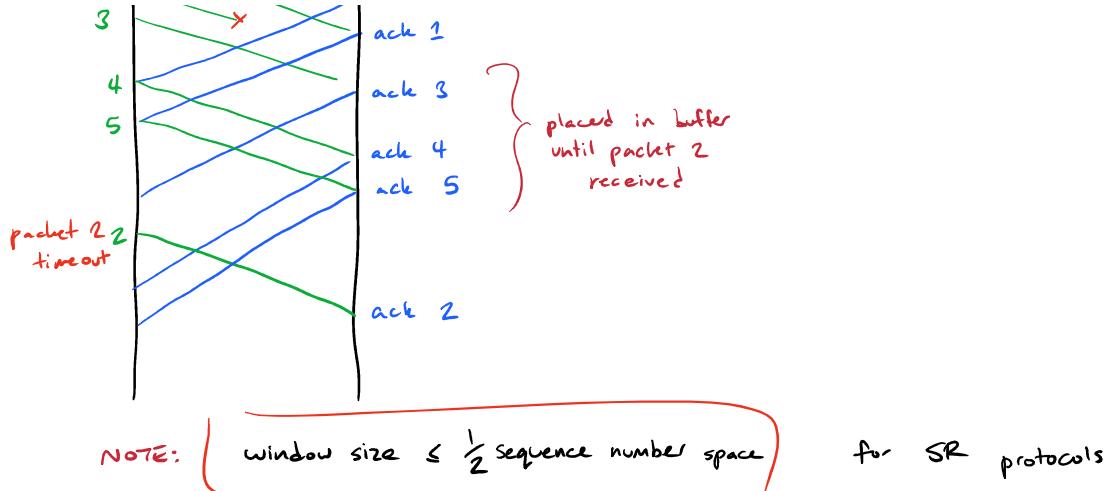
- wasteful b/c retransmitting all N packets when only one fails



SELECTIVE REPEAT (SR)

- receiver individually acknowledges correctly received packets
- sender only retransmits packets lost or corrupted
- out of order packets are buffered until missing packets received
- each packet has own timer & timeout





↳ otherwise sequence number rolls over and receiver doesn't know if receiving a retransmission or a new packet

TCP

- reliable data transfer

↳ flow control, sequence numbers, acknowledgements, timers, congestion control

↳ don't overwhelm receiver

↳ reliability

↳ equal link bandwidth for each connection

- connection-oriented

↳ 3-way handshake,

↳ maintain connection throughout session

↳ create connection
↳ acknowledge/send request

← 32 bit →	
source port #	dest. port #
sequence number	
ACK number	
hdr len	flags
checksum	receive window
options	urg data pointer
application data	

uses cumulative acks

20 bytes } no field that explicitly says what the entire length of the TCP segment is

↳ minimum header size, if options length = 0

shared w/ UDP:

- destination/source port #
- length (although only header length w/ TCP)
- checksum

RST, SYN,
FIN

QUESTION

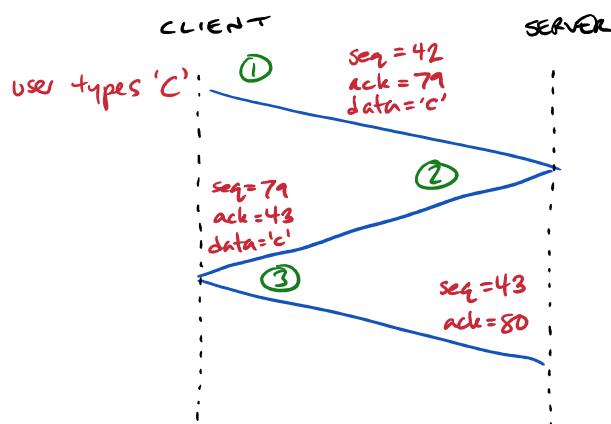


- destination/source port
- length (although only header length w/ TCP)
- checksum

SEQUENCE NUMBERS → identify each byte to be delivered
→ segments are given the sequence number of the first byte in their data

ACKNOWLEDGEMENT NUMBERS → contain the sequence # of the next byte waiting to be received

example using Telnet



- ① 42 and 79 were chosen randomly in handshake
 - ↳ represents first segment sent from server
 - ↳ acknowledging because that's what client is waiting for
 - ↳ represents first segment sent from client
- ② acknowledges it has received segment 42
 - ↳ because data 'c' was only one byte it tells client it now expects 42+1, byte 43,
 - and also echos back 'c'
 - ↳ this is the first segment being sent by the server, so the sequence # is 79, which was decided on handshake
- ③ acknowledges it received echo
 - ↳ because it received segment 79 which was one byte, it expects 80, so that's its ACK #
 - ↳ because it sent segment 42, which was one byte long, its new segment # is 43
 - ↳ even though it contains no data

RECEIVE WINDOW → used for flow control
→ indicates # of bytes willing to accept

HEADER LENGTH → specifies length of TCP header

OPTIONS FIELD → optional & variable length, typically not used

FLAGS → 6 bits, ACK, SYN, FIN, RST, PSH, URG

URGENT DATA POINTER → not really used
→ identifies location of urgent data

TCP CONNECTION VARIABLES

receive buffer

send buffer

receive window

last byte read

congestion window

↪ constraint on rate sender can send traffic into network

amount of unacknowledged data may not exceed minimum of cwnd or rwnd

$$\boxed{\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}}$$

$$\boxed{\text{sender rate} \approx \frac{\text{cwnd}}{\text{RTT}}}$$

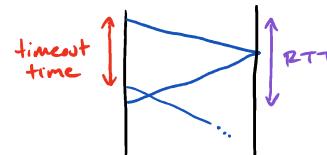
can control rate by adjusting cwnd

RTT ESTIMATIONS & TIMEOUT

↪ timeout must be longer than RTT otherwise it will send unnecessary retransmissions

↪ if too long, slow reactions to loss

↪ RTT is variable, so timeouts are computed and changed



To estimate RTT:

↪ TCP records RTT for one unacknowledged segment at a time called SampleRTT

↪ to smooth it out, an exponential!

moving average is kept track of
and updated like so

$$\text{EstimatedRTT} = (1-\alpha) \text{EstimatedRTT} + (\alpha) \text{SampleRTT}$$

↪ note: this is an assignment called with
every new value of SampleRTT,
it is not an equation

$$\alpha = \frac{1}{8} = 0.125$$

↪ recommended value

↪ the variability/deviation is also
kept track of as DevRTT

$$\text{DevRTT} = (1-\beta) \text{DevRTT} + (\beta) |\text{SampleRTT} - \text{EstimatedRTT}|$$

$$\beta = \frac{1}{4} = 0.25$$

↪ recommended

↪ using EstimatedRTT and DevRTT we keep
a running value for the timeout interval as such

$$\text{Timeout Interval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

↪ initially set to 1 second

↪ safety margin

TCP's RELIABLE DATA TRANSFER

keys:
error detection
receiver feedback
retransmission

→ TCP only uses one timer even if there
are multiple transmitted unacknowledged segments
↪ timers require considerable overhead

simplified model:

TCP responds only to three major events

Data received from application above

- ↳ create TCP segment w/ sequence # Next Seq Num
- ↳ start timer if not already running
- ↳ pass segment to IP
- ↳ set Next Seq Num = Next Seq Num + length (data)

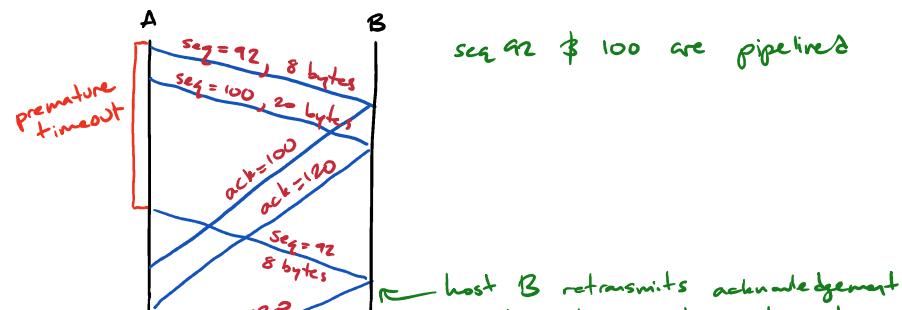
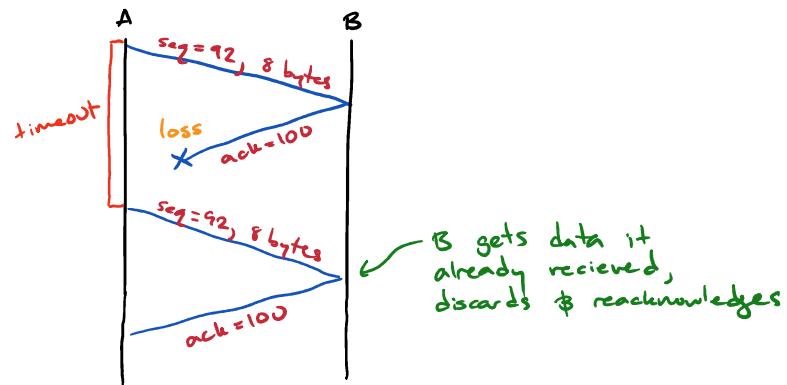
Timer timeout

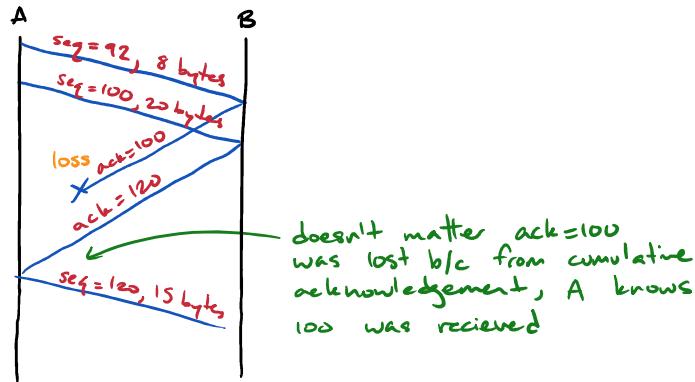
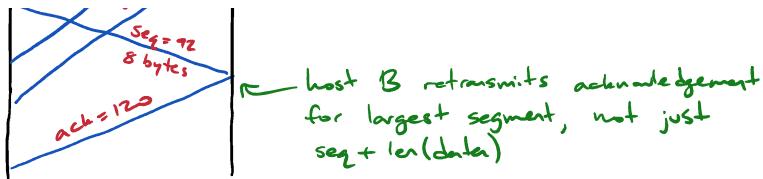
- ↳ retransmit unacknowledged segment w/ smallest sequence #
- ↳ start timer
- ↳ double timeout interval

Acknowledgment received

- ↳ if incoming ACK number > sequence # oldest unacknowledged byte
then the ACK is acknowledging a previously unacknowledged segment
- else
then the ACK is about something it already acknowledged, so ignore it

RETRANSMISSION SCENARIOS



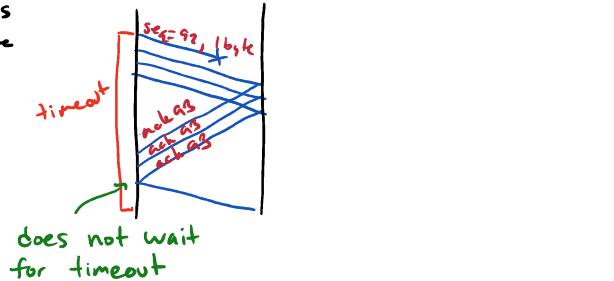


OTHER TCP BEHAVIOURS

FAST RETRANSMIT

- when segment is lost, the long timeout is a delay
- the sender can detect packet loss by duplicate acks

if 3 duplicate ACKS a loss
is probable, retransmit before
timeout

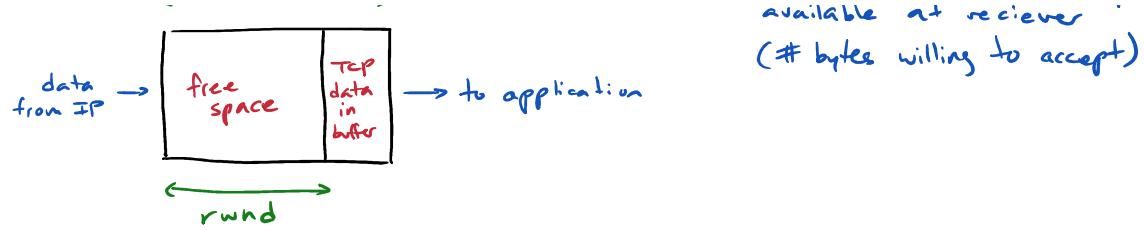


FLOW CONTROL

- avoids overwhelming receiver / overflow receiver buffer
- in TCP both sides specify receive window, rwnd



how much free buffer space available at receiver (# bytes willing to accept)



in order to not overflow:

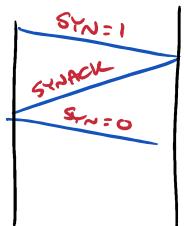
$$\boxed{\text{LastByteRead} - \text{LastByteRead} \leq \text{RcvBuffer}}$$

receive window is amount of spare room

$$\text{rwnd} = \text{RcvBuffer} - (\text{LastByteRead} - \text{LastByteRead})$$

CONNECTION MANAGEMENT

- ↳ three-way handshake to establish connection
- ↳ agree on connection parameters



CONGESTION CONTROL

- ↳ end-to-end congestion control (used by TCP)
- ↳ TCP receives no information about congestion state from IP

TCP keeps track of congestion window

$$\boxed{\text{LastByteSent} - \text{Last Byte Acked} \leq \min\{\text{cwnd}, \text{rwnd}\}}$$

amount of unacknowledged data at sender

$$\boxed{\text{TCP sender rate} \approx \frac{\text{cwnd}}{\text{time}}}$$

So TCP can control sending rate by modifying cwnd "self-clocking"

DATA AT SERVER

$$\text{TCP sender rate} \approx \frac{\text{cwnd}}{\text{RTT}}$$

so TCP can control sending rate by modifying cwnd "self-clocking"

↳ it will increase cwnd until loss event happens, timeout or 3 duplicate ACKs

slow start

cwnd=1, double cwnd every RTT

continue slow st. until

ssthresh exceeded, then increase only linearly

process until loss

if timeout cwnd=1, ssthresh = $\frac{1}{2}$ cwnd

if 3 dup ACKs cwnd = $\frac{1}{2}$ cwnd, ssthresh = $\frac{1}{2}$ cwnd