


# Dynamic Memory Systems Design

- **Introduction and Overview**
- **Dram Technology**
- **Conventional Asynchronous Dram Timing Analysis:**
  - **Ras\*/Cas\* Signalling,**
  - **Access and Cycle Times.**
- **Asynchronous Dram Controller for a 68000 Based System.**
- **Refreshing Concepts.**
- **Fast Page Mode Drams,**
- **Synchronous Dram (SDRAM)**
- **Double Data Rate Synchronous Drams (DDR-SDRAM)**



# Introduction and Overview

## Introduction

- Dynamic ram (DRAM) is an alternative to SRAM for random access, volatile R/W memory.
- Mainly used in high performance computers, such as PC's or workstations.

## Benefits of Dram over Sram are

- **Cost, Storage Capacity, Power and Packaging Densities :**

- Beyond a certain size it's much **cheaper** to build large memory systems around Dram than Sram.
- This is because a dram cell is typically 18 x **smaller** than the equivalent sram cell, requiring 1 transistor per bit rather than 6 (+interconnect)
- This results in **less power consumption** and more silicon available for memory cells leading to **higher storage densities** than to Sram.
- Packaged dram chips are also smaller than their equivalent Sram packages as they utilize a multiplexed address bus which also leads to a reduction in PCB size.

# Introduction and Overview

## Drawbacks of using Drams:

- **Complex Interfacing**: Dram require a dedicated controller that can take a CPU address and multiplex it (split into two halves) onto the Dram address lines. The timing of signals is also more complex.
- **Need for Refreshing**: Due to their inherent design characteristics, Drams have a nasty habit of forgetting their data if they are not periodically refreshed by external circuitry.

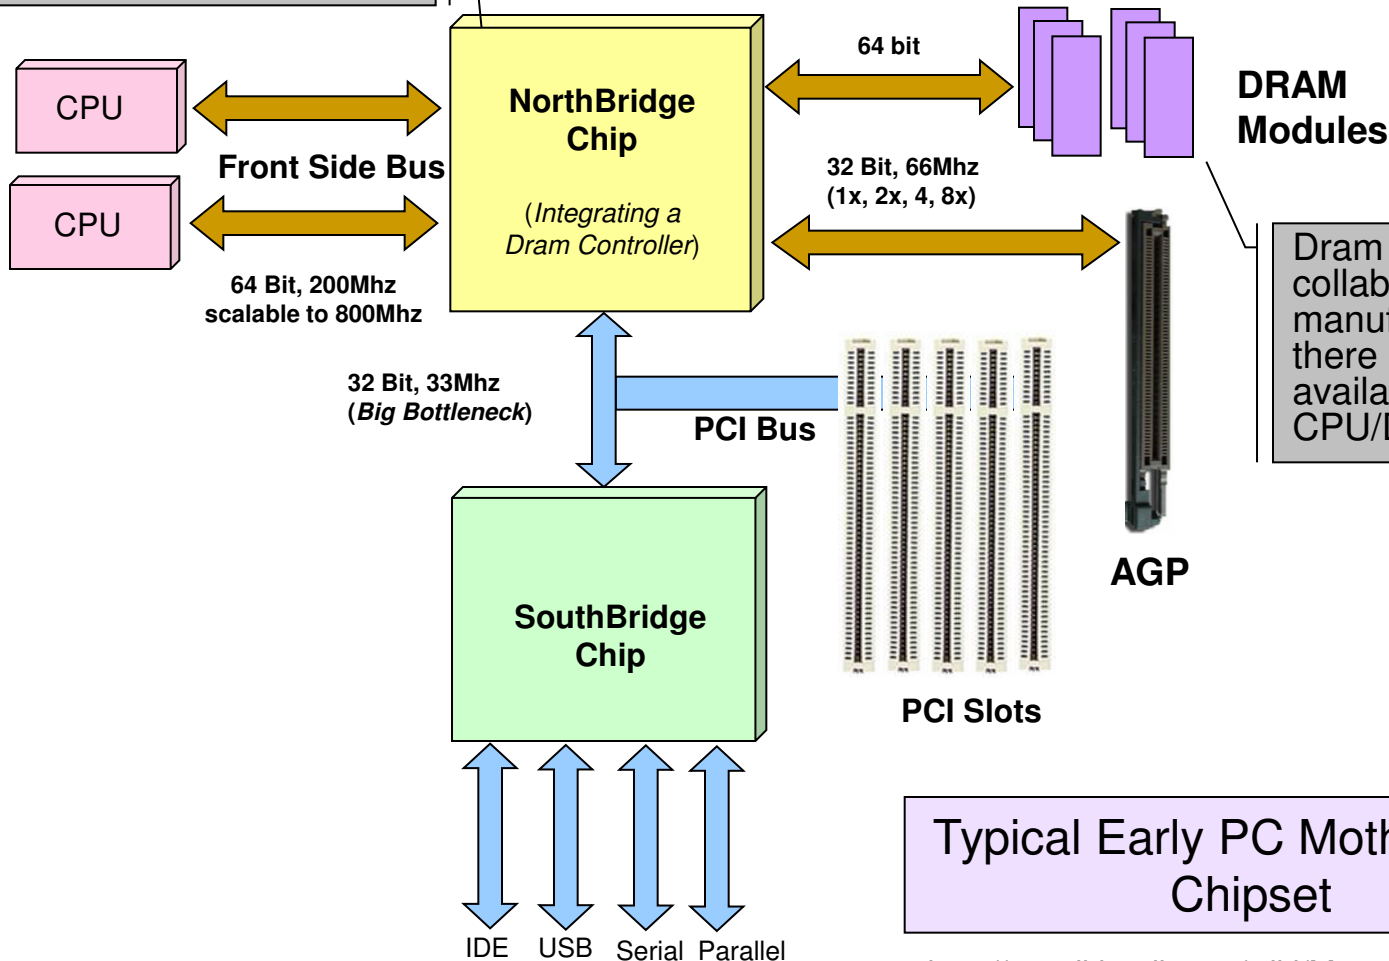
Additional refresh circuitry has to be built into the dram controller, increasing the cost and complexity of the resulting system.

- **Lower Speed**. Drams are generally slower than their equivalent Srams which is why most high speed computers use caches (made of sram) to store frequently accessed program/data. If you ran your computer directly from dram without a cache it would be pretty slow.

# Dram Controllers in Early PCs

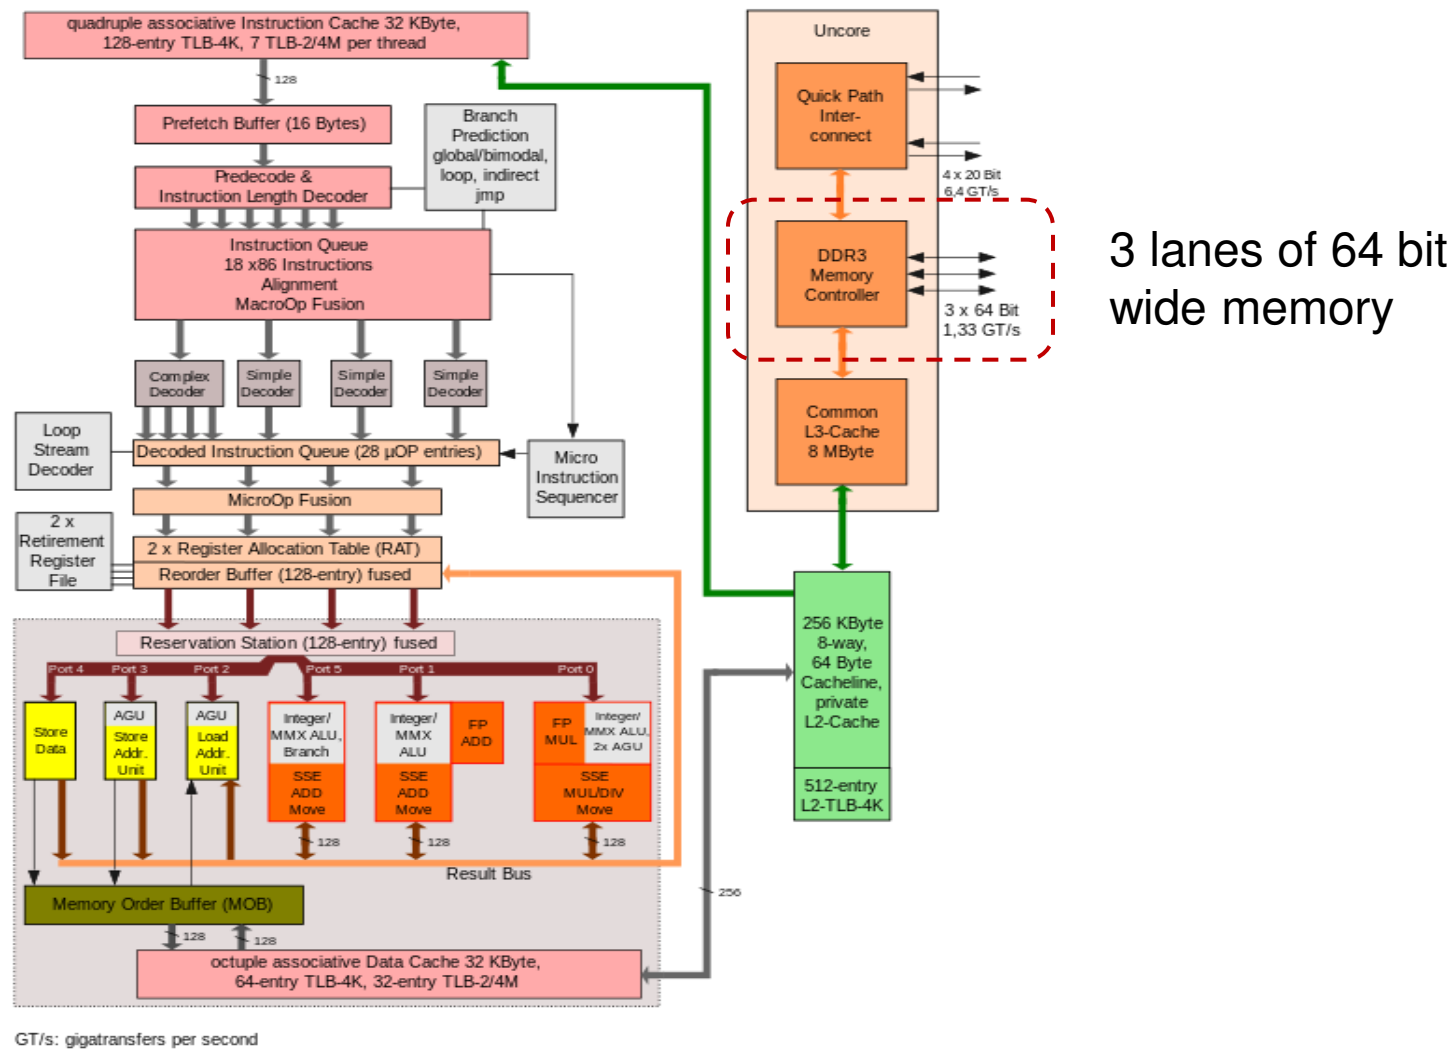
Dram technology almost always requires dedicated memory controllers for each combination of host CPU and Dram technology. Can't just hook up CPU directly to the memory chips

Dram controllers handle the complex task of generating the signaling, delicate timing, and refreshing required to get the best performance out of the latest generation of devices.



Typical Early PC Motherboard Chipset

# Dram Controllers now Integrated into I7 CPU

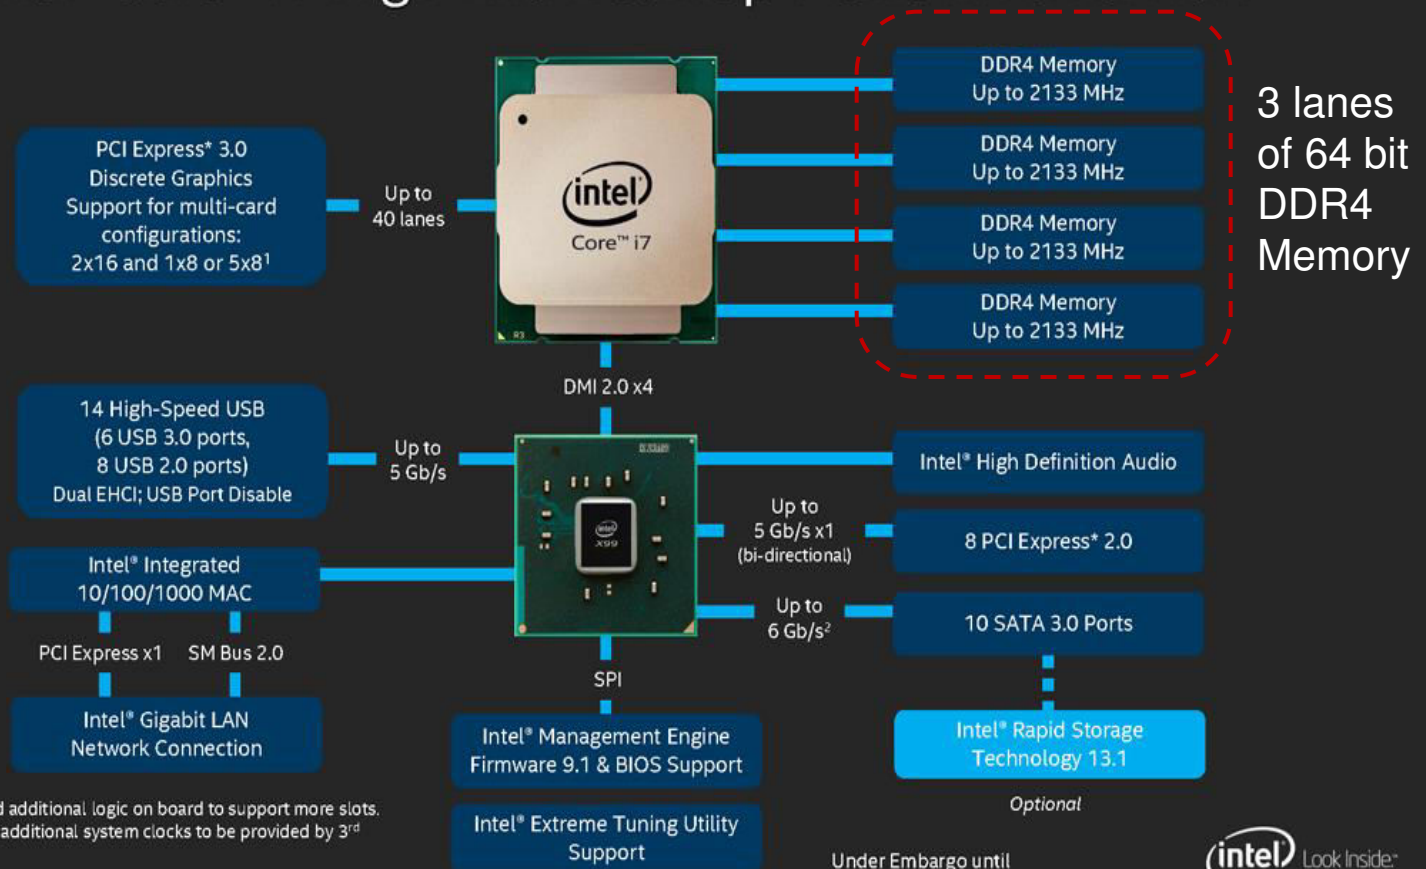


Intel I7 "Nehalem" core with integrated DDR3 memory controller

# New “Haswell” I7 Processor Motherboard

**DDR4** memory controller built into CPU – **Northbridge chipset** is “no more”

## Intel® Core™ i7 High End Desktop Platform Overview



<sup>1</sup> 3 slots available, but need additional logic on board to support more slots. 5x8 configuration requires additional system clocks to be provided by 3<sup>rd</sup> party components.


<sup>2</sup> All SATA ports capable of 6 Gb/s.

Under Embargo until  
9:00am PST, August 29, 2014



# Dram and Microcontrollers

- It is highly unlikely that you would ever use Dram for **8** or even **16** bit systems. The reasons for this are
  - Most microcontroller programs are stored in **Flash Eprom**, not loaded into static or dynamic memory so the **need** for large amounts of R/W memory is more limited.
  - 8 and 16 bit systems just cannot address large amounts of memory anyway and are aimed at simpler applications.
  - Even relatively large amounts of Sram are **inexpensive**. For example a **4Meg x 16** SRam interfaces directly to a CPU with minimal complexity.




- 128KB Flash EPROM
- 4KB EEPROM
- 32KB SRAM
- 8-ch, 10-bit, ADC
- 8-bit Timer
- 7-ch, 8-bit PWM
- 9 KBI inputs
- 56 GPIO
- 3 CAN Channels
- 2 SCI & 2 SPI Channels
- I2C Channel

6812 - 16 bit processor board from Axion Manufacturing




# Dram and Microcontrollers

- For more sophisticated 32 and 64 bit processors their larger address space means they run more sophisticated code (e.g. graphics and real time operating systems) making Dram more suitable.
- Today a number of 32 bit microcontrollers have an integrated a Dram controller making it easy to interface with external dram. e.g.
- Opposite is a cold **Coldfire** (68k based) development board from FreeScale with 64MB Dram + Graphics, network etc





## XCalibur4401 - Intel® Core™ i7 Processor Single Board Computer



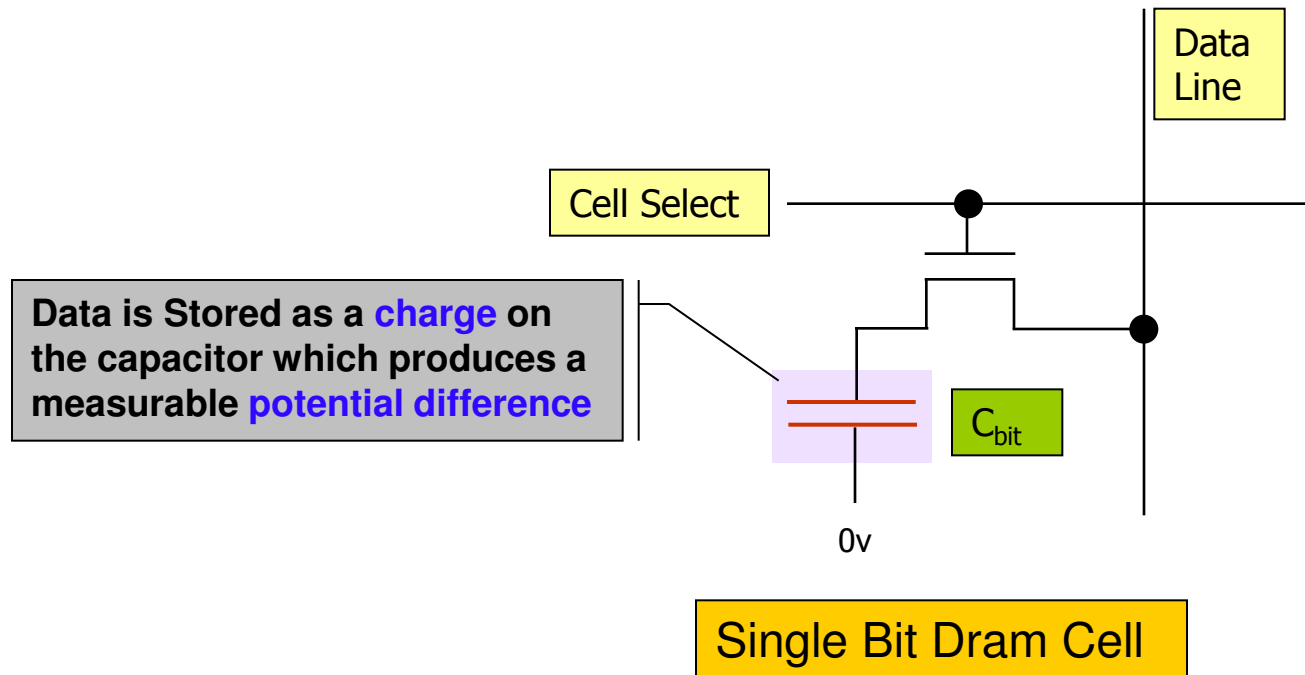
- Up to **16 GB of DDR3 - ECC DRAM**
- Up to **128 GB of Flash Memory**
- Three front panel **Gigabit Ethernet** ports
- Graphics Interface
- Two **USB 2.0** ports
- Four **SATA** ports
- Front and rear graphics ports
- Wind River VxWorks, Linux, Microsoft Windows, Neutrino, and LynuxWorks LynxOS operating systems

Example of an Intel I7 based industry processor board with Compact PCi bus interface

# Dram Technology: The DRAM Cell



## Dram Basics

- DRAM uses only **one transistor per bit** of storage as shown below.
- To understand the operation of this cell we have to look beyond pure **logic** and consider some **physics** associated with the manufacture of MOSFETs.
- A logic **1** or **0** is maintained, either by the *presence* or *absence* of a charge, across a capacitor  $C_{\text{bit}}$ . The charge gives rise to a voltage of about **250mV**.



# Dram Technology: The DRAM Cell


- Several common methods have been used to fabricate the inherent 'capacitor'.
- Perhaps the earliest is the 'trench capacitor' method illustrated below which etches a hole or a 'well' into the substrate and connects it to the Drain of the Mosfet.
- The well is lined with Silicon Dioxide and filled with Poly silicon because of its good dielectric properties (*i.e. makes a big capacitor value relative to most other materials*). Even so the capacitor is small in value, typically  $<10^{-12}$  farads i.e.  $<1\text{pf}$ .



Dram cells difficult to scale.  
Currently about 20nM, going  
small increases leakage  
affecting data retention time

# Dram Technology: The DRAM Cell


- Microscopic view of Intel's embedded dram capacitors clearly shown buried in substrate. This was used to make cache memory on latest I7 processor.




# Dram Technology: The DRAM Cell

- More recently, technology has created an inherent capacitor *buried* under the transistor, the so called 'Z-Cell', **1T** or **zero capacitor** dram cell, which takes up less space and increases the storage density on the chip.

Logic 1



Logic 0



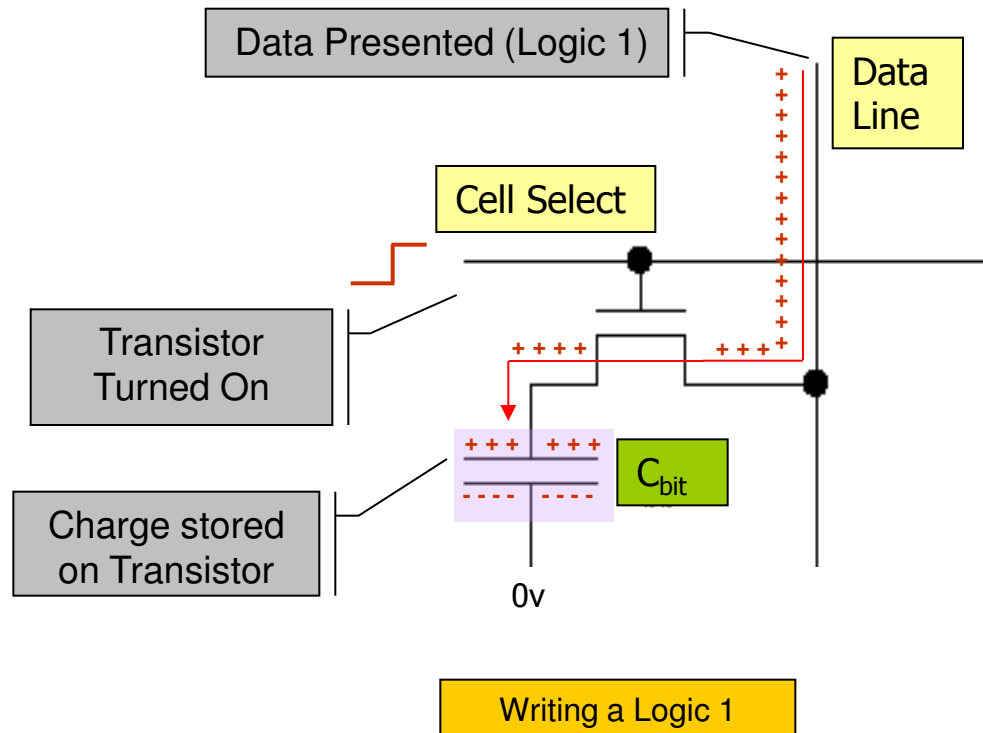
Z-RAM relies on the floating body effect, which was first encountered in the new silicon on insulator (SOI) process introduced in the early 2000s. This effect causes capacitance to form between the transistor and the underlying insulating substrate, and was considered a **problem** that needed to be solved in conventional digital systems. The same effect, however, allows a DRAM-like cell to be built using the transistor only, the floating body effect taking the place of the conventional trench capacitor. (<http://en.wikipedia.org/wiki/Z-RAM>)



# Dram Technology: The DRAM Cell

## Storing a Logic 1 in a Cell

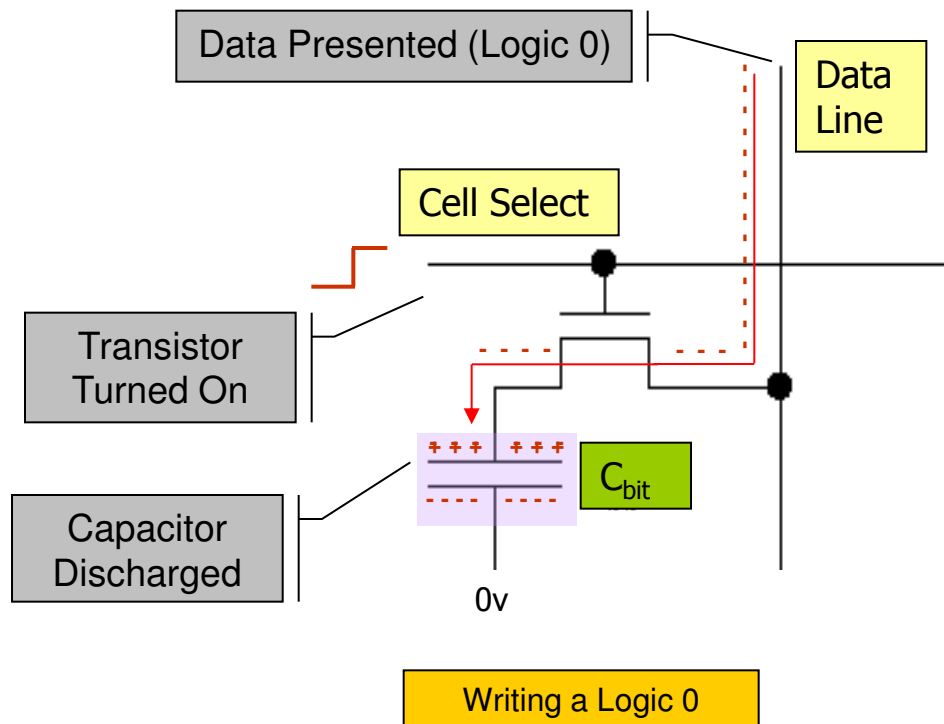
- A logic 1 is first placed onto the *Data* line.
- The *Cell Select* signal, (*from the chips internal address decoder*), is taken **high** when an address is presented to the chip. This turns on the transistor.
- Positive charge will flow from the data line and will be deposited onto  $C_{bit}$  creating a voltage potential across the capacitor.
- The Cell select line is then negated, leaving the charge trapped on the capacitor.



# Dram Technology: The DRAM Cell

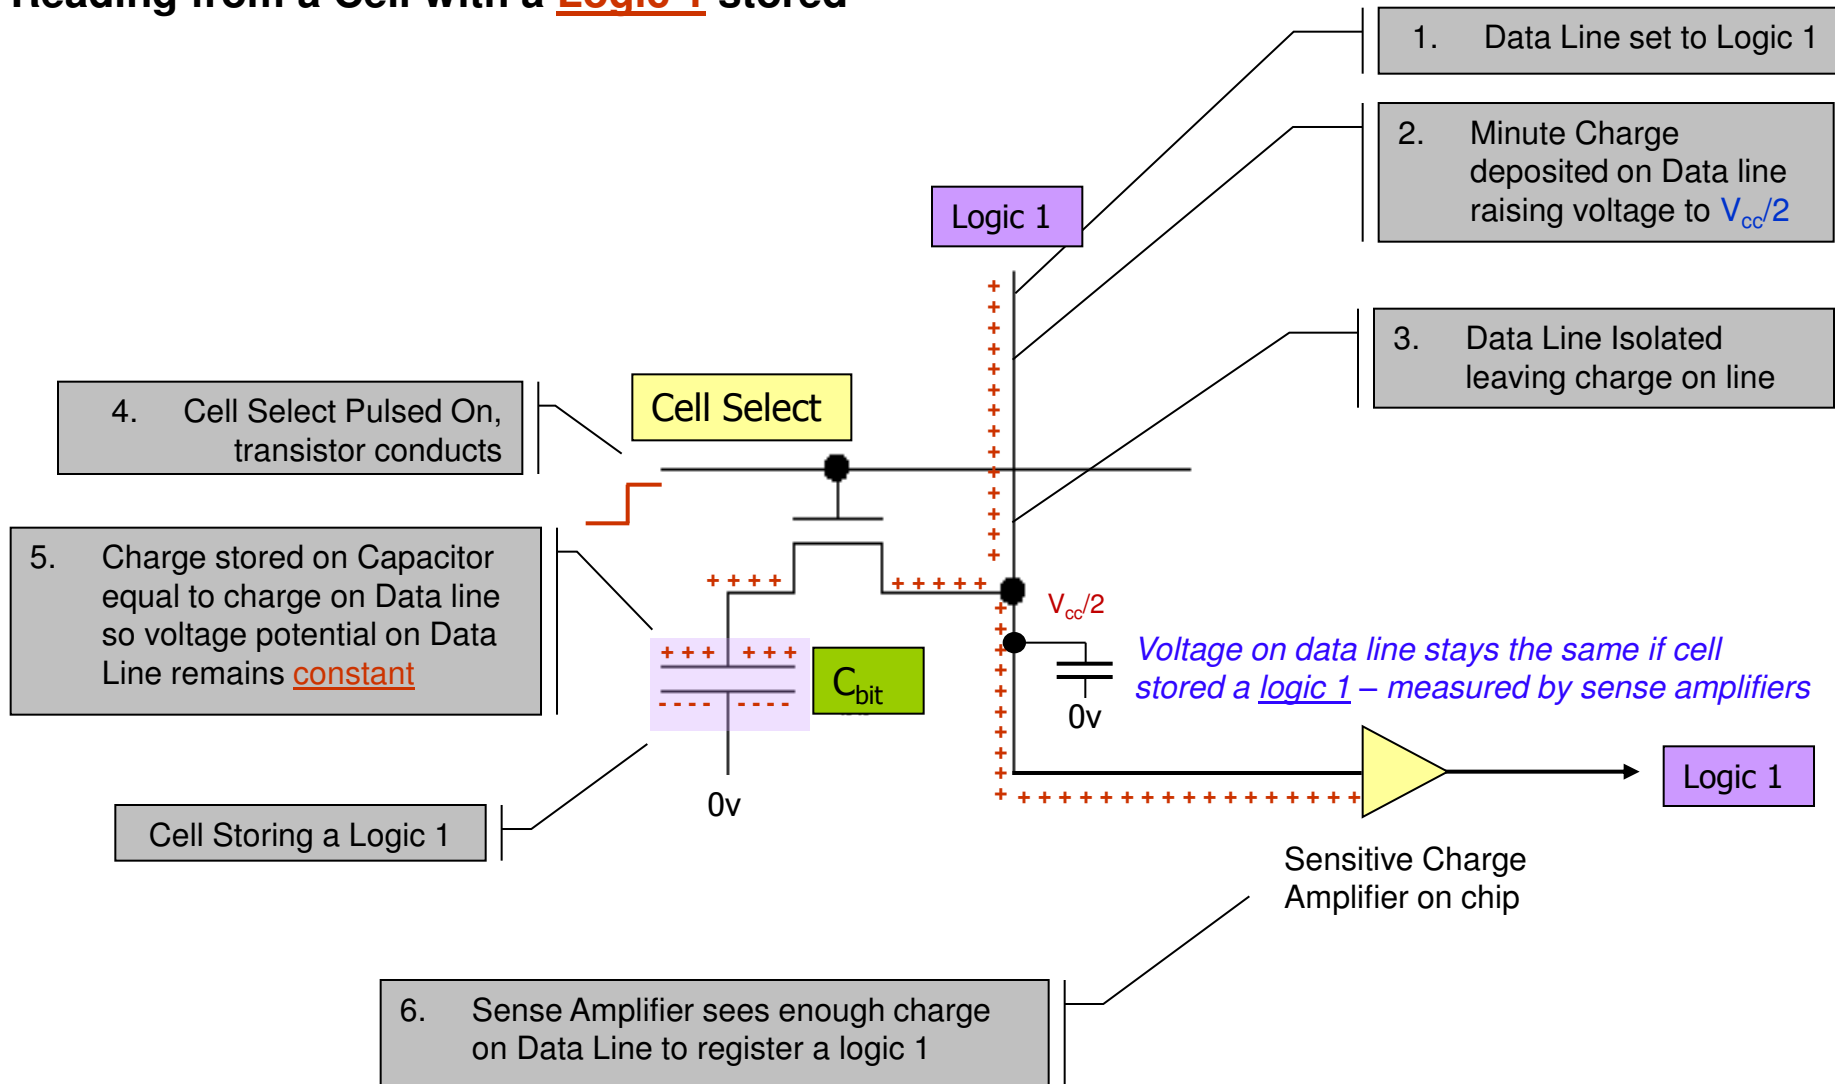
## Storing a Logic 0 in a Cell

- A logic 0 is placed on the *Data* line
- The *Cell Select* signal is taken **high** when an address is presented to the chip. This turns on the transistor.
- Any positive charge that might have been stored on the capacitor (*due to it previously storing a logic 1*) will be discharged creating no voltage potential across the capacitor.
- The Cell select line is negated, leaving the capacitor discharged.



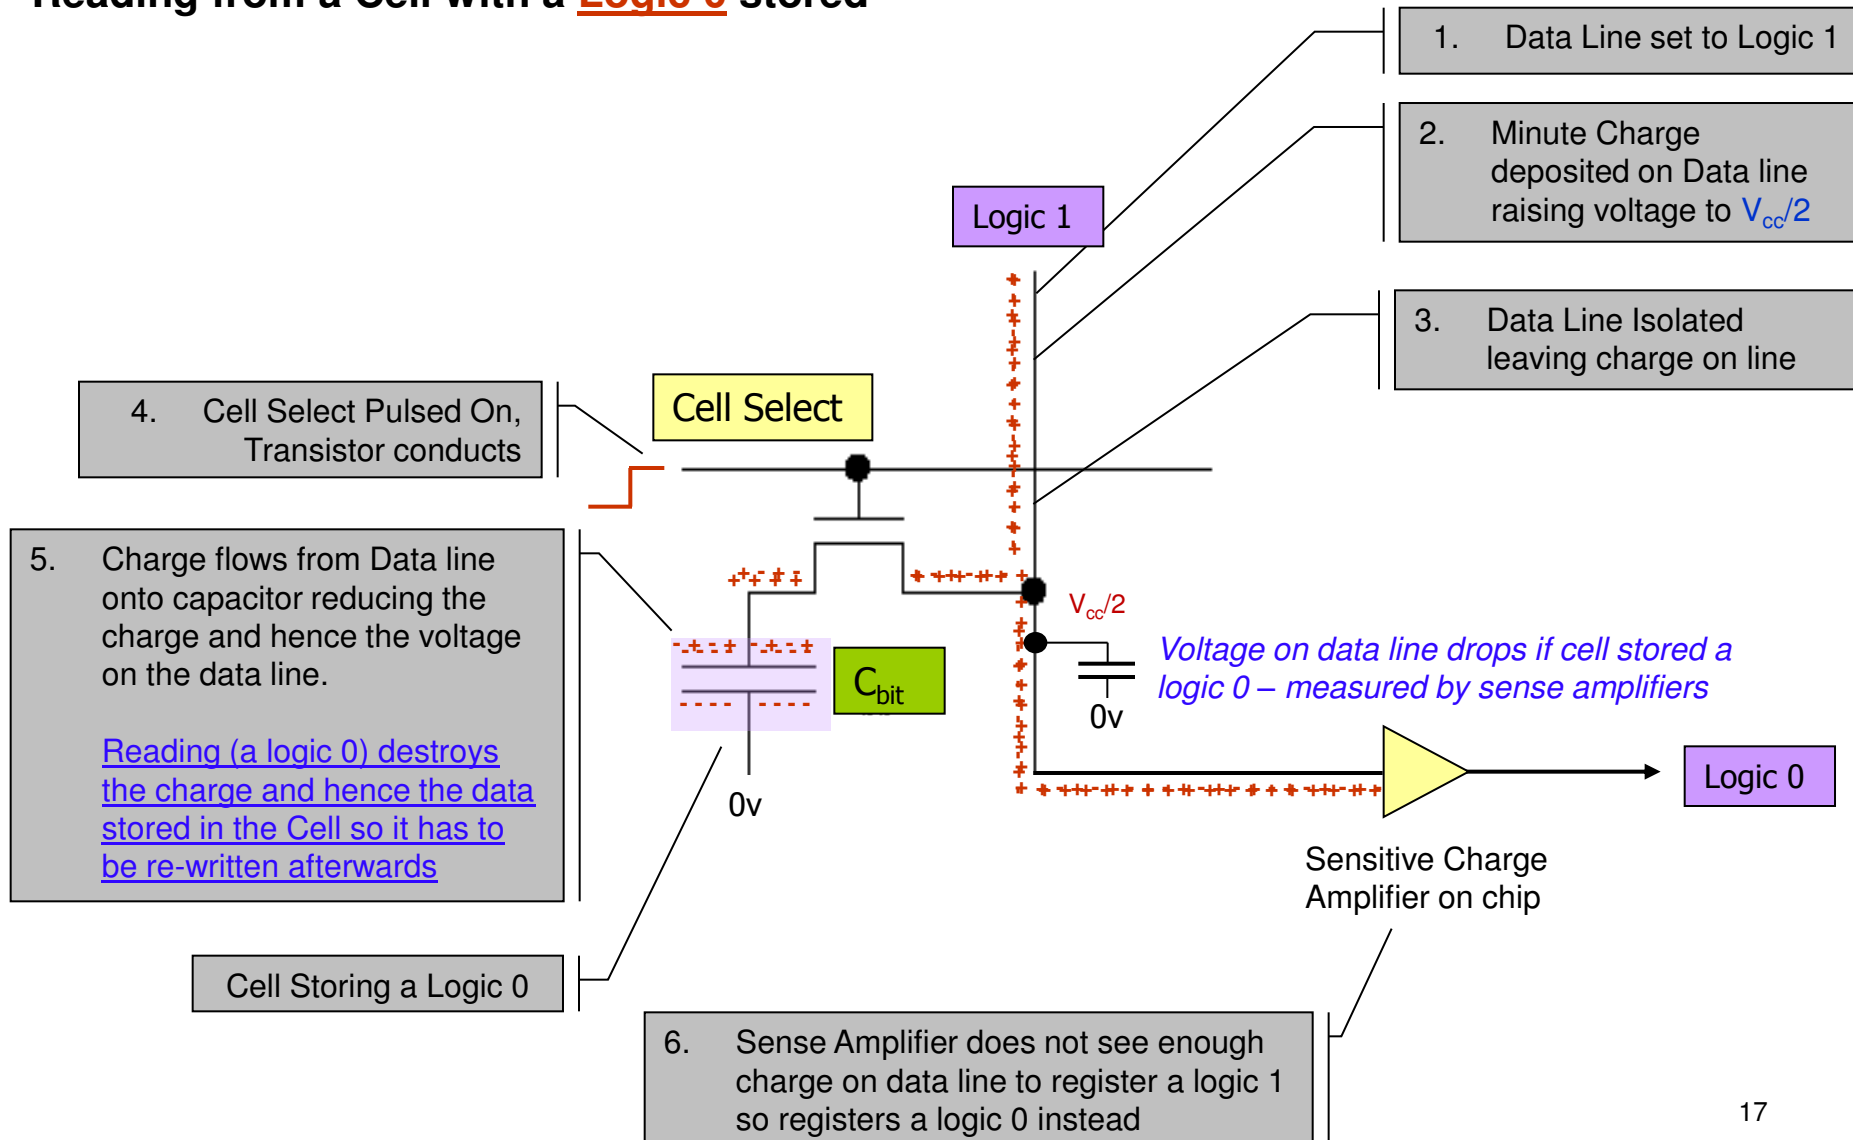
# Dram Technology: The DRAM Cell

## Reading from a Cell with a Logic 1 stored



# Dram Technology: The DRAM Cell

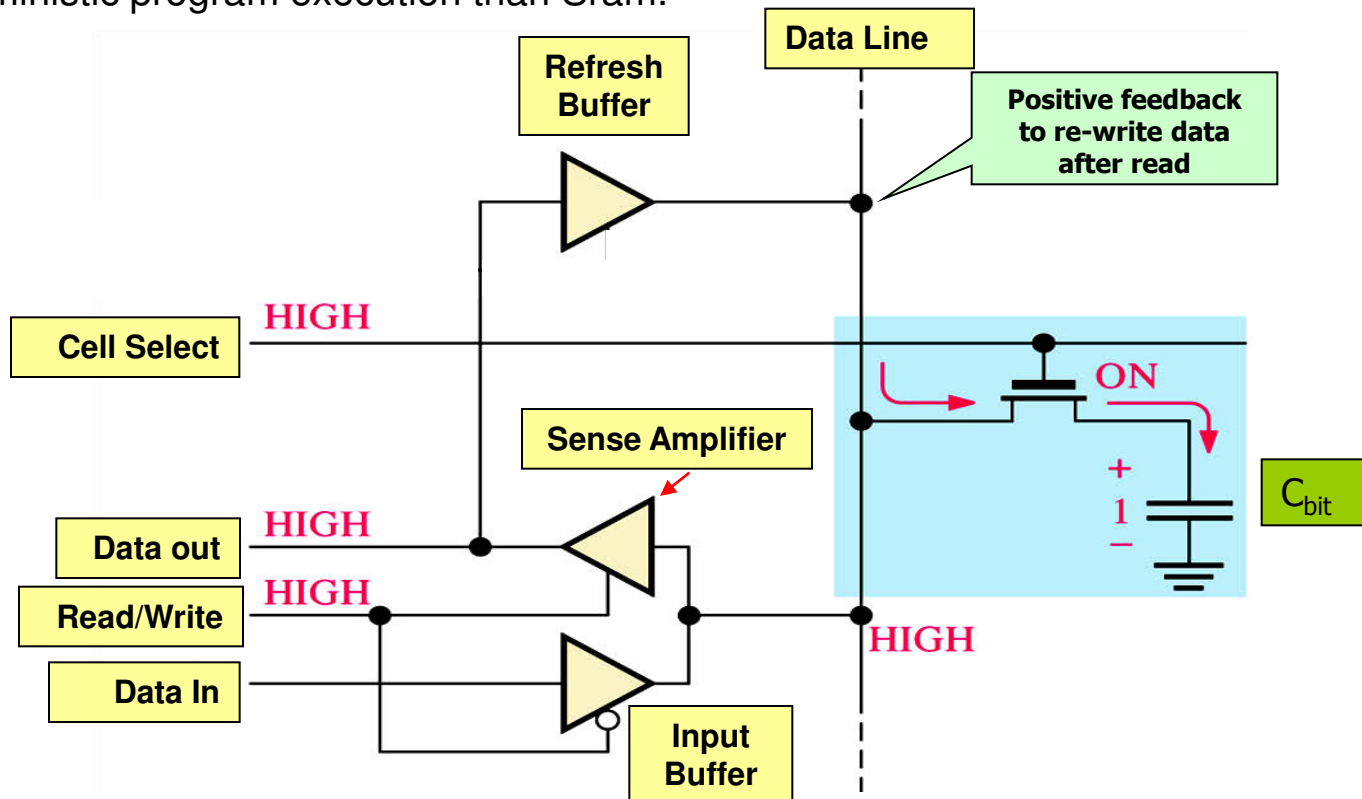
## Reading from a Cell with a Logic 0 stored



# Problems with Dram Technology

## Charge Leakage

- The minute charge stored on ( $C_{\text{bit}}$ ), tends to **leak** away with time leading to lost data and is also corrupted by the action of **reading** the cell.
- **Refreshing** must therefore be performed whenever a cell is read (*see feedback loop below reinforcing the charge stored on the capacitor*), and, in the absence of frequent reads, has to be performed **periodically** by **external hardware**.
- **During a periodic refresh** the CPU will **not** be able to **access** the chip making for less deterministic program execution than Sram.





# Problems with Dram Technology

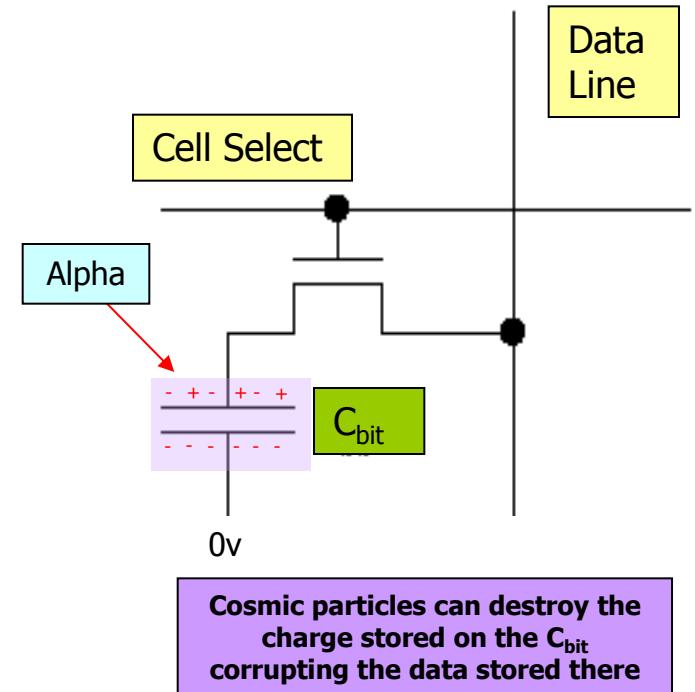
## Problems with Dram Cells: Read Cycle Times

- **Reading** is a relatively slow process because it requires that a precise measured charge be deposited on the Data line (*so that the sense amplifiers can detect the difference between a 0 and 1 stored in the cell – see page 14/15*) which takes time for the chip to generate internally.
- Thus there is a period of time after each access, which the Dram chip needs to **recover** and thus it **cannot be accessed** straight away, slowing down subsequent each access.
- This recovery time is known as the **Row pre-charge time** and is comparable to that of the **access time** of the device.
- As a consequence of this recovery time, the chips **cycle time** (i.e. the time between successive CPU accesses to the memory chip) is about **twice** as long as the **access time** and thus wait states may be required in very fast systems.

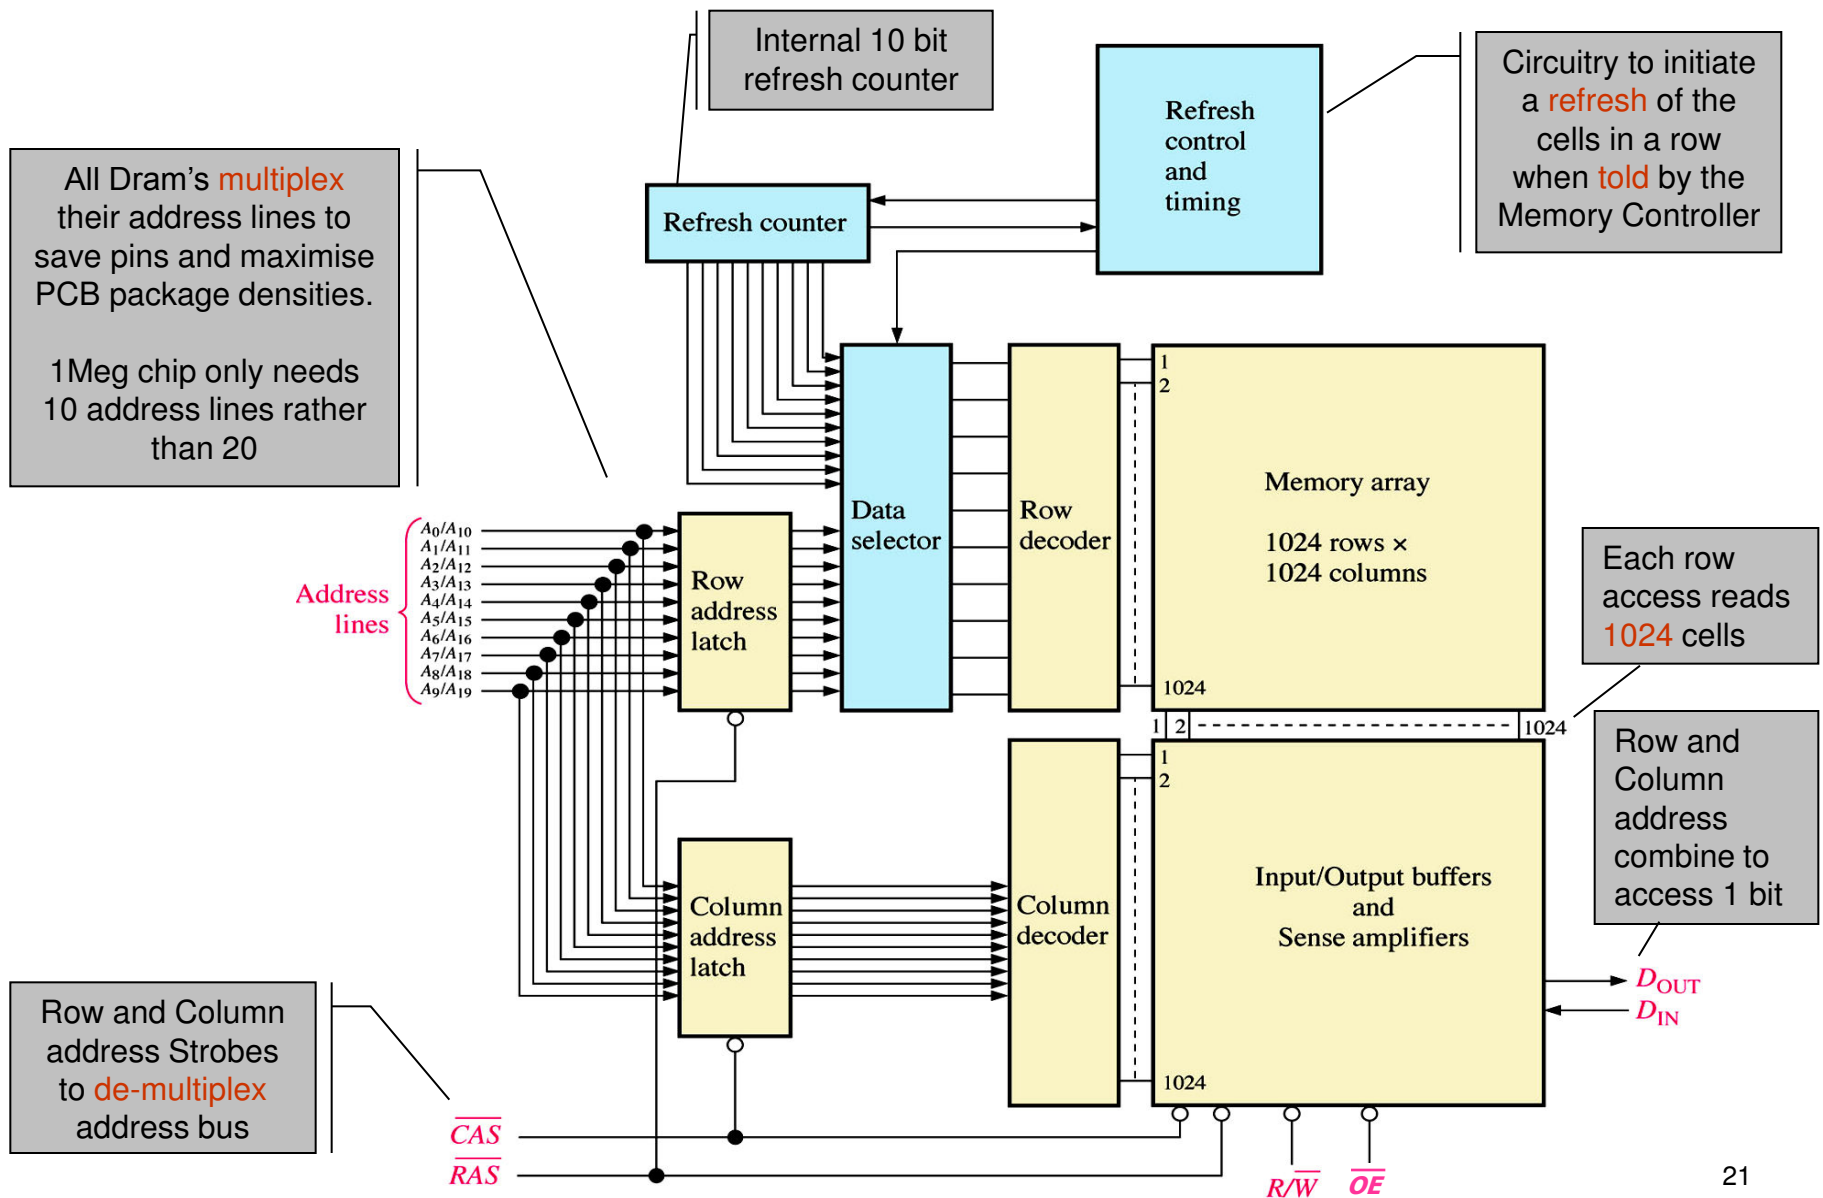
# Problems with Dram Technology

## Corruption of Data

- Dram cells can be **corrupted** by stray **Cosmic** particles passing through the device which can cause enough **ionisation** to corrupt the charge stored on the capacitor.
- This corruption is referred to as a **soft-error** because the damage to the cell is not permanent but leads to a **1** being incorrectly read as a **0**.
- In high integrity (e.g. **military**) or high availability (e.g. **server**) systems, this can be a real problem since the likelihood of soft errors arising is proportional to time.
- In such situations, additional external circuitry may be required to **detect** and **correct** the data.
- More recently special **ECC** memory chips have become available. These however are much more expensive (*see later lecture*).




- First generation Drams were '*Asynchronous*' devices with no clocked circuitry. An example architecture of a **1M x 1bit** device is given below.



# Example Asynchronous Dram (1M x 4 bit Device)

## OKI 514400

- Single 5 volt supply
- 1Meg x 4bit data
- 10 Address Lines
- 4 Bi-directional Data Lines
- 60ns Access Time
- Internal Refresh Counter
- Refreshing required for all 1024 rows every 128ms.



Pin Name	Function
A0–A9	Address Input
$\overline{\text{RAS}}$	Row Address Strobe
$\overline{\text{CAS}}$	Column Address Strobe
DQ1–DQ4	Data Input/Data Output
$\overline{\text{OE}}$	Output Enable
$\overline{\text{WE}}$	Write Enable
V <sub>cc</sub>	Power Supply (5 V)
V <sub>ss</sub>	Ground (0 V)

# A Simplified Asynchronous Dram Read Cycle

**Step 2:** The **RAS\*** signal is taken low. The device latches the row address internally

**Step 1:** A row address is multiplexed onto the Dram address lines

**Step 4:** The **CAS\*** signal is taken low. The device latches the column address internally

**$t_{RC}$  Cycle (min):** The time between successive accesses of the device

**$t_{RP}$  Row Precharge time**

RAS\*

CAS\*

Address

W\*

Data out from memory

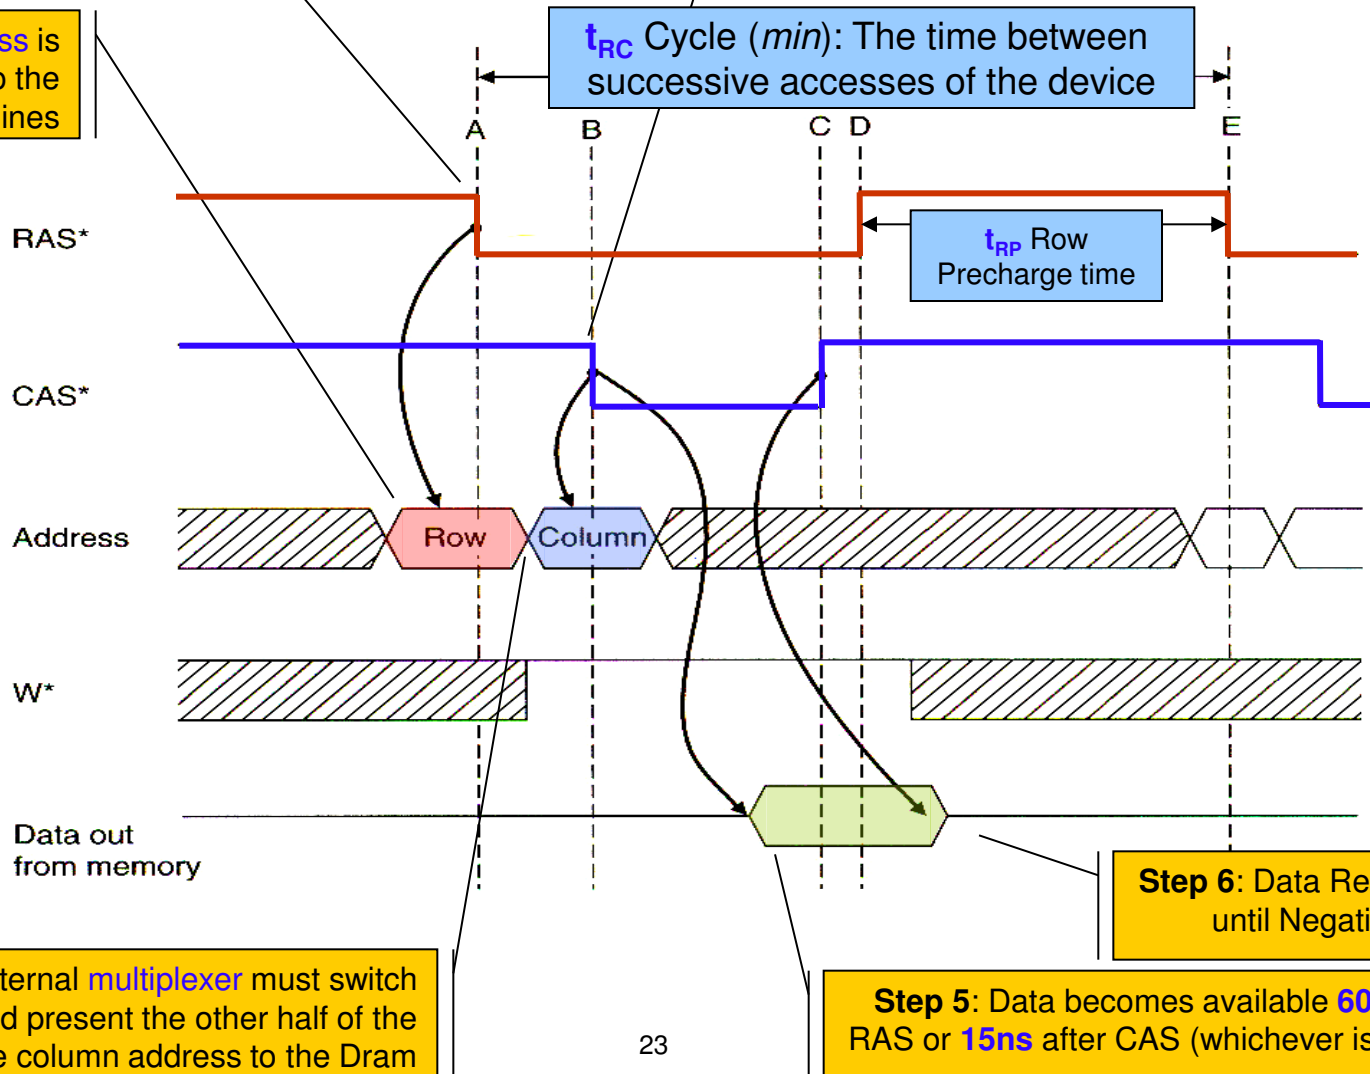
Row

Column

**Step 6:** Data Remains valid until Negation of CAS\*

**Step 3:** External multiplexer must switch over and present the other half of the address, the column address to the Dram

**Step 5:** Data becomes available **60ns** after RAS or **15ns** after CAS (whichever is slower)






# Dram Refreshing

## What, Why and How

- Refreshing arises because the **charge** on the **capacitor** of the dram cell **leaks away**.
- Refreshing '**tops-up**' the charge before the data is lost due to leakage.
- Few memory chips refresh completely by themselves. External hardware is usually required if only to trigger the device into performing the refresh at the correct time.
- Refreshing involves **reading every cell** in the chip (*which implicitly means that the charge on the cell gets topped up*) within a period specified by the manufacturer, (typically **64mS**)



# Dram Refreshing

- A **single refresh operation** will refresh **all cells** sharing the **same row address**, (1024 of them in a 1Meg chip) thus we do **not** have to refresh all cells *individually*.
- For a 1Mbit chip such as the **OKI 514400** comprising 1024 rows, this means that **every row has to be refreshed** at least once during the specified **refresh** period.
- Most modern Dram chips simplify the task of refreshing by having a **refresh counter** fabricated onto the chip to keep track of which Row is to be refreshed next (see page 19).
- During a refresh, the CPU must **not** be permitted to **access** the memory chip.

# Dram Refreshing

- Most Drams can be **physically** refreshed in **variety of ways**, depending upon the complexity of the controller you wish to build and the dram chip itself.

## Refreshing by **Reading** :-

As we saw in a previous lecture, reading a cell **automatically refreshes** not only *that* cell, but also **all other cells** sharing the same Row address.


Provided the CPU, during the execution of its program, accesses all row addresses within the chip within the given refresh period of say **128ms**, then all cells in the chip will be refreshed automatically.

The flaw in this approach is that there is no guarantee, given the way programs jump around within loops etc. that all row addresses will be issued within the refresh period so it's not a reliable method (*unless you do it as part of an Interrupt service routine designed to refresh this way*) but it work well in **graphics cards** – can you think why?

# Dram Refreshing

## RAS\* Only Refreshing


- The very 1<sup>st</sup> generation of Drams used **Ras\*** only refreshing. Here a **row** address is presented externally to the chip and **Ras\*** is asserted, without the **Cas\***. This is what distinguishes a refresh from a normal access.
- The **row** address is typically maintained via an **external counter**, which is **multiplexed** onto the dram address lines when a refresh operation is initiated, and incremented afterwards, i.e. more external circuitry.



# Dram Refreshing

## CAS\* before RAS\* Refreshing

- Later 1<sup>st</sup> Generation devices (such as the OKI 514400) had the row counter and multiplexer integrated onto the data chip (reducing external chip count).
- A new refresh protocol was introduced called, **CAS before RAS** (*the opposite of a normal access*).
- The internal row counter is auto incremented afterwards.







# Dram Refreshing

## Downsides of Refreshing

- Energy consumption: Each refresh consumes energy.
- Performance degradation: DRAM rank/bank unavailable while refreshed.
- QoS/predictability impact : (Long) pause times during refresh.
- Refresh rate limits DRAM capacity scaling (see graph below).



**Performance Impact  
due to Refreshing**



**Energy Impact due  
to Refreshing**

# Dram Refreshing

- **Observation:** Most DRAM rows can be refreshed much less often without losing data.
- **Key idea:** Refresh rows containing weak cells more frequently, other rows less frequently. This can be done through **Profiling**, i.e. test retention time of all rows and use this data in conjunction with an intelligent dram controller that only refreshes rows when needed.
- **Tricky** as retention time of drams changes with **time/temperature**, may need to profile for many days/weeks to find optimum refresh rate.

## CPU Signals

Circuit has to  
- Decode Address

- Generate  
**RAS/MPLX/CAS**  
timing for a normal  
CPU Bus Cycle and  
provide **DTACK** back  
to the CPU.

- Initiate periodic  
**Refresh** operations

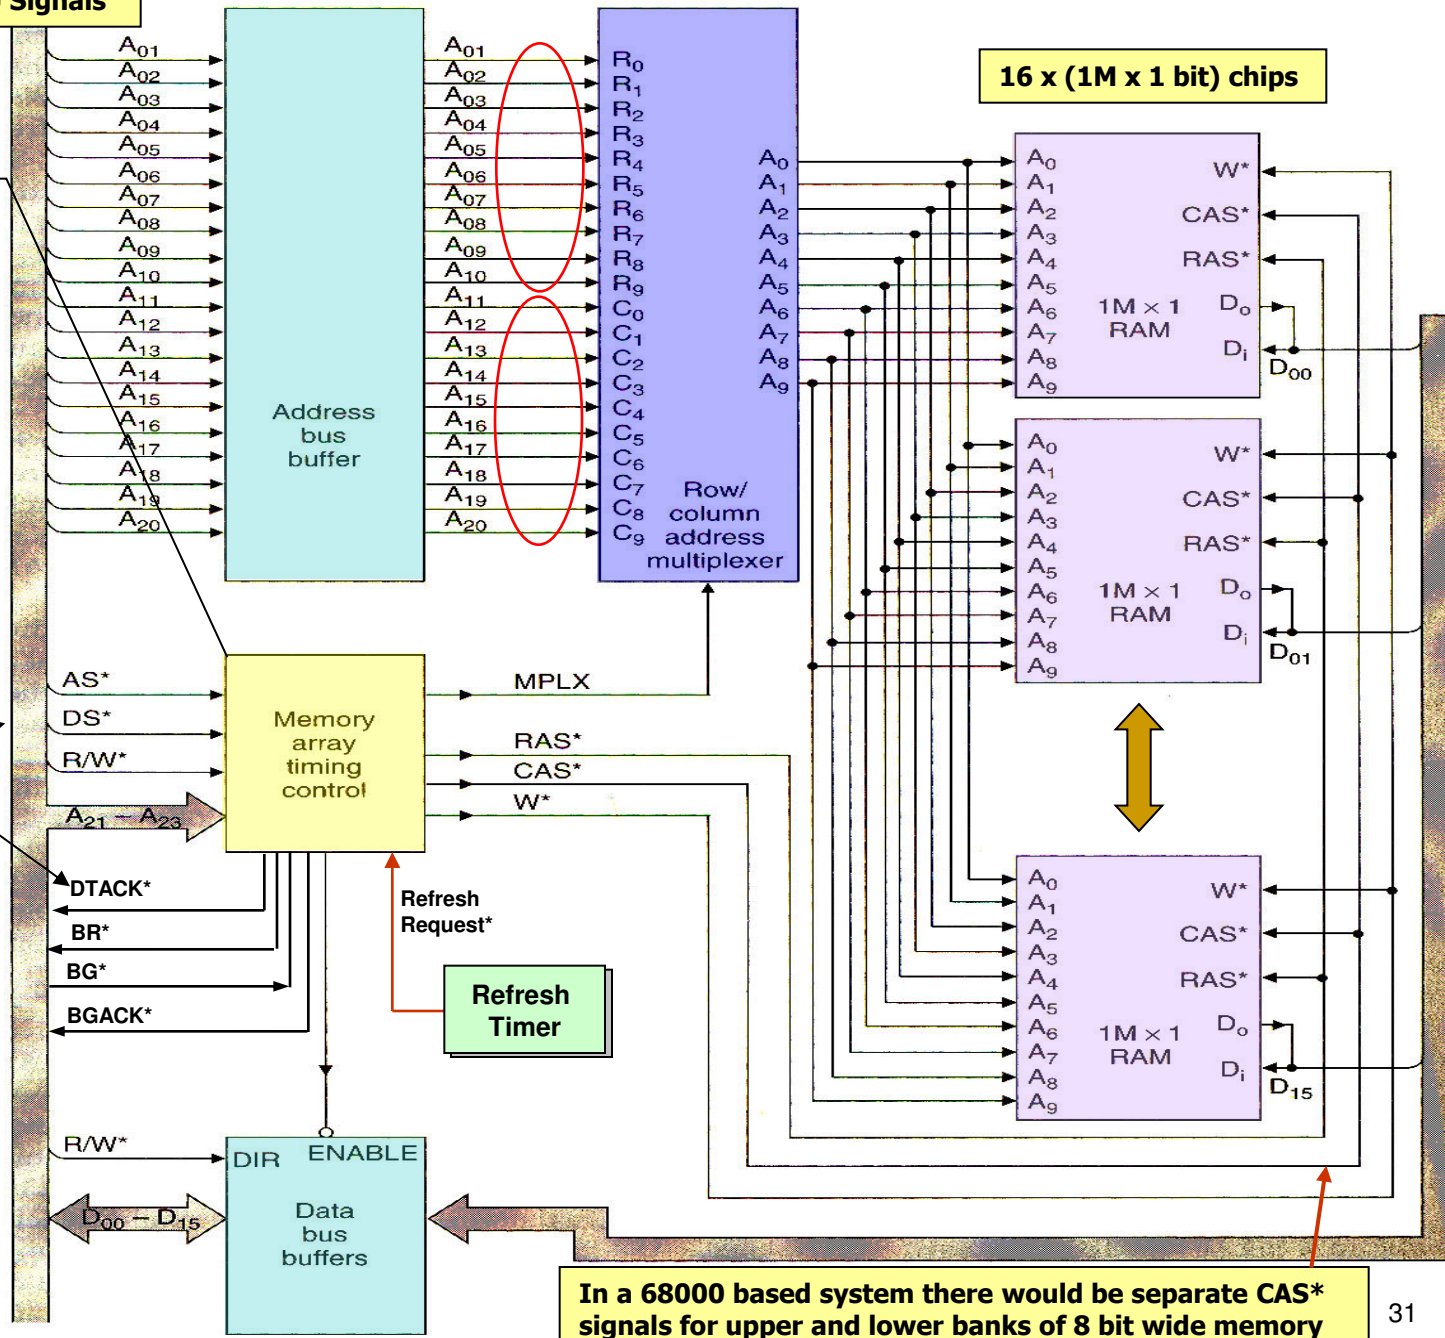
- Take control of the  
bus during a refresh

**DS\*** means **LDS\*/UDS\***

**DTACK\*** can control  
wait state generation

Bus request/bus  
grant Signalling to  
deal with refreshing

Takes control of the  
Address and Data  
bus from the CPU  
while a refresh  
operation is  
performed




In a 68000 based system there would be separate **CAS\*** signals for upper and lower banks of 8 bit wide memory

# Dram Vs CPU Speed

- In the early days of Processor and Dram development, CPUs could run code **directly** from Dram, because their speeds were similar.
- However, the illustration below shows that post 1980, CPU speed had raced ahead of Dram speed. Consequently, running programs or accessing data directly from Dram incurred significant delays via the associated **wait states** imposed by the slow Drams, so all that super fast CPU speed was wasted.

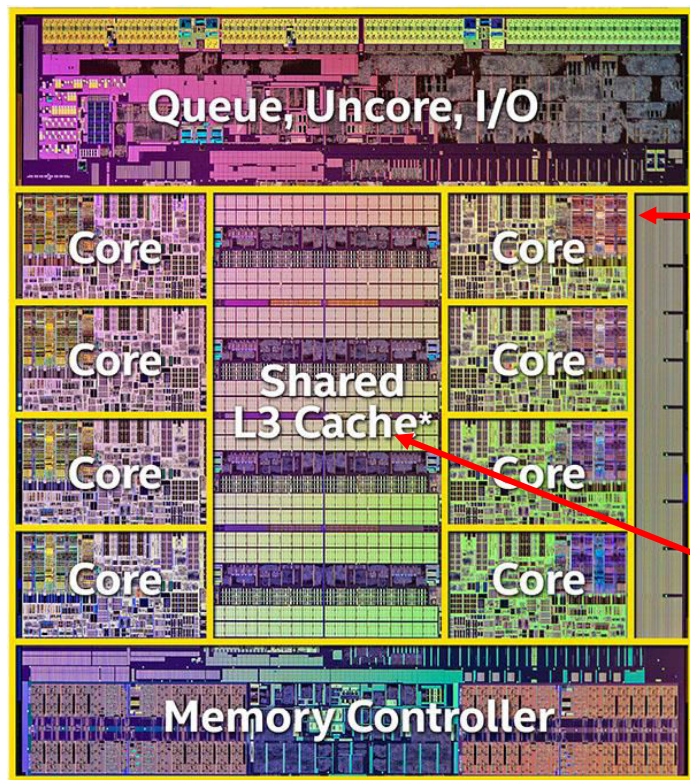
## ■ Processor vs Memory Performance





# Using Caches to Improve Memory Speed

- To combat this, processors were developed to utilize “on-chip” caches to store “local copies” of information maintained in Dram. A cache based CPU could access **data** and **instructions** at full CPU speed *without* incurring “wait state” penalties, provided of course the data/instruction had been loaded into cache in the first instance.
- The bigger the cache, the better the “hit” rate and the better the performance.



## (2015) 8 Core I7 processor with shared Cache and Memory controller

Each core has a 32kB Instruction and 32kB Data Cache (**Level 1 Cache** – full speed).

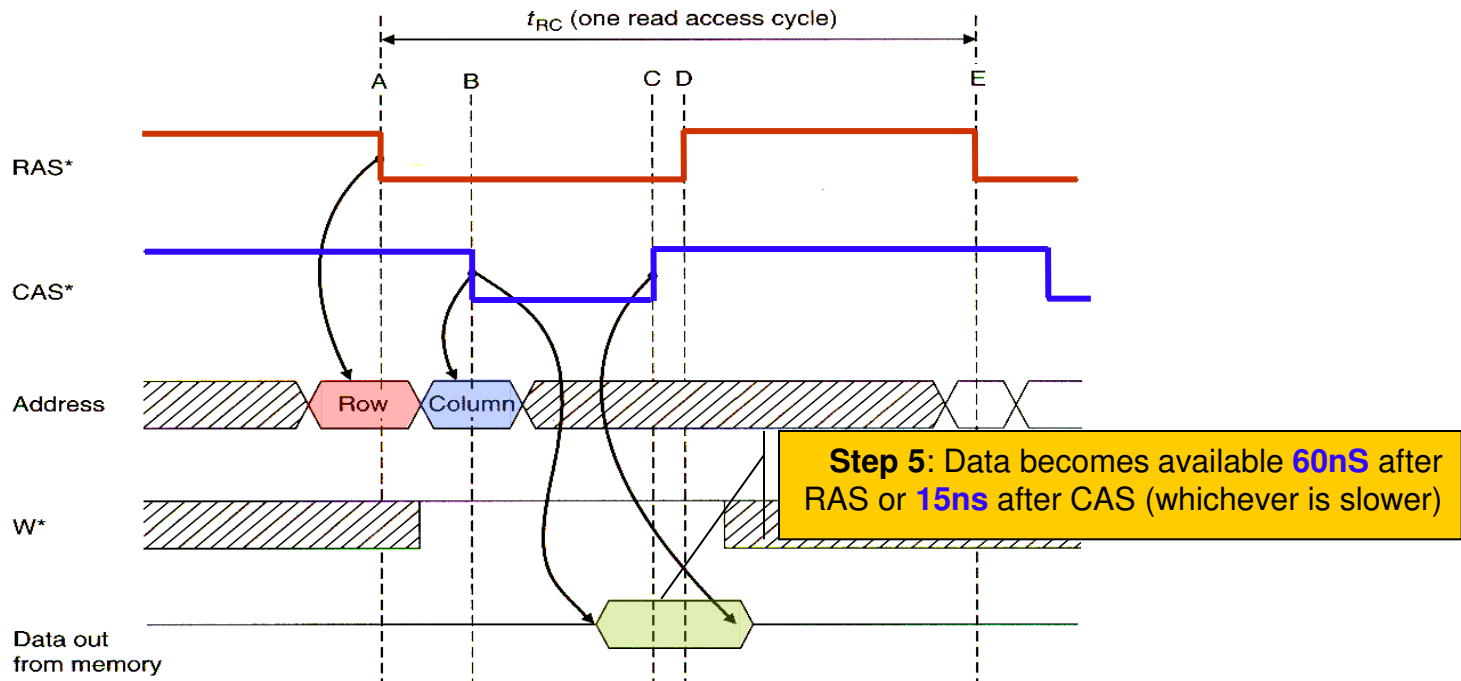
Each core also has 256kB of **Level 2 cache** which runs slower but is still faster than Dram and caches data and code local the that core.

On Chip 20MB shared Level 3 Cache (giving 2.5MB per core)

Cache/Dram controller

# Dram Developments - Fast Page Mode Drams

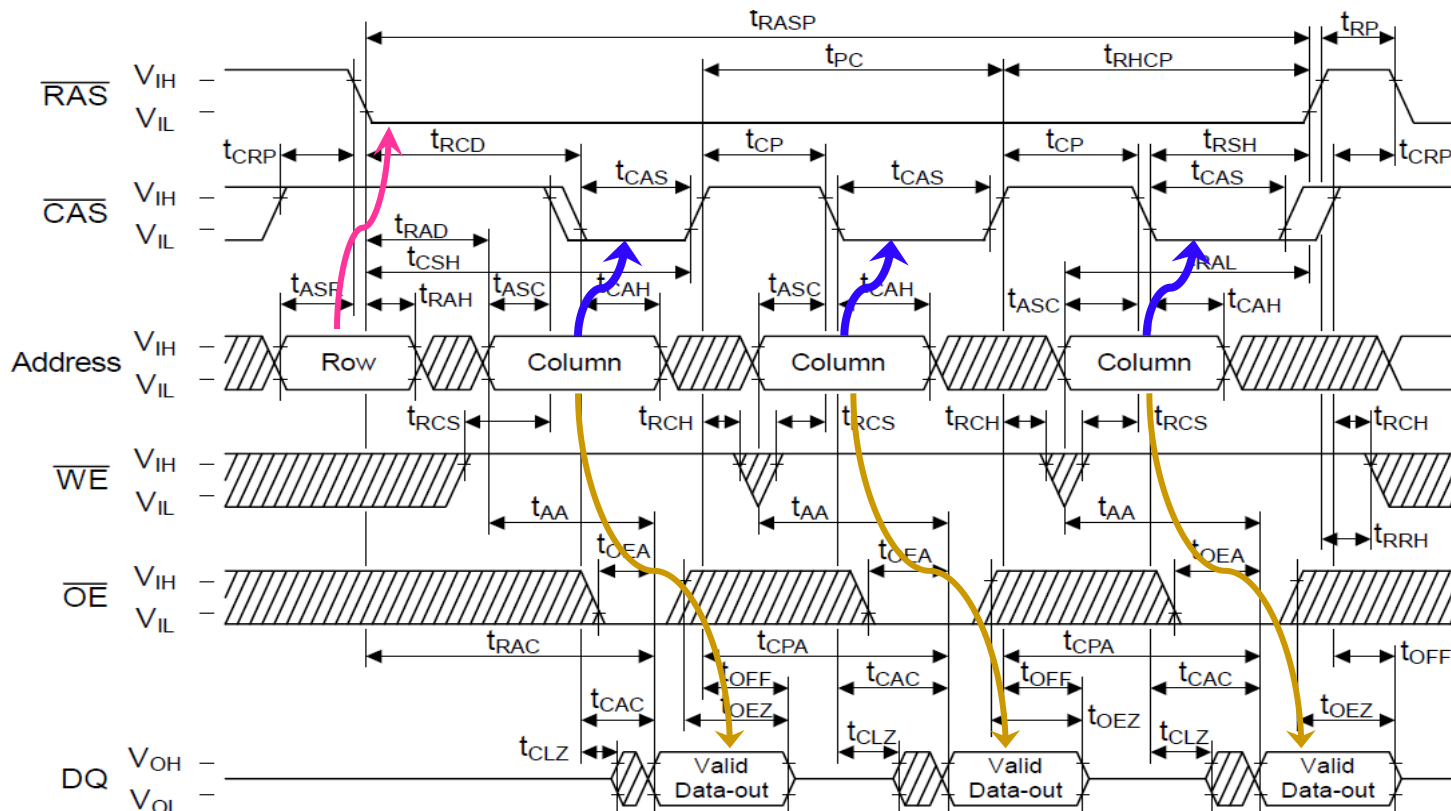
- The big problem then was how to get data from Dram into the cache as fast as possible, i.e. “fill the cache”. First generation Drams needed both a **Ras\*** and a **Cas\*** signal each time memory was access making them slow.



- However by noting that the access time from **Ras\*** (**60ns**) was significantly **longer** than that from **Cas\*** (**15ns**) a reduction in access times could be obtained by using data already in the Dram row buffer.
- This lead to the development of **Fast Page Mode Drams** in the late 90's

# Fast Page Mode Asynchronous Drams

- In page mode, the row buffer of the DRAM can be kept "open" by holding **RAS** low while performing multiple reads or writes with separate **CAS** pulses/addresses.
- A simplified timing diagram is given below illustrating the concept.
- Now a cache/memory controller could be designed to take advantage of this and improve the fill speed of a cache by perhaps **300%**. New data is available every **40ns** using this approach (vs every **110ns** using the RAS/CAS approach).



# Fast Page Mode Asynchronous 64k Dram Architecture

