


# Synchronous Dram (SDRAM)

- As we saw in the last lecture, a huge gulf had opened up between the Multi-GHz speeds of the fastest CPUs and the slow speed of conventional Async Drams. The graph below shows that CPU performance has doubled every **18 months**, but dram performance has only doubled every **10 years**.
- Fast Page Mode Drams** and **on-chip caches** helped, but dram latency has proved to be a real bottleneck to processor performance. Something had to change.

## ■ Processor vs Memory Performance




# Synchronous Dram (SDRAM)

- To increase performance, a new generation of **Synchronous** Drams were developed that could 'burst' multiple items of data from **successive addresses** (i.e. the contents of a dram row buffer) onto the data bus using a high speed **clock** synchronized to the CPU's memory controller. Bursting would help to fill the CPU cache much faster.
- For example, a dram with a **1024 byte row buffer** could have its row contents clocked out at say 1Ghz like a shift register.
- **Pipelining** improves this even more by allowing a sophisticated dram controller to issue new commands to a Synchronous Dram memory chip before the memory chip has finished dealing with the last command, i.e. it allowed for read/write requests to overlap in an attempt to keep the data transfer flowing
- To take advantage of this new technology required a new generation of more sophisticated **Dram Controllers**.

# Synchronous Dram (SDRAM)

- Simplified Synchronous Dram timing diagram.
- After presentation of a row and column address, data is pumped or “**burst**” out of the Dram (*via a shift register*) under the control of a high speed clock.
- The CPU then grabs the data with each clock cycle and transfers it to the

## SDRAM Read



Note that no new addresses are needed after the initial row and column address have been presented

A burst of 4 bytes of data. At 1GHz this equates to a new byte every 1nS vs 40ns with Fast Page mode dram. This is equal to 1GByte per second

# Synchronous Dram (SDRAM)

## Synchronous Dram Commands

- Unlike Asynchronous Dram, the new generation of SDrams require that a **command** is issued with each **POSITIVE** edge of the clock indicating the type of **operation** to performed, such as **activate**, **read**, **write**, **stop burst**, **refresh** etc.

These commands are entered via the **CS\***, **WE\***, **RAS\*** and **CAS\*** signals on the Sdram in conjunction with an address/bank value.

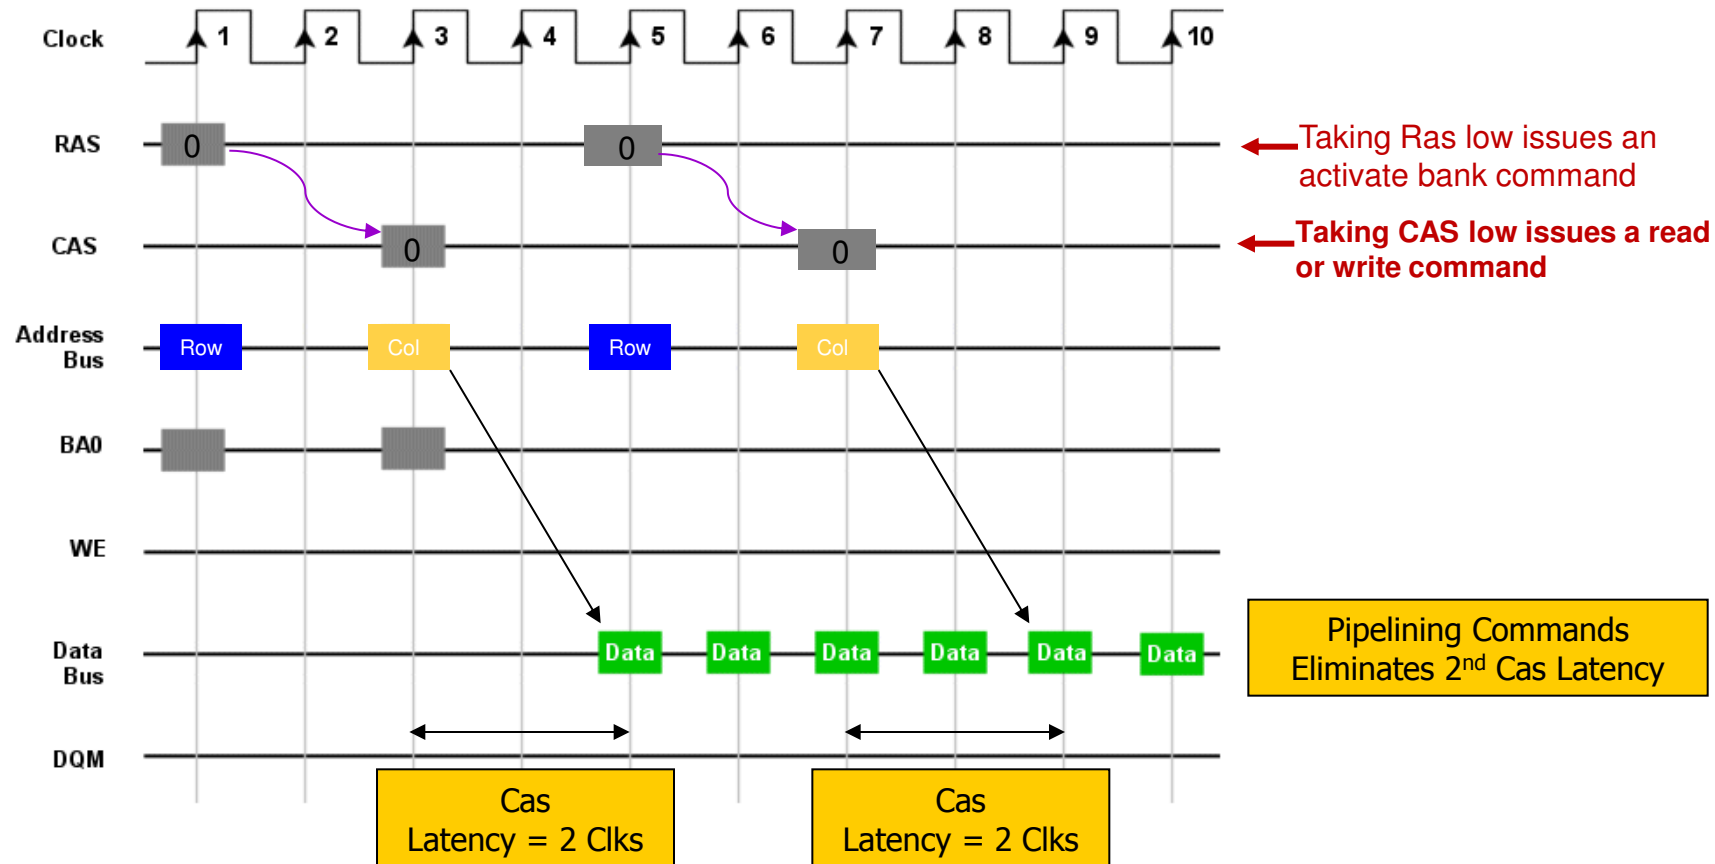
- Unless you intend to issue a new command, you must present a NOP command with each clock

NAME (FUNCTION)	CS#	RAS#	CAS#	WE#	ADDR
DESELECT (NOP)	H	X	X	X	X
NO OPERATION (NOP)	L	H	H	H	X
ACTIVE (Select bank and activate row)	L	L	H	H	Bank/Row
READ (Select bank and column, and start READ burst)	L	H	L	H	Bank/Col
WRITE (Select bank and column, and start WRITE burst)	L	H	L	L	Bank/Col
BURST TERMINATE	L	H	H	L	X
PRECHARGE (Deactivate row in bank or banks)	L	L	H	L	Code
AUTO REFRESH or SELF REFRESH (Enter self refresh mode)	L	L	L	H	X
LOAD MODE REGISTER	L	L	L	L	Op-Code

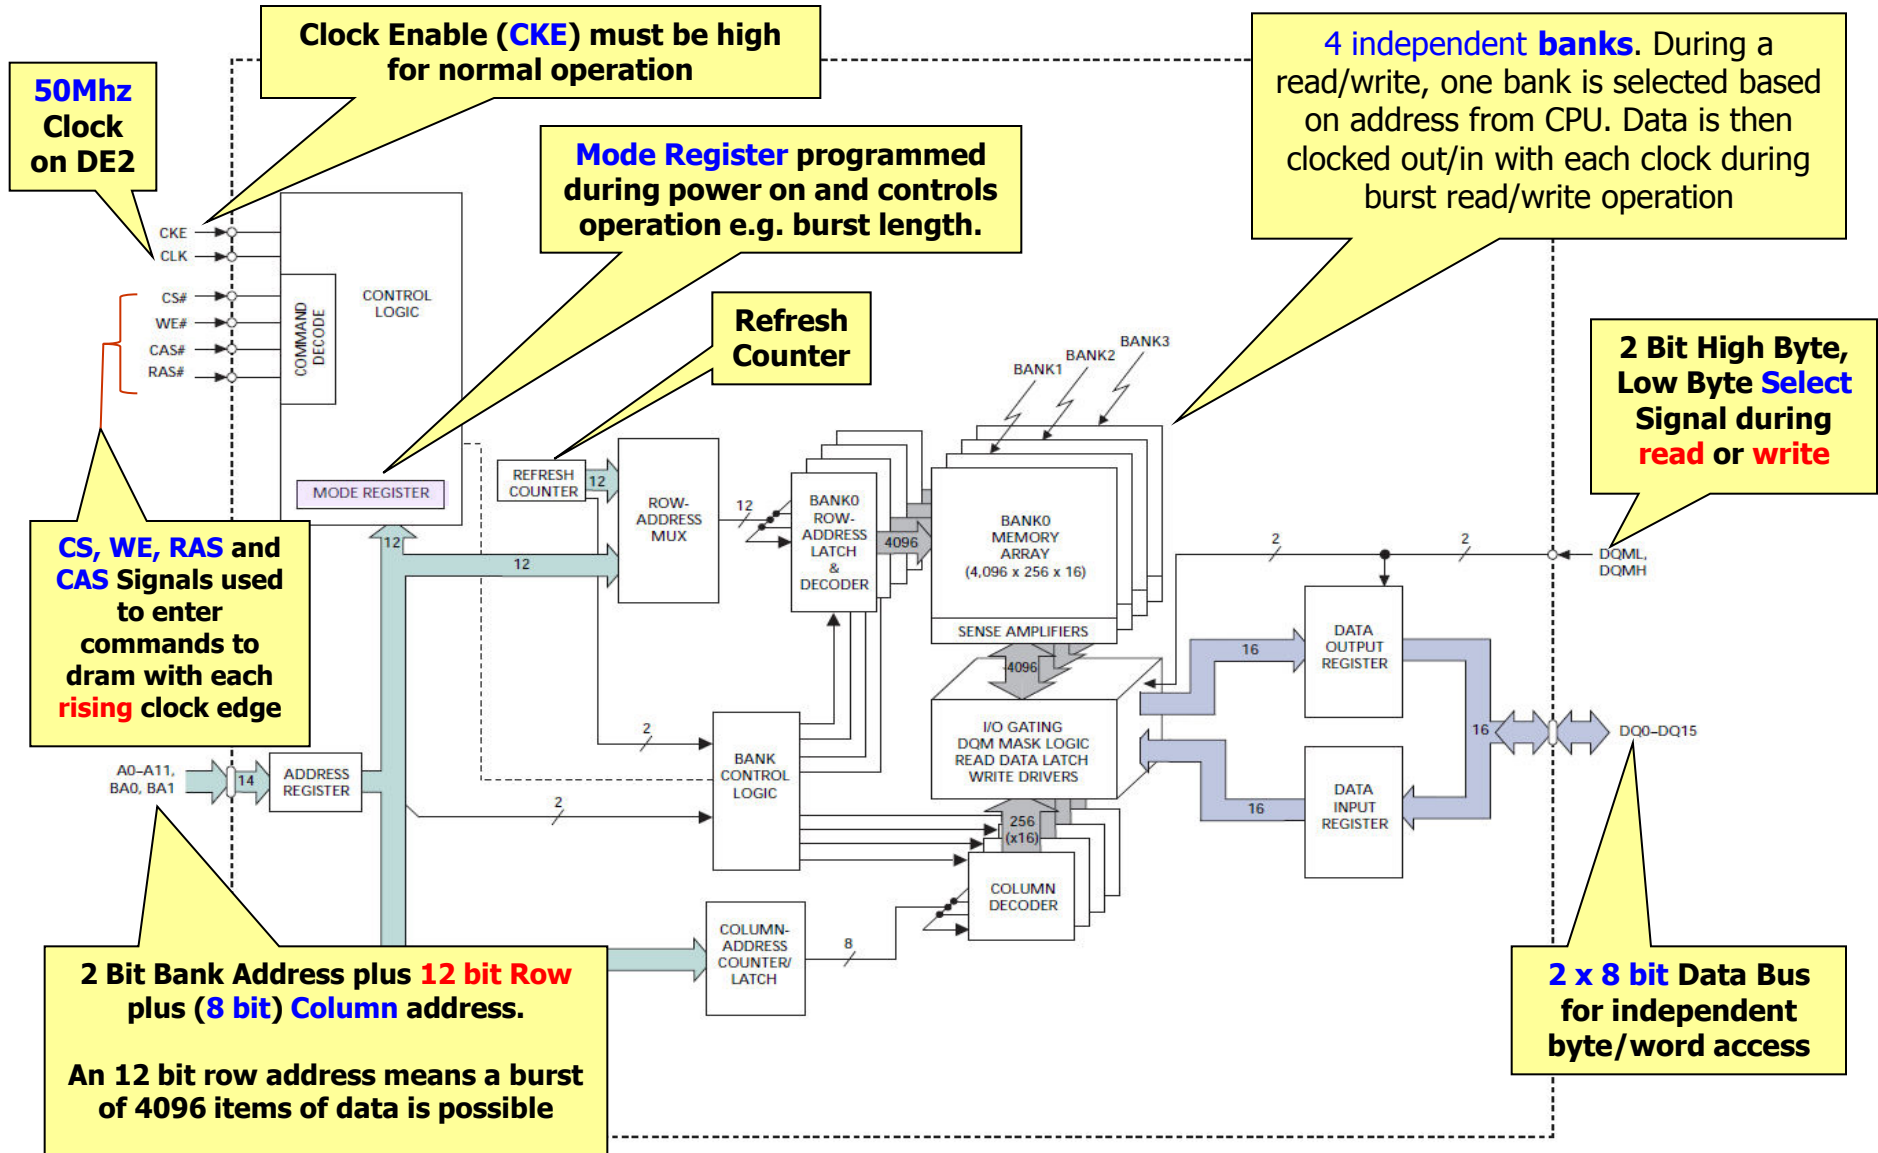
# Synchronous Dram (SDRAM)

## Pipelining Commands


- Pipelining allows a a dram controller to issue new dram commands before the previous one has completed, this allowing continuous data transfer without incurring subsequent CAS latencies. This is shown below where two consecutive bursts of 4 bytes are requested and occur back to back.



# Typical 4M \* 16 bit SDRAM Architecture (as per DE2)



# Typical 4M x 16bit SDRAM Pin Out



Access time from CAS is a fixed delay e.g. 5.4ns as shown below. With higher clock speeds (e.g. 143Mhz, we need to wait more clock periods after CAS (the **CAS Latency**) to build up the 5.4ns delay before we can get at the 1<sup>st</sup> item of data, but access to subsequent data is faster overall .

## KEY TIMING PARAMETERS

SPEED GRADE	CLOCK FREQUENCY	ACCESS TIME		SETUP TIME	HOLD TIME
		CL = 2 *	CL = 3 *		
-7E	143 MHz	–	5.4ns	1.5ns	0.8ns
-75	133 MHz	–	5.4ns	1.5ns	0.8ns
-7E	133 MHz	5.4ns	–	1.5ns	0.8ns
-75	100 MHz	6ns	–	1.5ns	0.8ns

12 Bit Address Bus:

A0-A11

16 Bit Data Bus:

DQ0-DQ15

2 Bit Bank Address:

BA0-BA1

2 Byte Select signals:

DQMH/DQML

*(The Byte select signals are used to activate one or other or both halves of the data bus. This mirrors the operation of the 68k's UDS\* and LDS\* signals)*


# SDram Chip Commands

- An SDram chip first has to be programmed before use with a **Load Mode Register** command.
- **Commands** are **presented** to the SDram chip by the SDram controller and sampled on each **rising edge** of the clock.
- Commands are interpreted by the signal levels on **CS#**, **RAS#**, **CAS#** and **WE#** (see below).
- **Activate commands** are issued to activate a **Bank** within the chip while at the same time presenting a **ROW** address to that bank (*this is analogous to the old asynchronous dram act of issuing RAS*).
- **Read/Write** commands require a **bank** and a column **address** which, when combined with the previous activate command uniquely identifies **bank/row and column location** within the chip (*this is analogous to the old asynchronous dram act of issuing CAS*).
- **Pre-charging** closes a bank/row in preparation for opening another with an **activate** command.
- **No Operation** (NOP) corresponds to a do noting or “carry on as before”.
- **Auto Refresh** operations can be **initiated** by the SDram Controller.

NAME (FUNCTION)	CS#	RAS#	CAS#	WE#	ADDR
DESELECT (NOP)	H	X	X	X	X
NO OPERATION (NOP)	L	H	H	H	X
ACTIVE (Select bank and activate row)	L	L	H	H	Bank/Row
READ (Select bank and column, and start READ burst)	L	H	L	H	Bank/Col
WRITE (Select bank and column, and start WRITE burst)	L	H	L	L	Bank/Col
BURST TERMINATE	L	H	H	L	X
PRECHARGE (Deactivate row in bank or banks)	L	L	H	L	Code
AUTO REFRESH or SELF REFRESH (Enter self refresh mode)	L	L	L	H	X
LOAD MODE REGISTER	L	L	L	L	Op-Code



## Mode register programmed with data on the lowest 12 bits of Address Bus: A11-A0



Program  
BA0, BA1,  
M11, M10 = "0, 0"  
to ensure compatibility  
with future devices.

M9	Write Burst Mode
0	Programmed Burst Length
1	Single Location Access

M8	M7	M6-M0	Operating Mode
0	0	Defined	Standard Operation
-	-	-	All other states reserved

			Burst Length	
M2	M1	M0	M3 = 0	M3 = 1
0	0	0	1	1
0	0	1	2	2
0	1	0	4	4
0	1	1	8	8
1	0	0	Reserved	Reserved
1	0	1	Reserved	Reserved
1	1	0	Reserved	Reserved
1	1	1	Full Page	Reserved

M3	Burst Type
0	Sequential
1	Interleaved

M6	M5	M4	CAS Latency
0	0	0	Reserved
0	0	1	Reserved
0	1	0	2
0	1	1	3
1	0	0	Reserved
1	0	1	Reserved
1	1	0	Reserved
1	1	1	Reserved

Burst Length and type depends upon the chip and how it is programmed.

Burst Length	Starting Column Address	Order of Accesses Within a Burst	
		Type = Sequential	Type = Interleaved
2	A0		
	0	0-1	0-1
	1	1-0	1-0
4	A1 A0		
	0 0	0-1-2-3	0-1-2-3
	0 1	1-2-3-0	1-0-3-2
	1 0	2-3-0-1	2-3-0-1
	1 1	3-0-1-2	3-2-1-0
8	A2 A1 A0		
	0 0 0	0-1-2-3-4-5-6-7	0-1-2-3-4-5-6-7
	0 0 1	1-2-3-4-5-6-7-0	1-0-3-2-5-4-7-6
	0 1 0	2-3-4-5-6-7-0-1	2-3-0-1-6-7-4-5
	0 1 1	3-4-5-6-7-0-1-2	3-2-1-0-7-6-5-4
	1 0 0	4-5-6-7-0-1-2-3	4-5-6-7-0-1-2-3
	1 0 1	5-6-7-0-1-2-3-4	5-4-7-6-1-0-3-2
	1 1 0	6-7-0-1-2-3-4-5	6-7-4-5-2-3-0-1
	1 1 1	7-0-1-2-3-4-5-6	7-6-5-4-3-2-1-0
Full Page (y)	n = A0-A12/11/9 (location 0-y)	Cn, Cn + 1, Cn + 2 Cn + 3, Cn + 4... ...Cn - 1, Cn...	Not Supported

# Interleaved vs. Sequential Access

Burst Length	Starting Column Address	Order of Accesses Within a Burst	
		Type = Sequential	Type = Interleaved
2	A0		
	0	0-1	0-1
	1	1-0	1-0
4	A1 A0		
	0 0	0-1-2-3	0-1-2-3
	0 1	1-2-3-0	0-1-2-3
	1 0	2-3-0-1	2-3-0-1
	1 1	3-0-1-2	3-2-1-0
8	A2 A1 A0		
	0 0 0	0-1-2-3-4-5-6-7	0-1-2-3-4-5-6-7
	0 0 1	1-2-3-4-5-6-7-0	1-0-3-2-5-4-7-6
	0 1 0	2-3-4-5-6-7-0-1	2-3-0-1-6-7-4-5
	0 1 1	3-4-5-6-7-0-1-2	3-2-1-0-7-6-5-4
	1 0 0	4-5-6-7-0-1-2-3	4-5-6-7-0-1-2-3
	1 0 1	5-6-7-0-1-2-3-4	5-4-7-6-1-0-3-2
	1 1 0	6-7-0-1-2-3-4-5	6-7-4-5-2-3-0-1
	1 1 1	7-0-1-2-3-4-5-6	7-6-5-4-3-2-1-0
Full Page (y)	n = A0-A12/11/9 (location 0-y)	Cn, Cn + 1, Cn + 2 Cn + 3, Cn + 4... ...Cn - 1, Cn...	Not Supported

The choice of **Interleaved** or **sequential** access depends on the type of cache controller on your CPU.

**Sequential** access is easy to understand. The CPU outputs a (column) address for say a read, and the sdram responds with the data at that address followed by the data from **successive** address locations up to the programmed burst length in the Mode register, cycling back to the first item where necessary.

For example, a read from location **3** with a burst length of **8** gives the contents of location **3, 4, 5, 6, 7** before wrapping back to **0, 1** and finally **2**.

**Interleaved** access is very difficult to comprehend at first sight, i.e. what order do we get the data out of the sdram for a particular start/column address and a programmed burst length? It starts off being the same as Sequential, i.e. the address the CPU issues is always the location returned first. However, the data from subsequent locations is determined by **EX-ORing** the Column address bits from the CPU with the internal Sdram burst counter.

For example take a column address (**A2,A1,A0**) of **101** and a burst length of **8**. If we EX-OR **101** with **000**, then **001**, **010**, **011**, **100**, **101**, **110**, and finally **111** we get, **5-4-7-6-1-0-3-2**. **Intel** CPU cache controllers prefer this mode that's why it exists.

# Programming the Mode Register

- Programming the **Mode Register** is performed a.s.a.p. after power on by the SDRAM controller, before the CPU accesses the memory for the first time.
- The Zentel SDRAM chip on the **DE2** board can have a clock latency of **2 clock** periods for a clock running at a relatively slow **50Mhz**.
- Because the **68k** does **not** have a cache and therefore does **not** issue or deal with burst requests, read and write burst lengths should be set to **1** (the choice of *sequential or interleaved access therefore becomes irrelevant*).
- The Zentel / DE2 chip has **12 address** lines so the mode register is set with the following data on address lines **A11..A0** = **001000100000**. i.e. **hex 220**

11	10	9	8	7	6	5	4	3	2	1	0
Reserved	WB	Op Mode	CAS Latency	BT	Burst Length						

- This configures the chip as follows:-

- Read/Write burst length = 1 location,
- Sequential access,
- Cas latency = 2 clock periods

- To **Load the Mode register** we present the above value on the sdram address lines, set **BA0** and **BA1** to be **"00"** and take **CS=0**, **WE=0**, **RAS=0** and **CAS=0** (see page 7). 11

# Power On Initialisation

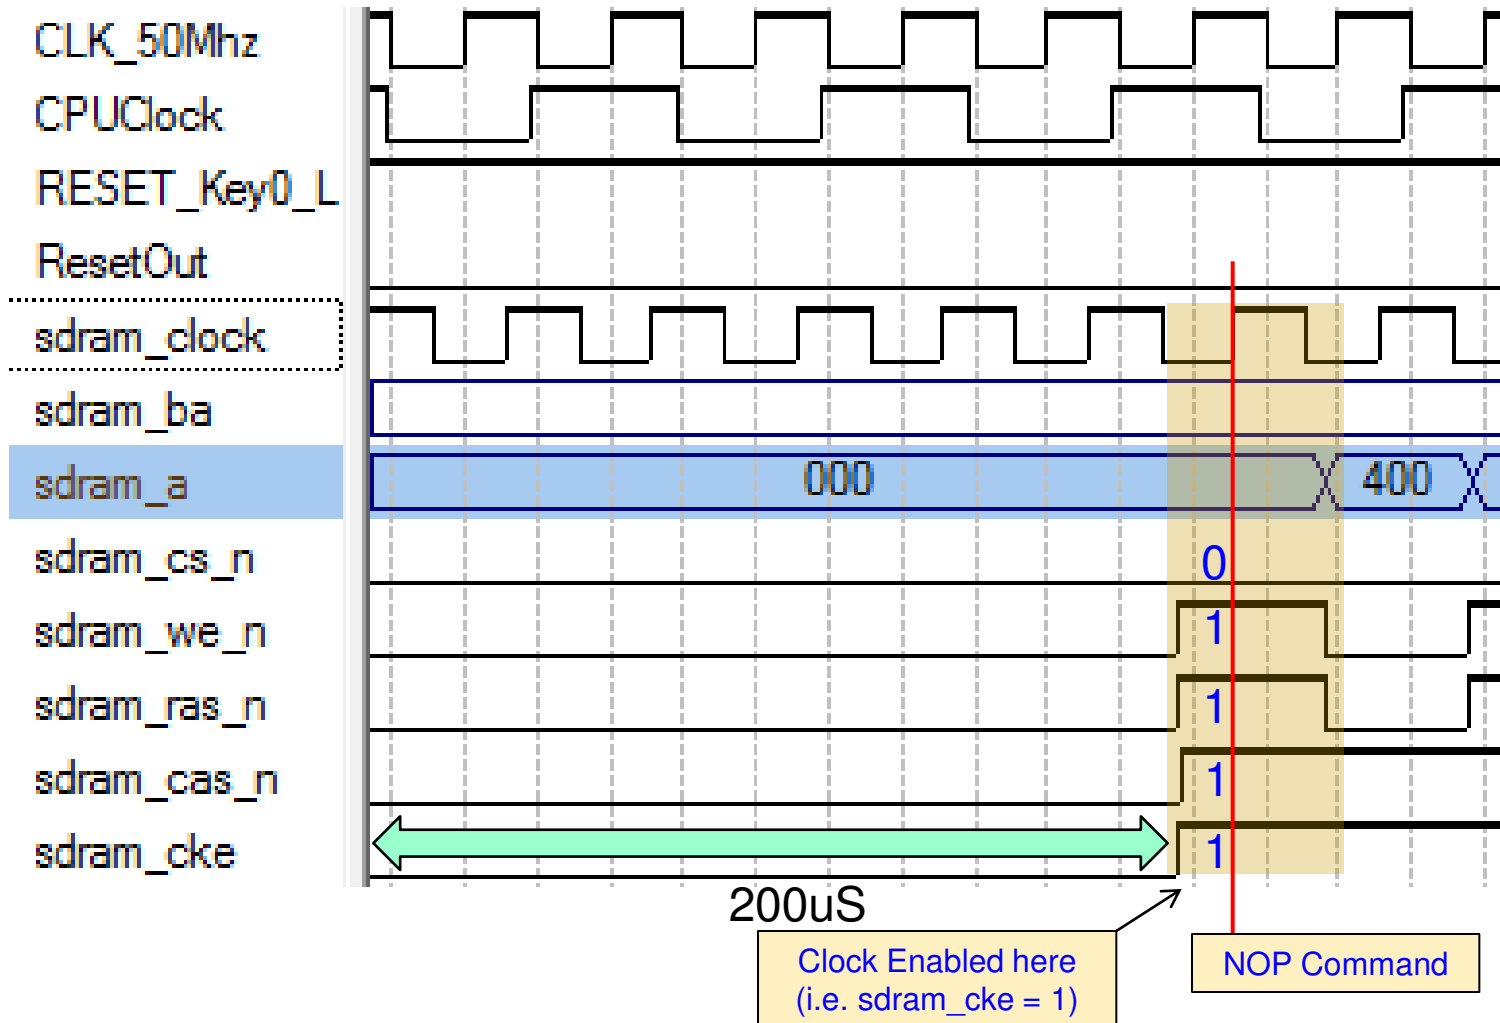
- The **power on** and **initialisation** sequence for the ZENTEL SDRAM is crucial.
- The data sheets say, or imply, from timing diagrams the following sequence.

1. Apply power and a stable clock but maintain the SDRAM **Clock Enable (CKE)** signal in the **INACTIVE** (low) state for **200us**. *In practice the 200us is not essential due to the long power up phase of FPGAs.*
2. After the **200us** period, take **CKE** high (active) and commence with at least 1 **NOP** command to the SDRAM.
3. Next issue a **Pre-charge** command for **all banks** of the devices. This means issuing a **pre-charge** command while at the same time setting address line **A10 = 1** (as per Page 7).
4. Issue 1 **NOP** command after above **Pre-charge** command
5. Next issue 8 or more **Auto-Refresh** commands (Page 7) **with** 3 NOP commands between each Auto-refresh.
6. **Next program** device operation with a **Load Mode Register** Command.
7. Finally issue 3 **NOP** commands.
8. The device is then **ready** for normal use, i.e. read/write operation.

- A simulation of the power on initialisation for an SDRAM is shown over the page, for the working version of the DE2

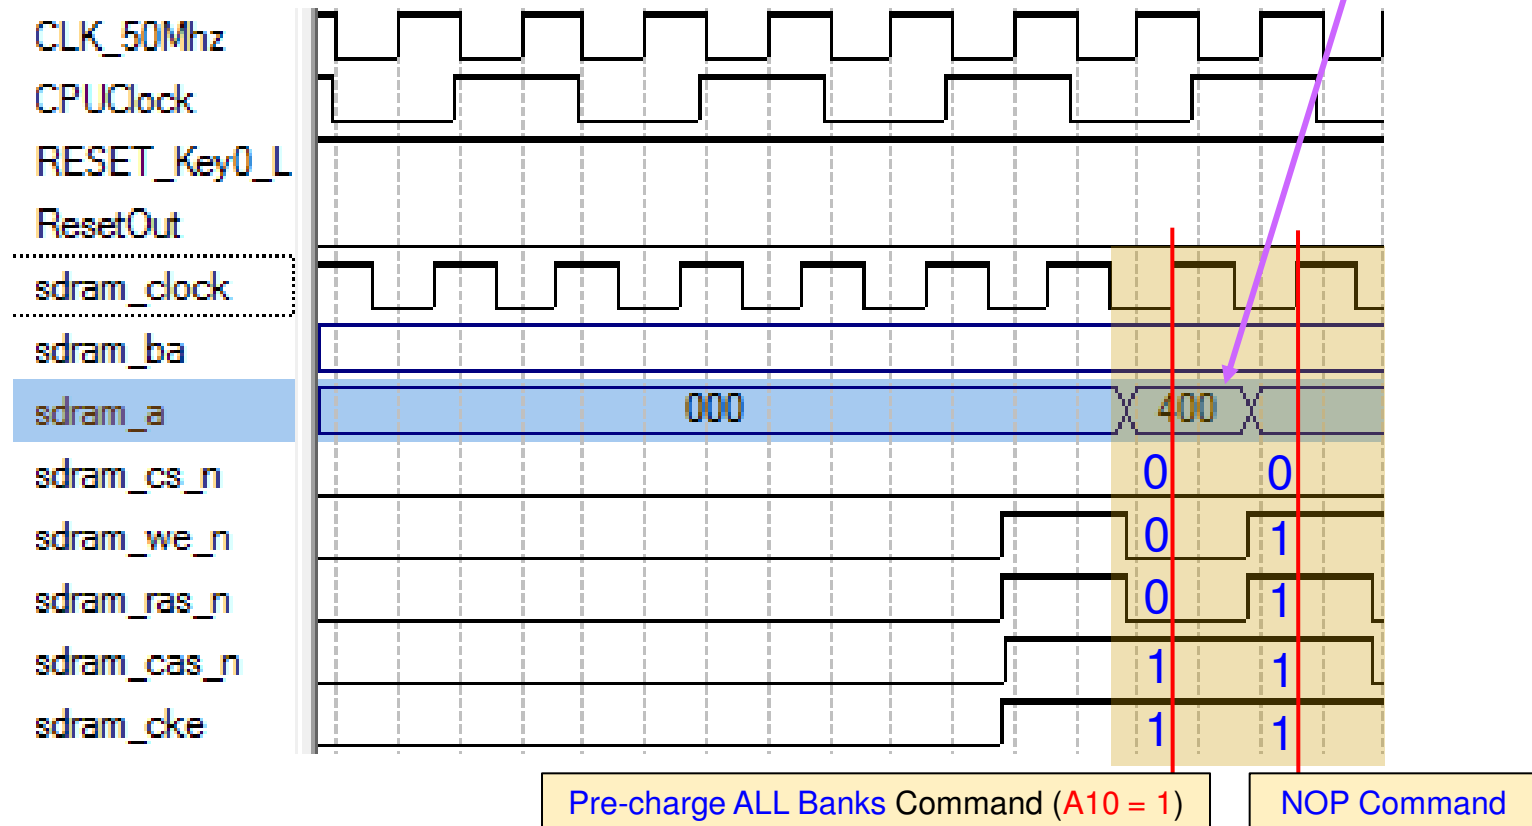
# Power On Initialisation

- Here is the 1<sup>st</sup> and 2<sup>nd</sup> step where **SDram CKE** is taken to logic 1 after **200uS** and we start issuing a **NOP** command, i.e. **CS**, **WE**, **RAS** and **CAS** are all logic 1.



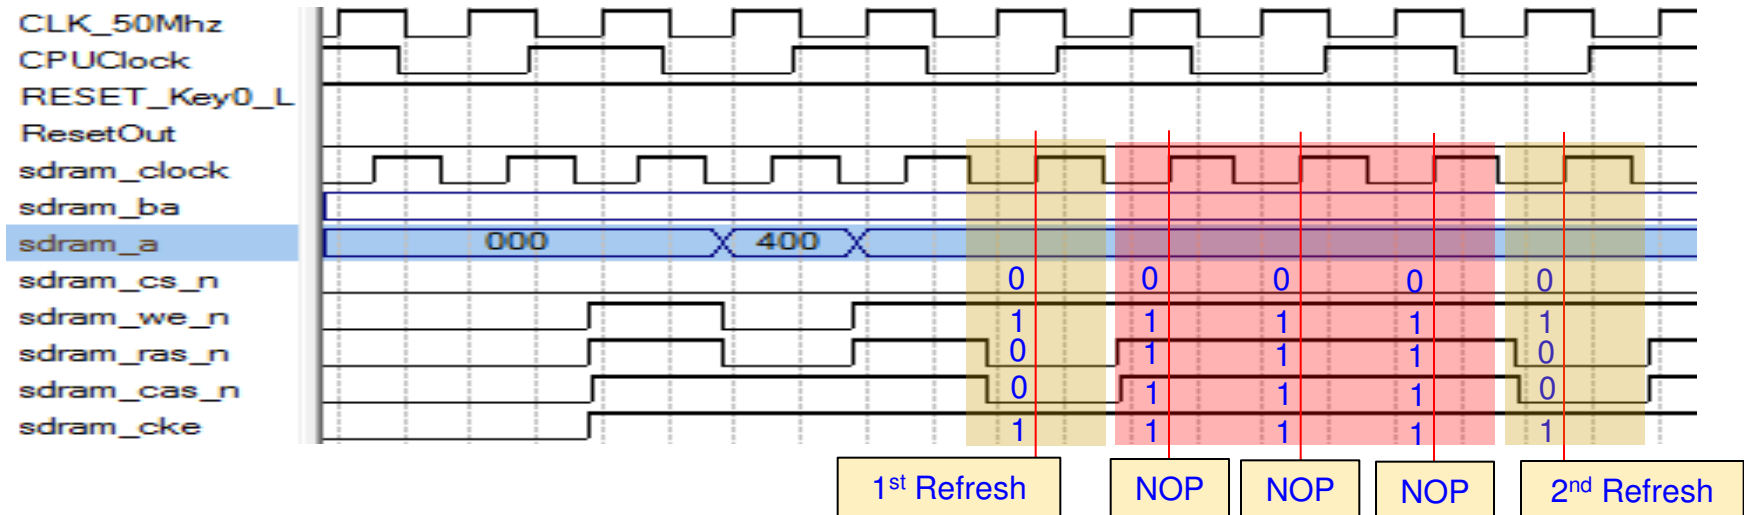
# Power On Initialisation

- Here is the 3<sup>rd</sup> and 4<sup>th</sup> steps where we issue a **Pre-charge** all banks command (i.e. a *pre-charge command* – Page 7 where **CS=0, WE=0, RAS=0 and CAS=1**). This deactivates all banks and rows within the chip. There must be a delay of at least **18ns** after the **Pre-charge** command (the *timing parameter*  $t_{RP}$  – page 38 in the Zentel data sheet) before a new command so for safety follow with 1 **NOP** command (at 50Mhz a new command is issued every 20ns).
- Note **A10 = 1** for **pre-charge ALL banks command**. This explains the hex value **400** we can see on the **12 bit** address bus **A11-A0** below, i.e. **0100 0000 0000**.



# Power On Initialisation

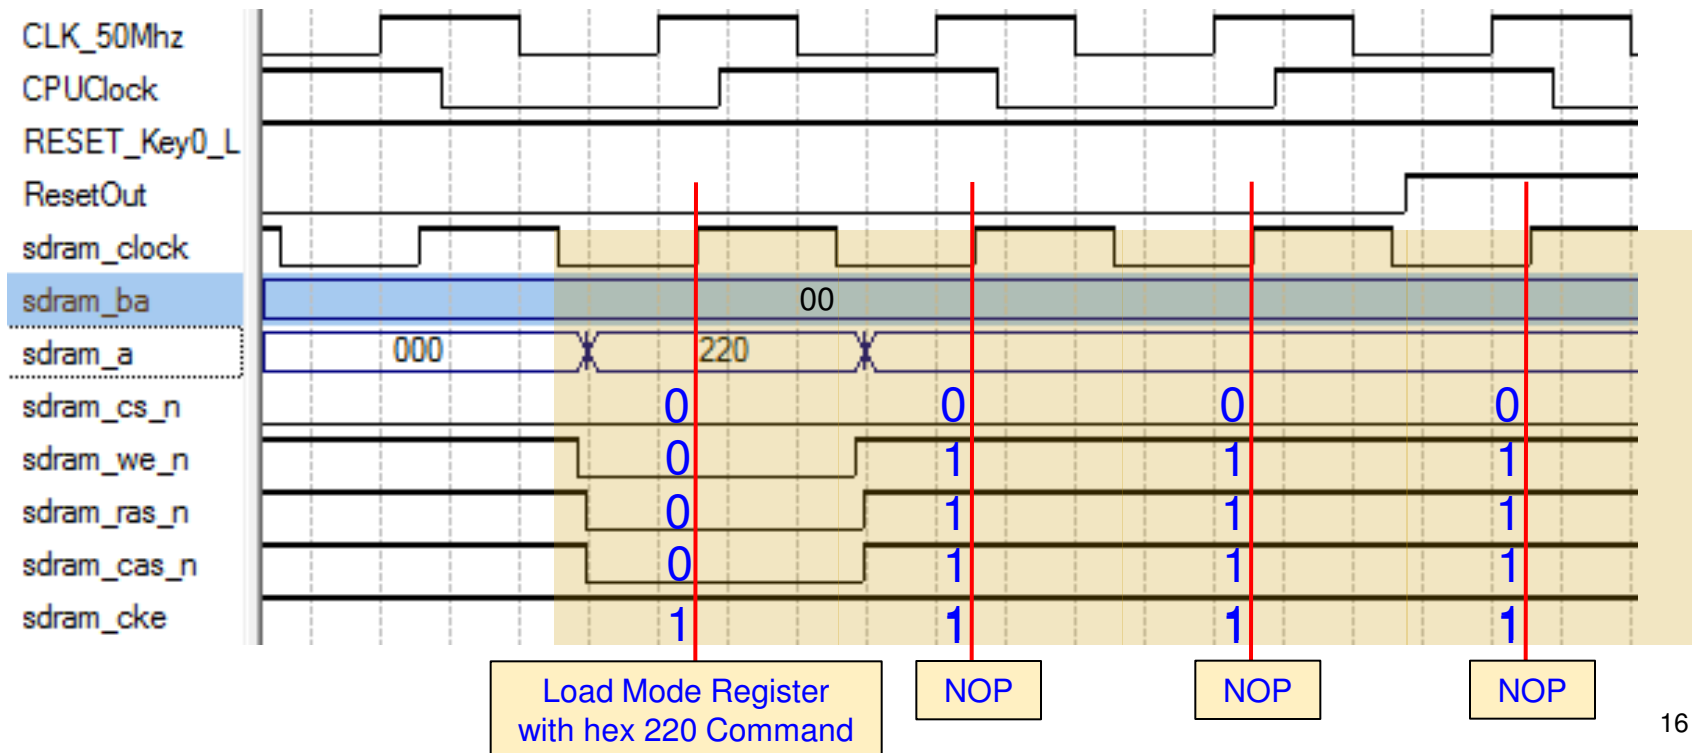
- Here is the 5<sup>th</sup> step where we issue 8 or more refresh commands (see Page 7) i.e. **CS=0, WE=1, RAS=0** and **CAS=0**. The Zentel chip data sheet suggests a min of 8 refresh commands, whereas most other manufacturers chips seem to require only 2.
- I issued 8 in total (*you can only see 2 refresh commands in the screen shot below*).
- Refresh commands on the Zentel chip of the DE2 board require a minimum of 58ns of delay (Time  $t_{RC}$  on Zentel data sheet) *which for safety margin means at least 4 clock periods @ 50Mhz* between them and the next command (*whatever that command may be*), so **always** issue 3 **NOP** commands after **each** refresh command before issuing any further commands to the Dram.



- The last step is to **program the mode register**.

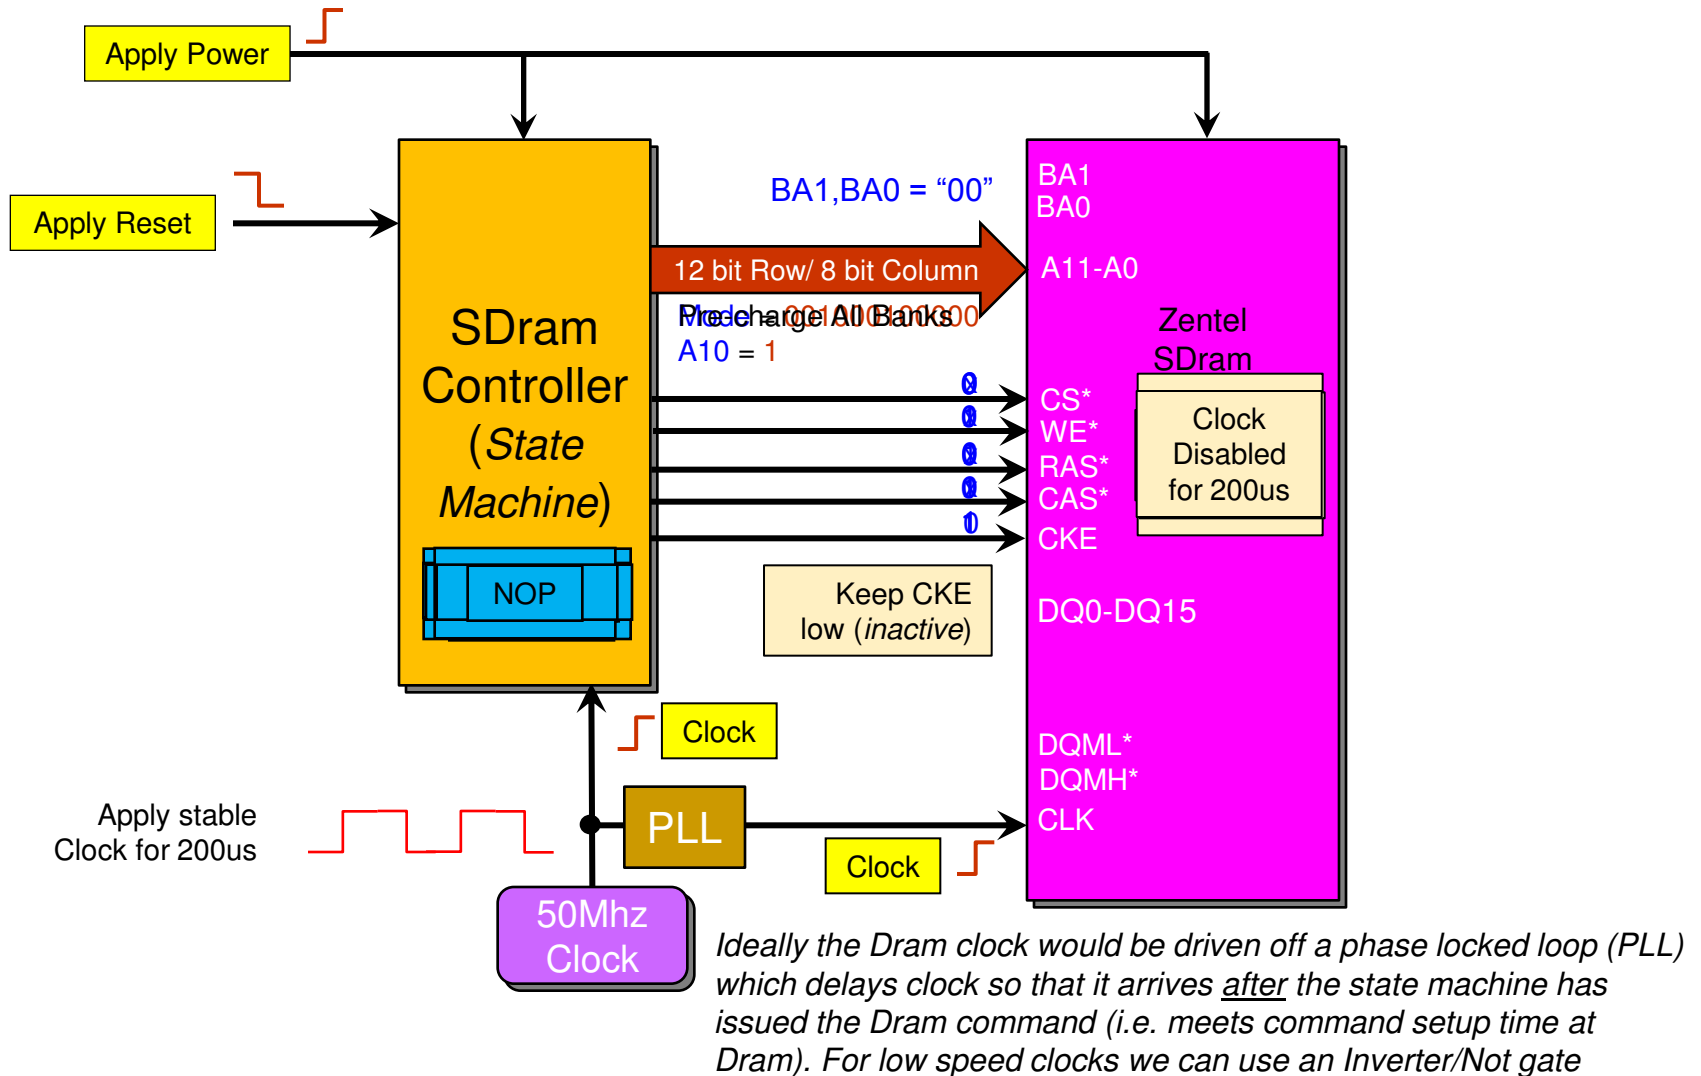
# Programming the Mode Register

- You can see the timing diagram associated with [loading the mode register](#) here. The programming occurs on the rising edge of the SDRAM clock as shown below, where hex [220](#) (see *Page 10*) is presented on the address lines and **CS=0, WE=0, RAS=0** and **CAS=0** and **CKE=1**.
- NOTE:** Bank Address should be “00” when loading the mode register. This is something the Zentel Chip mentions but not all SDRAM chips seem to need this.
- At least **3 NOP** commands must be issued after loading the mode register before the device is ready to accept read or write commands.





# Animation showing Reset and Power On Initialisation



Note this animation shows the sdram controller issuing just 2 refresh commands (followed by 3 NOPs). The Zentel chip says 8 refresh commands are required.

# Power On Initialisation

- How do I issue these commands?
- You design a state machine driven off the same clock as the SDram to drive the SDram signals **CS**, **WE**, **RAS**, **CAS** and **CKE** signals.
- After **reset** this state machine runs through a set of initialisation states to program the SDram chips and then enters an **idle** state (where it issues **NOP** or **refresh** commands to the dram) ready to accept **read** and **write** requests from the processor.
- By driving the sdram clock off a time delayed state machine clock (i.e. the output of the PLL - which can be fine tuned), we can get the state machine to output new commands (i.e. values for **CS**, **WE**, **RAS**, **CAS**, **CKE** signals) on one **edge** of the SDRAM clock, thus allowing plenty of time to meet the set up times at the SDram pins prior to its clock arriving at the SDram
- This has been done on the template project file you have. However because the clocks speeds on the DE2 are slow, we can get away with a simple inverter/NOT gate to delay the dram clock by half a clock relative to the sdram controller state machine (i.e. plenty of time **50Mhz = 20ns**, so half a clock delay is **10ns**, that is the sdram clock is delayed by 10ns relative to Sdram controller clock). We'll look at the timing later.

# Issuing a CPU Address to the SDram Chip

- **Reading** and **Writing** to the SDram chip both require the presentation of a CPU address to the chip.
- The SDram latches the address internally as separate bank, row and column addresses upon application of an **Activate**, **Read** or **Write** command.
- We thus have to split the address lines coming from the **68k** into groups of **Bank**, **Row** and **Column** address bits. So how many address lines do we allocate to the Bank, Row and Column addresses?

We look at the data sheet for the Dram chip we are using, i.e. the Zentel chip on the DE2 and study it's architecture.

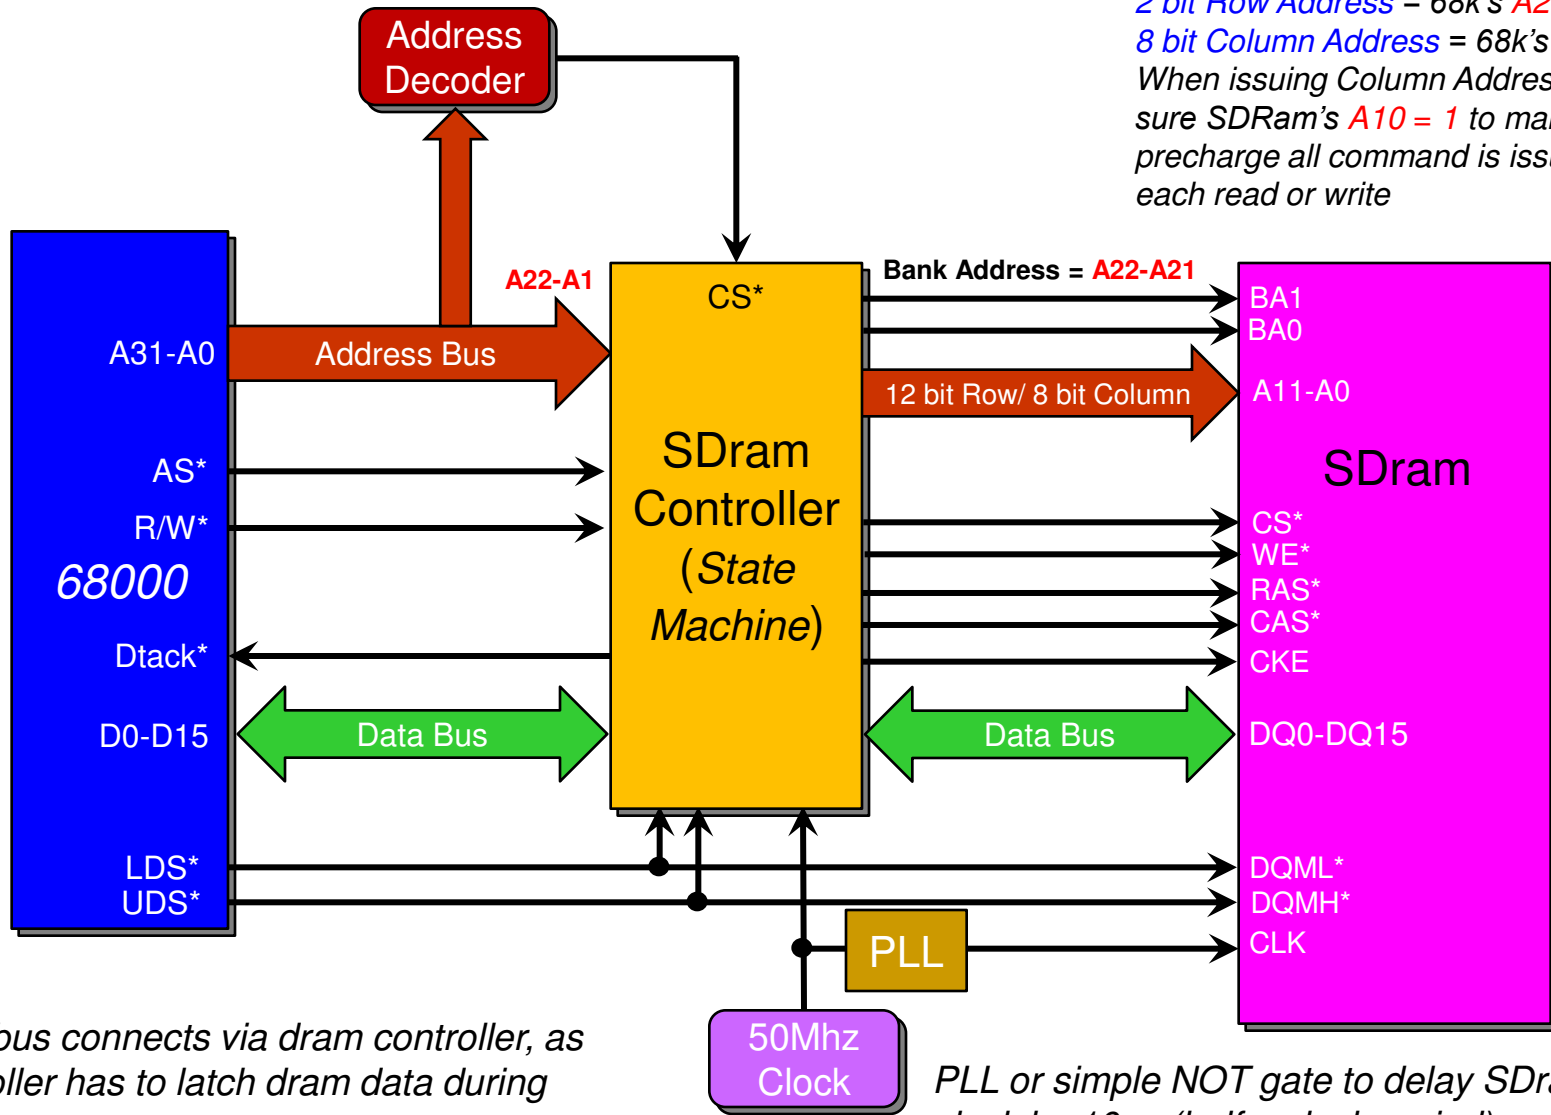
- The Zentel Chip is an **4M Word** chip organised internally as **4 banks** of **1M x 16** bits.
- This means that it requires a **22** bit address (**2** bits for the **Bank** and **20** bits to specify a unique **Row/Column** location within the **1M** space of the chosen bank).

# Issuing a CPU Address to the SDRam Chip

- Although it's not important in this application (*because we do not use burst mode for read or write operations*), it makes sense to allocate **A22-A21** of the 68k's address bus as the **Bank** address inputs to the Sdram, i.e. map **A22-A21** to **BA1-BA0**.
- We can then use **A20-A1** as the remaining address lines to present as **Row** and **Column** addresses (*but any mix of address lines should work*).
- Internally, each bank in the Zentel chip is organised into **4096** rows of **256** columns (see page 5), thus we need to break down **A20-A1** into a **12 bit Row address** and an **8 bit Column address** and present each to the SDRam via its **12** bit address bus.
- It's sensible then to present the 68k's address lines **A20-A9** as the **row** address (i.e. map them to **A11-A0** on the SDRam) and **A8-A1** as the **column** address (*i.e. map them to **A7-A0** on the SDRam*). This implies an address multiplexer of some kind inside the dram controller (*this is quite easy in VHDL*).
- **NOTE:** Early SDRams required that you issue a **pre-charge** command between a **read** or **write command** (*when for example switching to a new bank or a new row*), but more recent devices like the Zentel chip can perform an **auto pre-charge** after any read or write provided you set **A10** on the SDRam to logic 1 when presenting the **Column address** for a **read** or **write** operation.

# Interfacing the Soft Core 68k to 16 bit Wide SDRAM

2 Bit Bank Address = 68k's **A22-A21**  
 2 bit Row Address = 68k's **A20-A9**  
 8 bit Column Address = 68k's **A8-A1**  
 When issuing Column Address make  
 sure SDRAM's **A10 = 1** to make sure a  
 precharge all command is issued after  
 each read or write



Data bus connects via dram controller, as  
 controller has to latch dram data during  
 read

PLL or simple NOT gate to delay SDRAM  
 clock by 10ns (half a clock period)

# Writing to the SDram

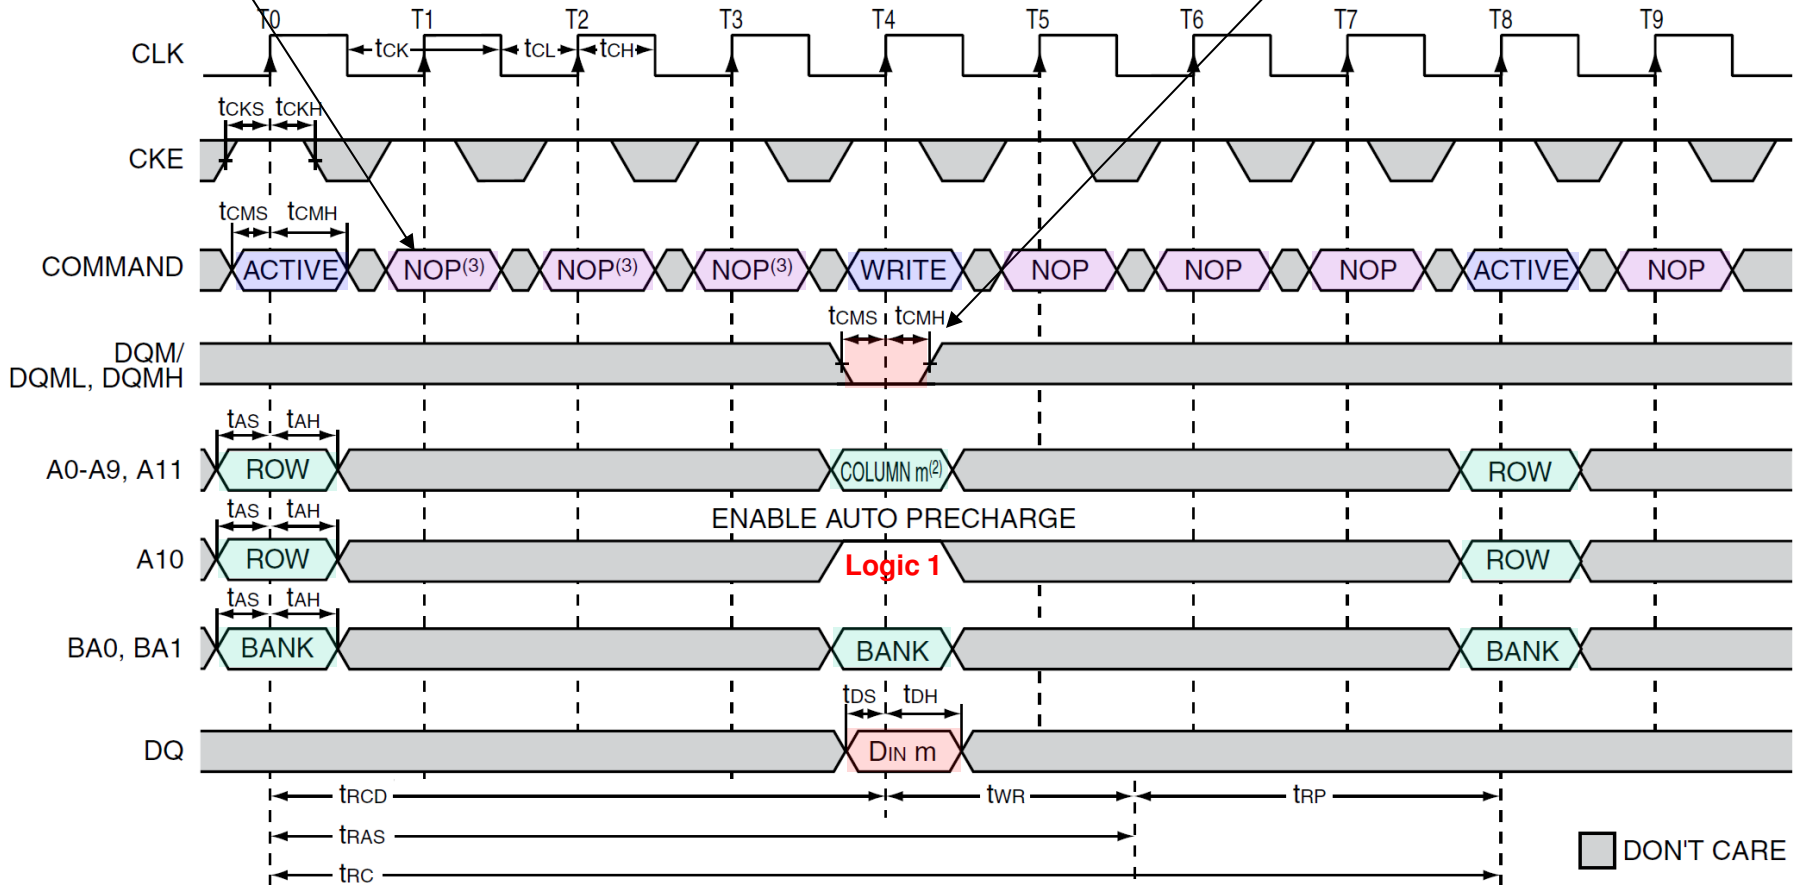
- In descriptive terms, the protocol for **writing** to a location in the SDram goes like this.
- 68k begins a **write** when **R/W** and **AS\* goes to 0**, which can trigger SDram controller to present a **Bank/Row** address to the SDram and issue an **Activate command** (**CS=0, WE=1, RAS=0, CAS=1**) to activate a **Bank/Row** within the chip (*Page 7*).
- In the case of the 68k, we cannot give the SDram the **write** command immediately after an **Activate** command because the **UDS\*/LDS\*** signals from the **68k** will not be valid at that point (*even though R/W indicates that a write is in progress*) so we have to insert SDram **NOP** commands until either **UDS\*** or **LDS\*** (*or both*) go low. This will ensure that the width of the 68k's data transfer is known to the SDram via the **DQMH\*/DQML\*** lines when a **write** command is actually issued.
- Once **UDS\*/LDS\*** become valid, controller can present a **Bank/8 bit Column** address (*remembering to set A10 = 1 at this point for an auto pre-charge after the write*) and issue a **Write** command (*Page 7*) i.e. **CS=0, WE=0, RAS=1, CAS=0**
- The data we present on the SDram during the **write** comes from the SDram controller – which in turn is taken from the 68k data bus (*see previous slide*) and will then be stored into the bank/row/column address specified. The timing is shown overleaf. This means we **must** turn **on** (i.e. **drive** – not **tri-state**) the SDram controller **data out** signals during a write to drive data **into** the SDram data lines.

# Writing to the SDram

NOP commands only necessary to meet min delay between ACTIVATE and WRITE command. Can be 0 for a slow clock or more than that for a fast clock

Setup and hold times of all signals relative to clock = **1.5ns**

Byte select signals **DQMH/DQML** must be valid at time of write



Min row activate to read or write command

$t_{RCD} \geq 18\text{ns}$  or 1 clk @ 50Mhz

Min time between reads/writes

$t_{RC} \geq 58\text{ns}$  or 3 clk @ 50Mhz

Time from end of write to activating a new row

$t_{RP} + t_{WR} = 3\text{ clks}$  @ 50Mhz


Min row active Time

$t_{RAS} \geq 40\text{ns}$  or 2 clks @ 50Mhz

**SINGLE WRITE - WITH AUTO PRECHARGE**

# Writing to the SDRAM

- You can see the simulated timing here where the 68k writes \$AA553322 to location 00860000/1/2/3 (in this system it's dram not static ram). **AS\_L** going low means we can issue **Activate** command. The address maps to a **bank address** of 0 a **row address** of hex 300 and a **column address** of 400 (it's a column address of 0, but **A10 = 1** (for the auto pre-charge command—see page 19) makes it 400).
- Your dram controller can generate **DramDtack\_L** pretty much as soon as it has presented the **Activate** command and must keep generating it until the 68k finishes its bus cycle (**AS\_L** goes high).



Notice **NOP** command(s) inserted between **ACTIVATE** and **WRITE** command. We issue these while waiting for the 68k's **UDS\***/**LDS\*** signals (and hence **DQMH\***/**DQML\*** at the SDRAM) to become valid to indicate data width **before** issuing the **WRITE** command.



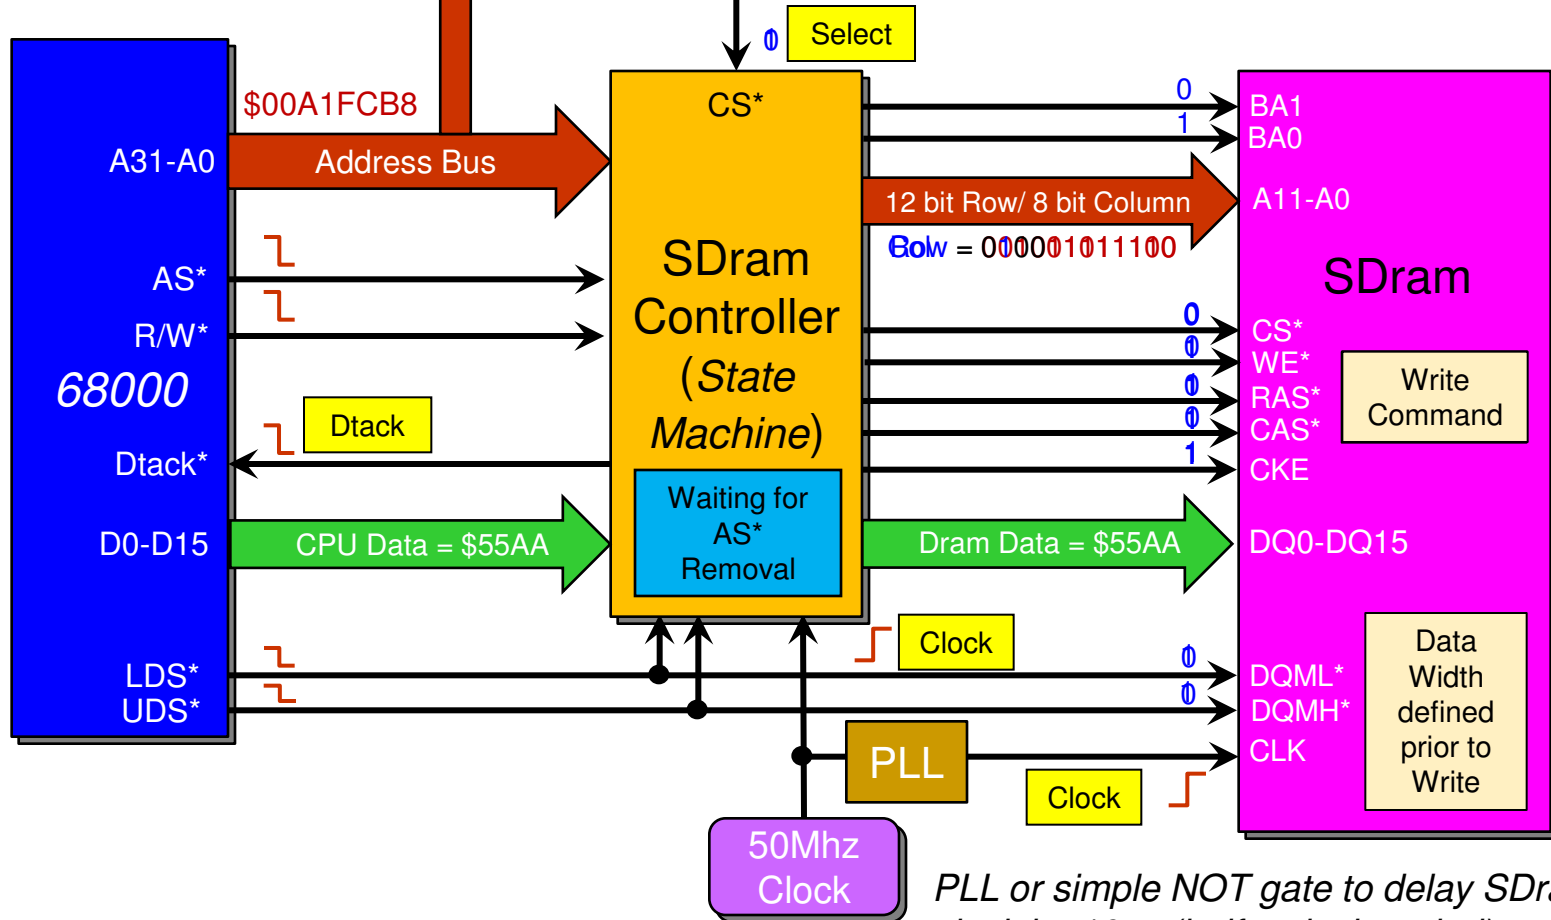
# Animation of Write Operation

Example Instruction

`move.w #$55AA,$A1FCB8`

Address : 1010 0001 1111 1100 1011 1000

2 Bit Bank Address = 68k's A22-A21  
 2 bit Row Address = 68k's A20-A9  
 8 bit Column Address = 68k's A8-A1  
 When issuing Column Address make sure SDRam's A10 = 1 to make sure a **pre-charge all command** is issued after each **write**

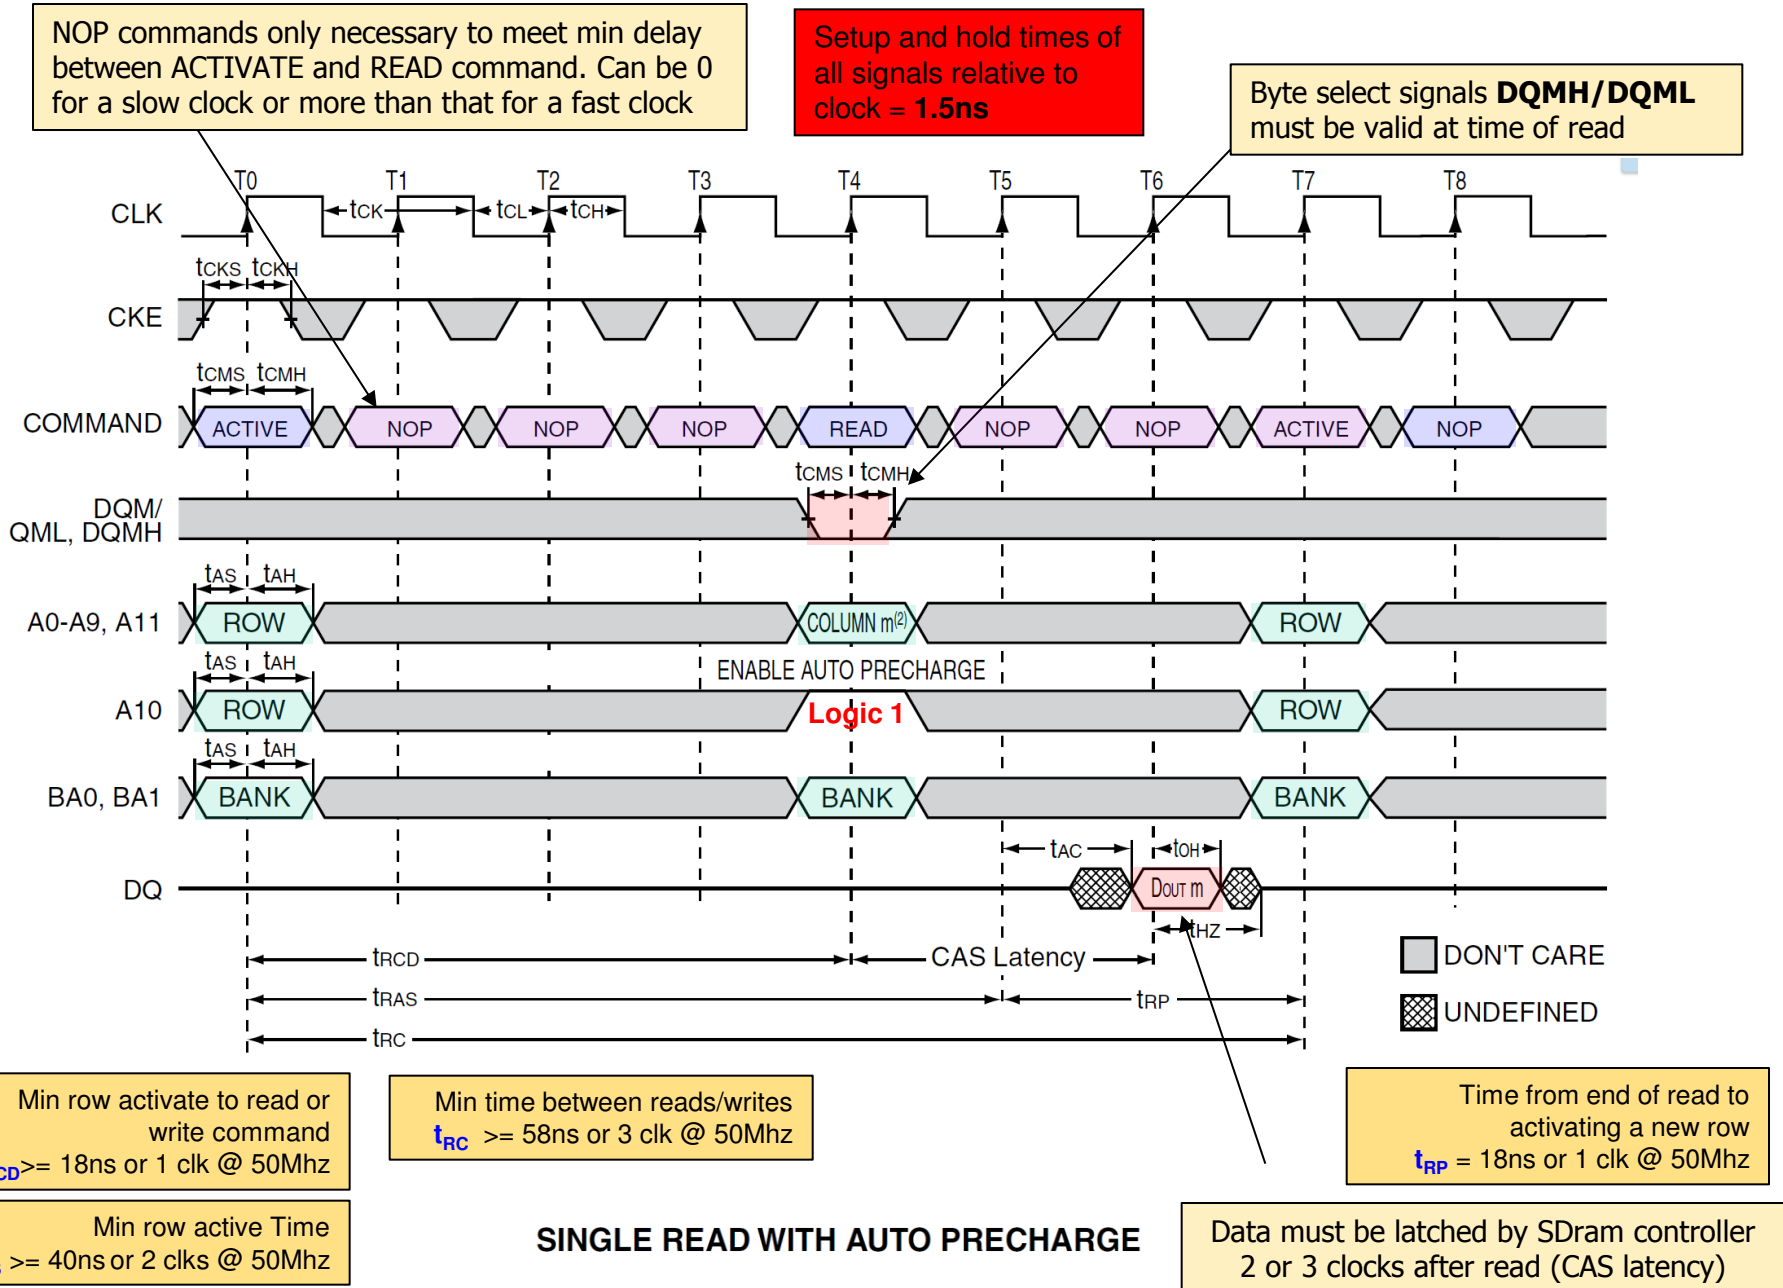


PLL or simple NOT gate to delay SDRAM clock by 10ns (half a clock period)

# Reading from the SDram

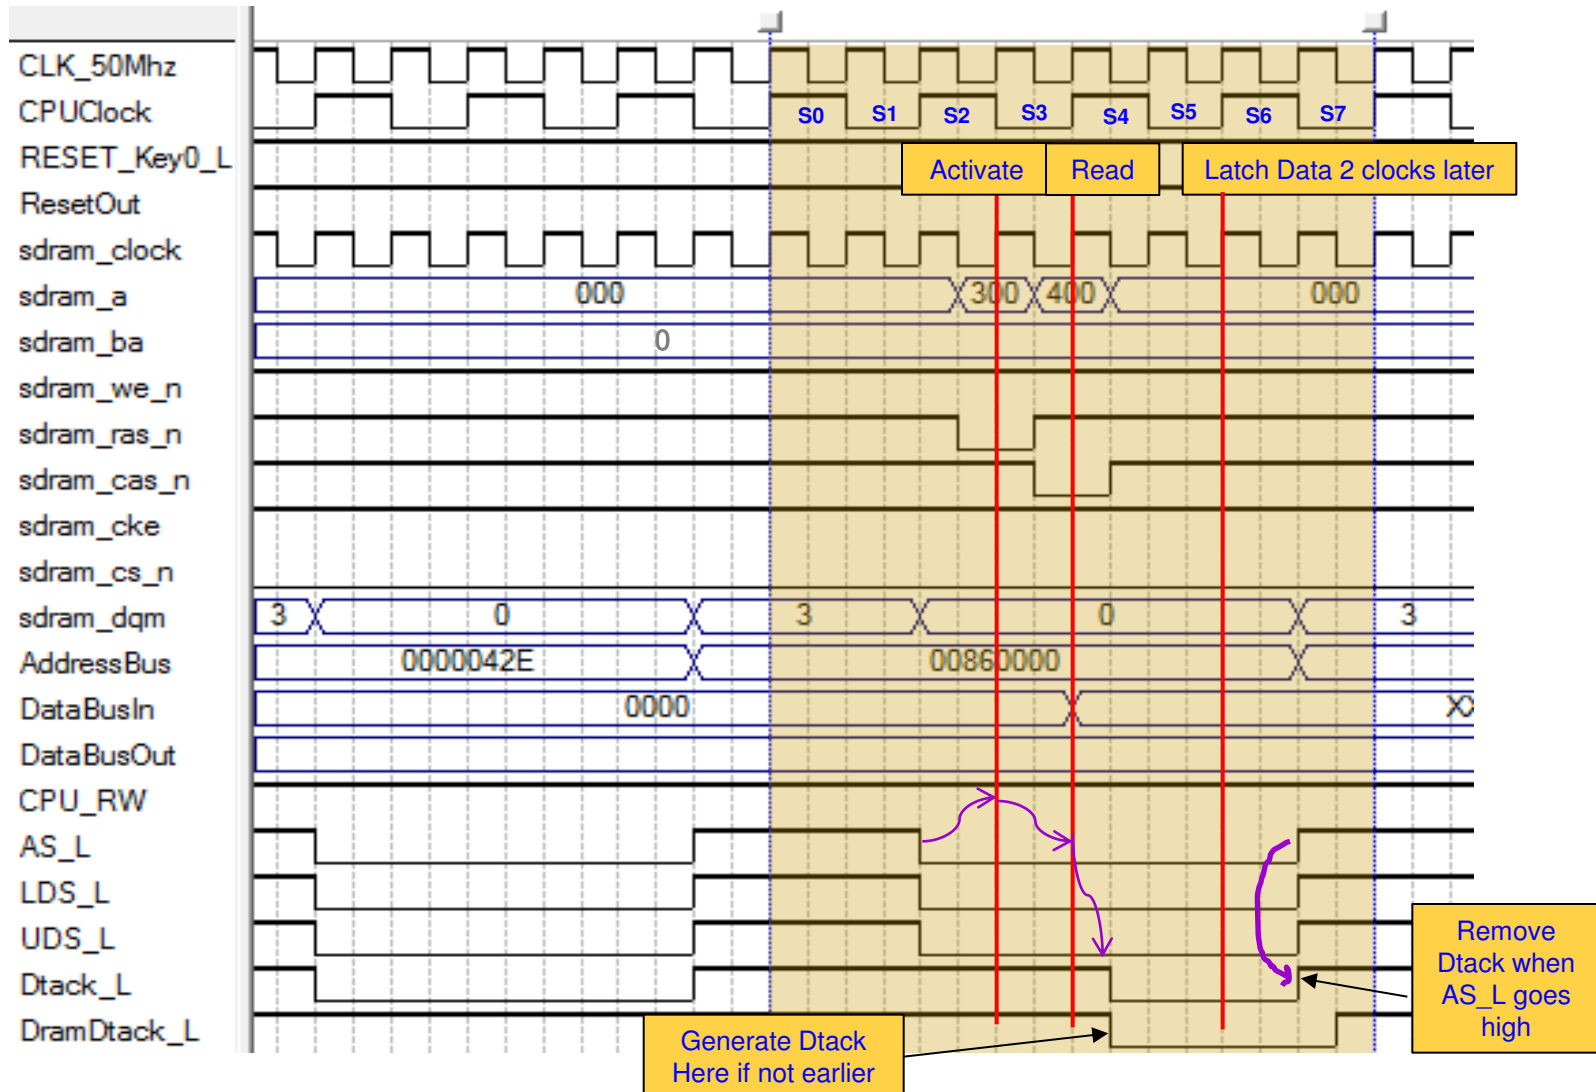
- Reading from the SDram involves presenting an address to the chip and issuing an **Activate** command, followed by a **READ** command. In descriptive terms, the operation to read goes like this.
- When CPU issues Address Strobe (**AS\* = 0**, **R/W = 1**), present a **Row** and **Bank** Address to the Sdram and issue the **Activate** command to activate a bank/row within the chip (Page 7) i.e. **CS=0**, **WE=1**, **RAS=0**, **CAS=1**. You can also generate **Dtack\*** here if your SDram design is fast enough to not require wait states on the 68000.
- Because **UDS\*/LDS\*** is driven **early** in the bus cycle for 68k **read** operation, we can present a **Column** and **Bank** address (*remembering to set **A10 = 1** at this point for an **auto pre-charge***) and issue a **READ** command immediately after the **Activate** command (Page 7) which means setting **CS=0**, **WE=1**, **RAS=1**, **CAS=0**.
- **IMPORTANT**, your SDram controller **must** tri-state it's data out lines to the SDram during a **read** operation so as not to create conflict at the bi-directional data lines of the SDram chip which will be driving data out at this time.
- The SDram will output it's data after the programmed **CAS latency** period of **2 clock cycles** has expired. Your SDram controller **MUST** latch this data internally and keep presenting it to the 68k, since the SDram chip will **remove** the data on the next clock. This is IMPORTANT (see next timing diagram) !!

# Reading from the SDRAM




# Reading from the SDRAM

- You can see the simulated **READ** timing here where the 68k reads a word at location **00860000**. This maps to a **bank** address of **0** and a **row** address of hex **300** and a **column** address of **400** (*actually it's a column address of 0, but  $A10 = 1$  makes it 400*).



# Reading from the SDRAM

- You can see the simulated **READ** timing here where the 68k reads a word at location **00860000**. This maps to a **bank** address of **0** and a **row** address of hex **300** and a **column** address of **400** (*actually it's a column address of 0, but  $A10 = 1$  makes it 400*).



## Reading from the SDram

- **Dtack\*** should only be generated when you are sure your data will be ready to give back to the 68k at the start of **S7**. Too early and the 68k may finish the access before the dram controller has latched the data from the dram.
- After experimentation and debugging, it was found to work if you issue Dtack\* in the state immediately after issuing the **Read** command (see *previous illustration*)

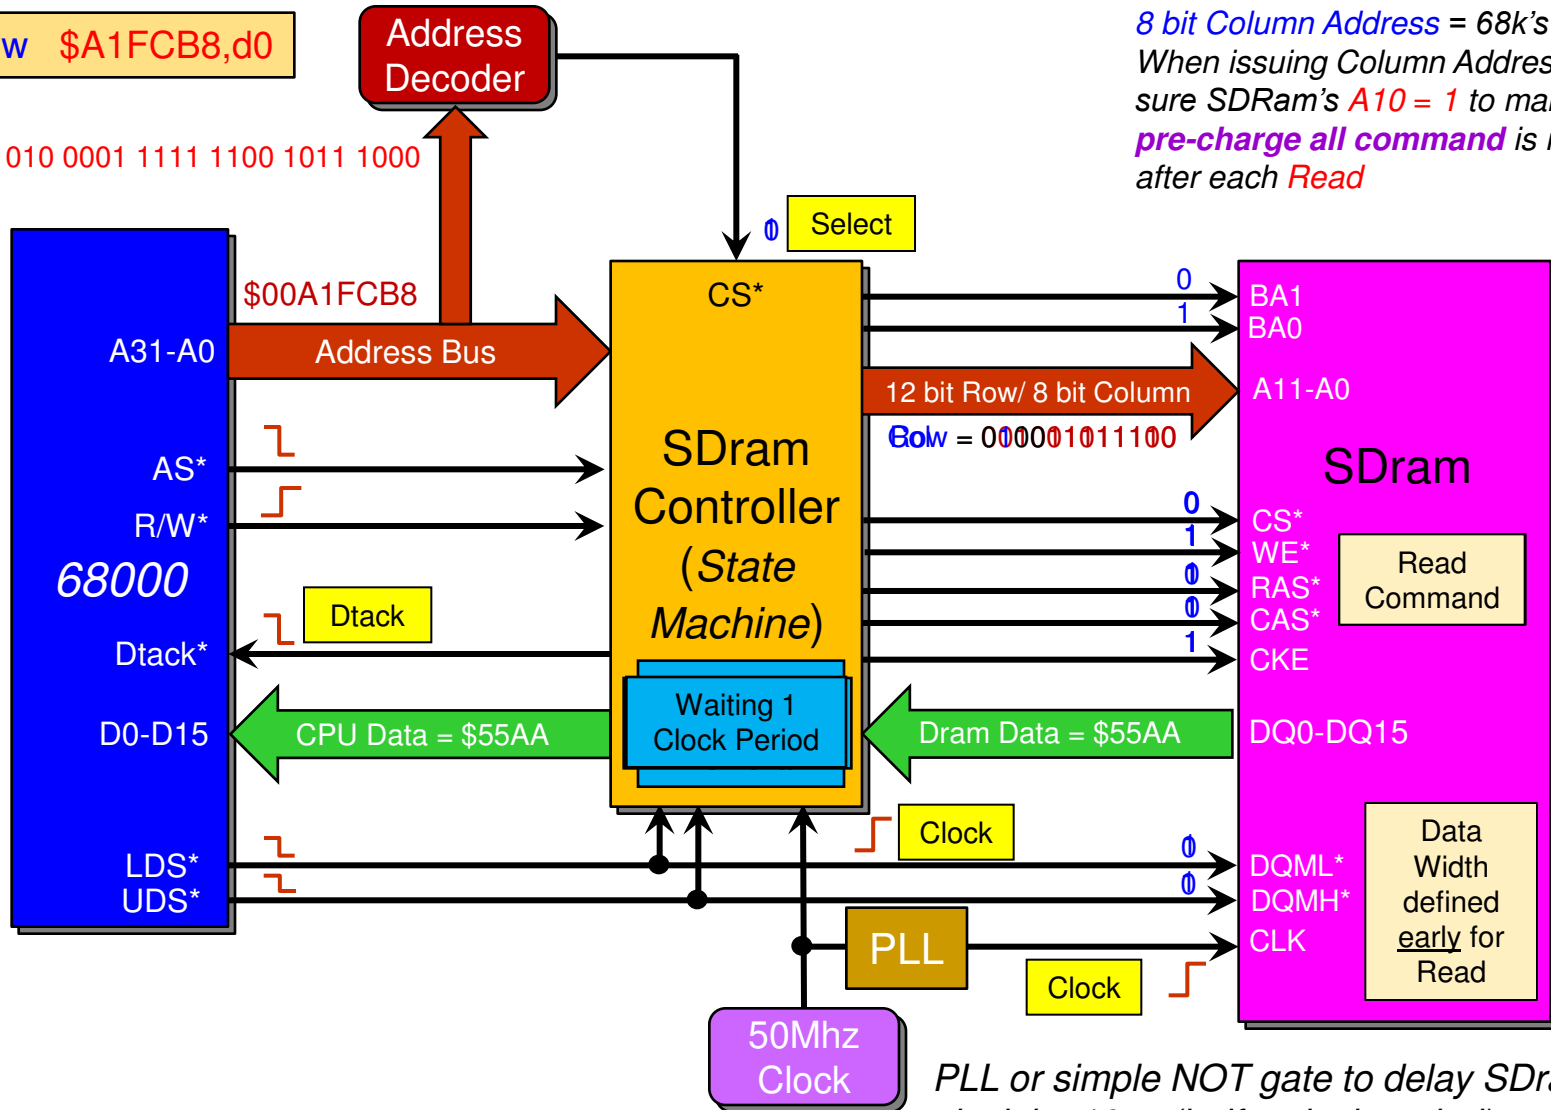
# Animation of Read Operation

Example Instruction

`move.w $A1FCB8,d0`

Address : 1010 0001 1111 1100 1011 1000

2 Bit Bank Address = 68k's A22-A21  
 2 bit Row Address = 68k's A20-A9  
 8 bit Column Address = 68k's A8-A1  
 When issuing Column Address make sure SDRam's A10 = 1 to make sure a **pre-charge all command** is issued after each **Read**



PLL or simple NOT gate to delay SDRAM clock by 10ns (half a clock period)

# SDram Refreshing

## What, Why and How

- Refreshing arises because the **charge** on the **capacitor** of the dram cell **leaks away**.
- Refreshing '**tops-up**' the charge before the data is lost due to leakage.
- Few memory chips refresh completely by themselves. External hardware is usually required if only to tell the dram when to perform it.
- Refreshing involves **reading every cell** in the chip (*which implicitly means that the charge on the cell gets topped up*) within a period specified by the manufacturer, (*typically 64-128mS*)



# Asynchronous Dram Refreshing

- A single refresh operation will refresh all cells sharing the same row address, thus we do not have to refresh all cells individually.
- For a 4 Meg location device such as the Zentel chip on the DE2 this means that every row has to be refreshed at least once during the specified 64mSec period and a row address is a 12 bit value, i.e. there are 4096 rows in that device.
- Most modern Dram chips simplify the task of refreshing by having a refresh counter fabricated onto the chip (see page 5) to keep track of which Row is to be refreshed next.
- **Important:** during a refresh, the CPU must not be permitted to access the memory chip. This can be achieved by delaying D<sub>tack</sub> back to the 68k until the chip has refreshed and then perform the delayed 68k's read or write request (*driven off AS\**).

# Asynchronous Dram Refreshing

## Hardware Refresh Modes

- **Burst Mode:**

Here all **Rows** in the chip are refreshed in one **continuous burst** of refreshing, one row after the other as fast as possible.

The downside is that the CPU will not be able to access the memory until **all rows** have been refreshed, leading to less predictable execution timing from a program, but most high performance CPUs have a **cache** these days, which means the CPU can probably carry on executing during a refresh.

- **Interleaved Mode:**

This is the most common arrangement, here the refreshing is interleaved between normal CPU accesses to the memory as it leads to more predictable program execution. For example, a refresh request for one row in the dram can be initiated every **64/4096 ms**, or every **15.6us** for a device like the **Zentel chip**.

- You will need a refresh timer in your dram controller to count clock cycles at **50Mhz** and issue a refresh request to your Dram controller state machine no later than every **15.6us**.

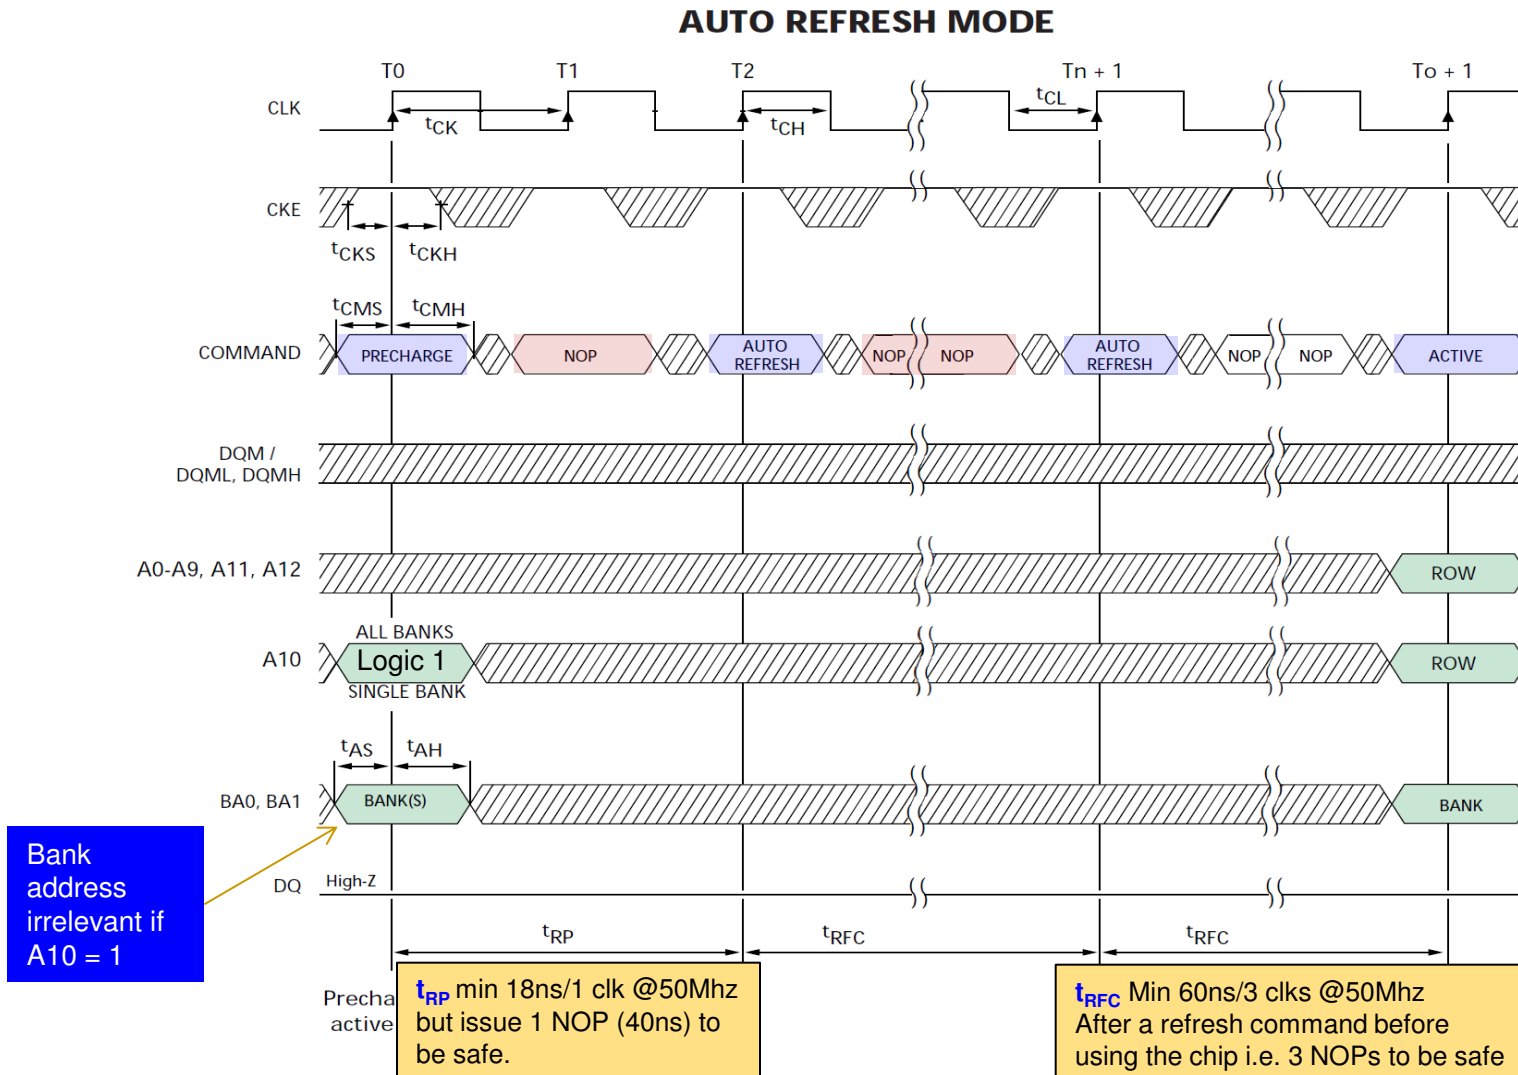
# Synchronous Dram Refreshing

- The Sdram chip has a command to tell it to perform an internal refresh of the next row (*it keeps track of which row is next using an internal counter*).  
Page 7, Refresh command means **CS=0, WE=1, RAS=0, CAS=0**
- Refreshing works like this

1. Issue a **pre-charge all banks command** (**A10** must = **1**) which means that the bank address is *irrelevant*.
2. Issue **1 NOP** command (i.e. at least 18ns – see next page)
3. Issue an **Auto-Refresh** command
4. Issue **3 NOP** commands (i.e. at least 60 ns or 3 clocks after last command – see next page)
5. Dram is now ready for 68k to access.

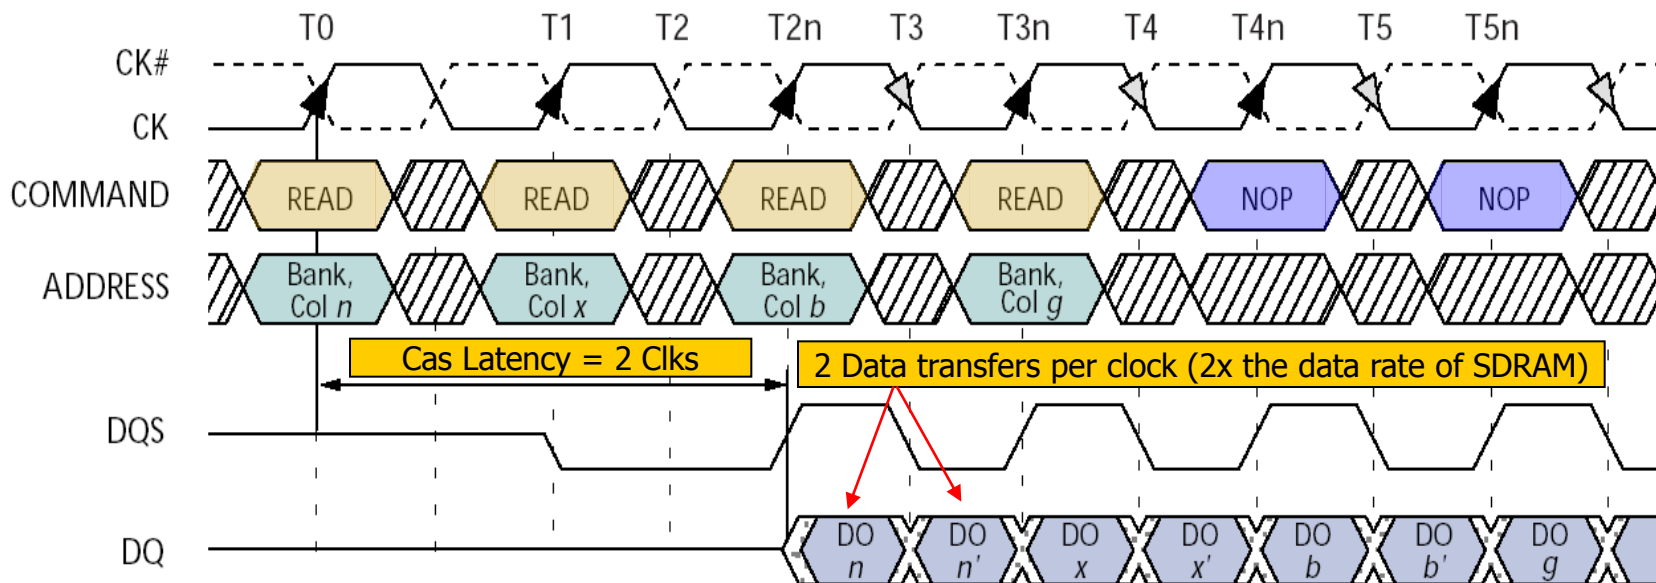
# Synchronous Dram Refreshing

- The timing diagram is shown here



## Double Data Rate Synchronous Dram (DDR SDRAM)

- This technique improves upon ordinary SDRAM by transferring data on both the *rising* and *falling* edges of the clock signal. **DQS** acts as a data strobe to help latch the output data during read operations and input data during writes.
- For a **133Mhz** Clock and a byte wide chip, **266Mbytes** of data per second can be transferred hence the name **DDR-266** Ram.
- This makes it twice as fast as SDRAM after the initial **CAS\*** latency, although in practice (*given the presence of good CPU **caches***) the speed improvement as observed by the user over SDRAM is not large (maybe **5%**).
- Today **DDR3-33 Ram** (**167 MHz** Clock) and **DDR-400** (**200Mhz**) chips are available.
- The timing diagram below shows random column address access to the chip with data for column **n** and subsequent data **n'** available on each edge of the clock.



# DDR SDRam DIMM Modules

- DDR SDRAM chips are most commonly used inside the PC where they are often packaged onto a small **daughterboard** module where typically **8** devices operate in **parallel** giving **64 bit wide data paths** and helping to achieve high bandwidth.
- These modules are advertised with names such as **PC2100**, **2700**, **3200** etc. ram which derive their name from the **bandwidth** of the module, e.g. **PC2100 = 2100 MB/sec**.
- This is related to clock speed and the width of data. For example, using **133MHz DDRram** with 64 bit wide data transfers on both rising and falling **edge** of the clock gives us  **$133 * 2 * 8 = 2100$  MBytes per second**.



# DDR2 and Beyond

- DDR2 keeps the same relatively slow (*and cheap*) 'core' memory speed used in 133Mhz SDRAM and DDR 266 ram, with roughly a **15 ns latency or access time** i.e. **2 clock** cycles at **133 MHz** to get at the **first** item of data in exactly the same time, but by doubling the **clock rate** of DDR ram, DDR2 ram could get at *subsequent* data in the same row even faster.
- However because of the slow core memory speed, initial CAS latency is now **4** clock cycles to get at the first item of data so they are labelled **4-1-1-1** etc.
- So **DDR2** ram clocked at **533Mhz** transferring data on both edges giving us **DDR2-1066 memory chips and PC2-8500** Memory modules.


Standard name	I/O Bus clock	Data transfers per second	Module name	Peak transfer rate
DDR2-400	200 MHz	400 Million	PC2-3200	3200 MB/s
DDR2-533	266 MHz	533 Million	PC2-4200 PC2-4300 <sup>1</sup>	4266 MB/s
DDR2-667	333 MHz	667 Million	PC2-5300 PC2-5400 <sup>1</sup>	5333 MB/s
DDR2-800	400 MHz	800 Million	PC2-6400	6400 MB/s
DDR2-1066	533 MHz	1066 Million	PC2-8500 PC2-8600 <sup>1</sup>	8533 MB/s

- **DDR3** uses lower voltages and even higher clock frequencies (**1Ghz**)
- **DDR4** is expected to use **1.2v** and clock frequencies **>2Ghz**



# Intel and Micron's Hybrid Memory Cube Consortium

- Stack Dram chips vertically with connections made through TSV (*through silicon vias*). Needs very precise alignment during manufacture  $< 3\mu\text{M}$ .
- 160GB/sec bandwidth – 1 possible replacement is for DIMM memory modules.
- Vastly reduced power.






# Intel and Micron's Hybrid Memory Cube Consortium

- Intel's Xeon Phi "*Knights Landing*" (3-TFLOPS) processor developed for massively parallel processing, is rumoured to feature up to 72 cores with connections ready to accept 8 hybrid memory cubes yielding 16-Gbytes of on-chip memory.
- It has enough in-package memory and off-chip DDR (max 384 Gbytes) to run massively parallel applications without adding fancy swapping and caching techniques.



Micron 2GB memory cube



Intel's Xeon Phi "*Knights Landing*"  
CPU

# Improving Memory System Performance

- Essentially there are four ways to improve the performance of a memory system
  1. The most obvious is to raise the IO **clock speed** of the device so that data can be accessed more quickly (as with DDR3 Ram which run at up to 1GHz).
  2. Increase data width e.g. **64bit**, **128bit**, **256bit**, **512bit** wide memory etc. to get more data per clock.
  3. Have multiple banks of memory e.g. multiple hybrid memory cubes.
  4. In areas such as graphics cards, where massive performance is required, these techniques are combined. The latest **AMD Radeon R9 290X** (dual GPU) utilises **2 x 512 bit wide** memory data paths (**128** bytes transferred per clock) with GDDR5 Ram (*essentially DDR3 optimised for graphics applications*) clocked at **1.25Ghz** to achieve transfer rates in excess of **320 GBytes/sec** (i.e.  $1.25 * 2 * 128$ ). Which is enough for 600 frames per second at 1080p resolution