


Error Detection and Correction for High Integrity Memory Systems

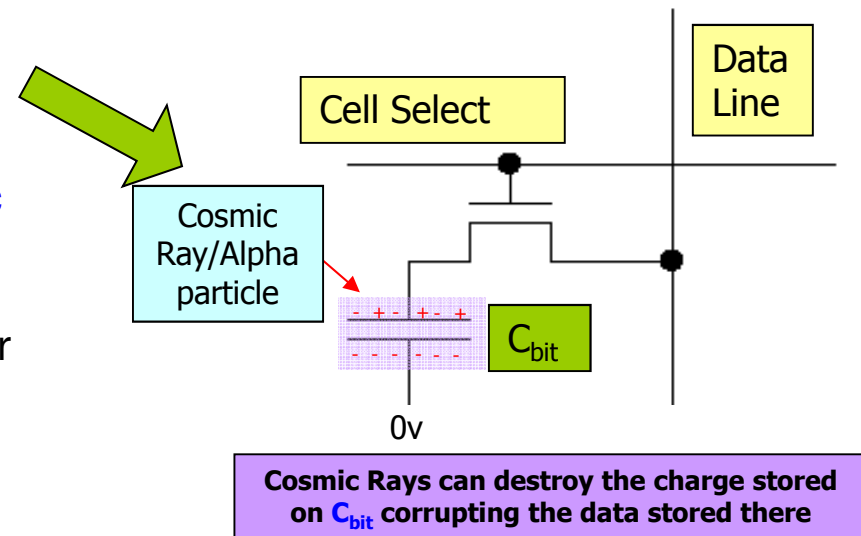
- Limitations to Dram Reliability.
 - Causes of Dram Failure.
- Simple Error **Detection** Schemes – **Parity**.
- Simple Error **Detection** and **Correction** Schemes - **Hamming Codes**.
 - Single Bit Error Detect, Single Bit Error Correct Codes (**SED/SEC**).
 - Double Bit Error Detect, Single Bit Error Correct Codes (**DED/SEC**).
 - Implementation in Memory Systems.



Error Detection and Correction in Memory Systems

Transient or Soft Errors in Dram Cells

- We know that Dram is the main type of memory used in **large, high availability, high reliability** computer system (*e.g. servers*).
- However, a significant problem with **DRAM** Cells is that they are **not** 100% reliable.
- DRAM data is **delicate**, because it is maintained as a **minute charge** stored on a capacitor. Consequently it doesn't take much to upset that delicate balance between a **logic 1** and a **logic 0**.
- This makes Drams susceptible to '**transient**' or '**soft**' errors where the data can become corrupted by unforeseen external influences (*nothing to do with failing to refresh them or any bad design for example*).
- Such errors can occur as the result of external radiation – Cosmic Rays etc. These are called **soft errors** as they do not permanently damage the chip, writing to it "**fixes**" the data.



Error Detection and Correction in Memory Systems

- **Soft errors** can be a real problem in **high availability, high integrity** memory systems, not just from an operational point of view (*i.e. the memory fails and the system may need rebooting*) but because they occur at **unpredictable times** and the failure is hard to **replicate** and **pinpoint**.

The hardware and software engineers often disappear on a “**wild goose chase**” trying to track down a bug in the software or the Dram controller which often isn't the cause of the problem.

(How many times do we blame Microsoft when our PC falls over?)

- **Swapping** dram chips doesn't improve matters as the faults are **not** down to a **physical failure** of the chip, i.e. it not due to **wear and tear** or **manufacturing defects**.
- Rather the failure is attributable to the **physics** of the capacitors used in the devices and is thus an **inherent problem** that users of Dram have to live with.
- Dram manufacturers have taken steps to “harden” their Drams to **shield** them better, but the move to higher density and lower voltages has seen a rise in soft error rates for **all** memory technologies: **Sram** and **Dram** (as well as multi-level cells in **Flash** memory)

Error Detection and Correction in Memory Systems

- In your average **Home PC**, soft errors in Dram memory systems are not a *significant* problem, because data/programs etc. tend **not** to be left lying around for long enough for the errors to occur or **accumulate** to the point where they become a problem.

This is because the memory in your home PC is forever being loaded with new data/programs, so soft errors, if they do arise, tend to be **overwritten** (*and hence eliminated*). Likewise we often shut down/reboot our PC when it is inactive.

Note: those who put their PC into **standby mode** when it's not in use are likely to notice the problem most, followed by those who **hibernate** their PC (*which saves and reloads any corrupted data to disk*).

- In **high availability systems**, where servers remain running 24/7 and require **99.9999%** availability, it can be a significant problem.
- A system with 1 GByte of RAM can expect an error every two weeks; a hypothetical 1TByte system would experience a soft error every few minutes. (http://www.tezzaron.com/about/papers/soft_errors_1_1_secure.pdf)

Error Detection and Correction in Memory Systems

What are the Causes of Dram Soft Errors?

- Mostly these are caused by **two** things:-
 - A cell being struck by a stray **Cosmic Ray** from outer space ionizing and corrupting the delicate charge stored there. **Alpha particles** emerging from the material packaging of the chip.
 - A cell being **sensitive to patterns** of charge stored on **neighbouring** cells, i.e. the charges on neighbouring cells can **all interact** (or **leak** into each other) and cause a bit to **'flip'** unintentionally.

(these were things that the first generation of Drams were very susceptible too until manufacturers learned to improve fabrication techniques and make them more resilient).

Note: Read/write errors are also becoming an issue in **Flash** memory cells, particularly where we attempt to store more than **1 bit per cell**, for the same reasons above, Error Detection and Correction may be needed here also.

Error Detection and Correction in Memory Systems

How do we deal with these problems?

1. A Memory Controller *could* Detect faulty data read from memory and inform the CPU via some kind of hardware 'exception'.

Verdict : Simple and cheap.

i.e. the dram controller does not try to correct the error, but simply informs the Operating System (via a CPU exception) that the access is unsafe. The application making the access will most likely be shut down or the system rebooted.

2. Detect the error and attempt to correct it "on-the-fly" in hardware so that the system can continue operation.

Verdict: More complex and expensive.

- The choice boils down to cost and how important the integrity and availability of your system is.

Error Detection and Correction in Memory Systems

Error Detection Schemes (*no correction*) – Single Bit Parity Checking

- **Single Bit Parity** schemes have been around for years and are able to detect the occurrence of **single bit errors** (and in fact **all errors** where an **odd** number of bit errors have occurred) in each memory location, but such a simple scheme cannot *correct* any of the errors.
- Single Bit Parity schemes utilise an additional **single bit of storage** for each and every memory location in memory and is thus relatively cheap and easy to implement.
- In the case of a **byte wide memory** system, such a scheme would require **9 bits** of storage for each and every memory location, i.e. **8 bits of data + 1 bit of parity** per location.

Error Detection and Correction in Memory Systems

Simple (i.e. cheap) **Parity** Schemes for error detection (*not* correction)

- With a simple Parity scheme, each byte, word or whatever width of data the memory system is organised as, is “tagged” with an additional single **Parity** bit before it is written to memory.
- This **parity** bit could be designed to be *even* or *odd parity*.
- That is to say, that the **total number of logic 1's** in the stored word, **including the parity bit itself**, is either an **even** or an **odd number**.
- With an **Even** Parity scheme, the parity bit is set to a value that would result in the total number of '1's stored in the memory location being an **even number**.
- With an **Odd** Parity scheme, the parity bit is set to a value that would result in the total number of '1's stored in the memory location being an **odd number**.

Error Detection and Correction in Memory Systems

Parity Bit Generation

- To generate the **Parity** bit **P** for an **even** parity scheme, we simply **add** the data bits together **ignoring any carry**.
- This turns out to be the same as simply **Exclusive-OR'ing** the data bits $b_0 - b_7$.
- The resulting parity bit is stored alongside the data in that location.
- **Odd** parity is simply the inverse of **Even** parity i.e. the **EX-NOR** of bits $b_0 - b_7$.

EXOR Gate Truth Table	
A B	Parity Bit
0 0	0
0 1	1
1 0	1
1 1	0

Result is a 1 when number of 1's in A and B is ODD

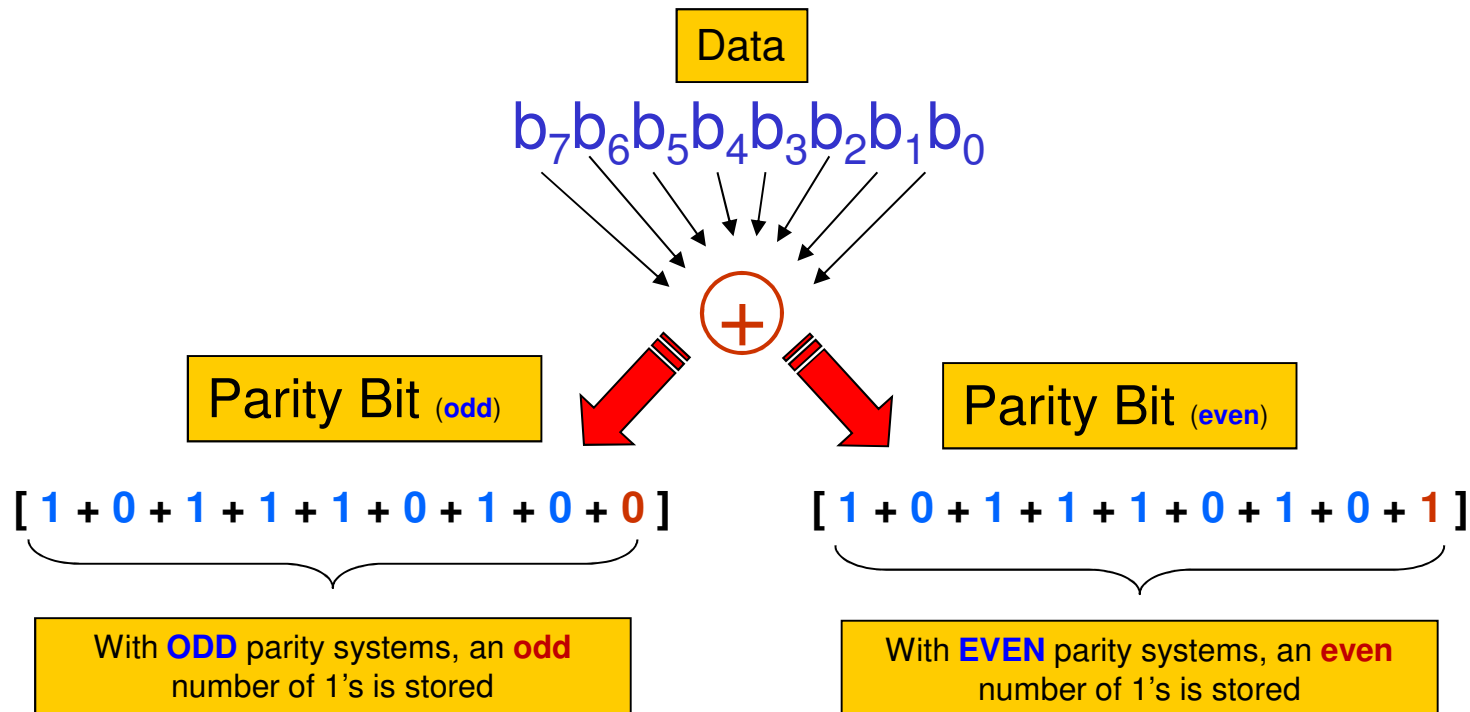
EX-NOR Truth Table	
A B	Parity Bit
0 0	1
0 1	0
1 0	0
1 1	1

Result is a 1 when number of 1's in A and B is EVEN

Error Detection and Correction in Memory Systems

Example 8 bit Data : [10111010] (Number of '1's in data = 5)

- For an **even** parity scheme the parity bit **P** will be set to **1**, resulting in the stored word (*data plus parity*) containing **6** logic '1's, i.e. an **even** number of logic 1's.
- For an **odd** parity scheme, **P** will be set to **0**, resulting in the stored word containing 5 logic '1's i.e. an **odd number** of logic '1's



Error Detection and Correction in Memory Systems

Problem

- Generate the value of the **Even** Parity bit for the following two items of **12** Bit Data:-

[101101001110]

[010110010101]

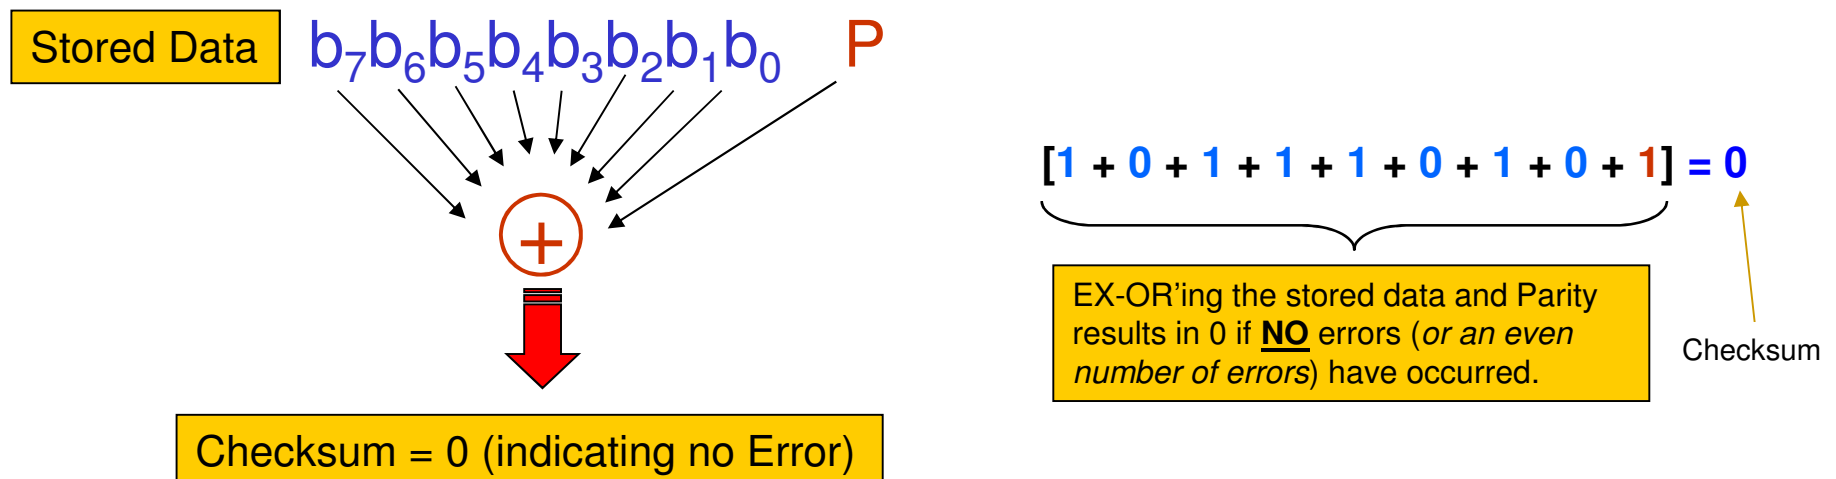
Two Solutions:

- **Add** up the data bits in binary (*ignoring the carry*) and the result is the value of the Parity Bit stored with the data.
- Alternatively, **count** the number of **1**'s in the data bits and assign a value to the Parity Bit such that the total number of **1**'s stored to an **even** number.

Error Detection and Correction in Memory Systems

Detecting an Error During a Read Operation

- In effect what we are doing with an EVEN parity scheme is introducing an **extra**, redundant bit of data such that when the whole word (**8 bit data plus 1 parity bit**) is **read back** and all the bits (**data plus parity**) are **Added** (i.e. **EX-OR'ed**) together, the **result** (called the checksum) should always be 0.
- If any **single** bit of data (or even just the parity bit itself) has been **corrupted**, the calculated **Checksum** will be 1 (i.e. the checksum bit will be **flipped**) and it is this that allows us to **detect** an error during a read operation. In fact any number of **odd** bit errors will result in a **Checksum** = 1, but an even number of bit errors (2,4,6 etc.) will **not be detected** at all since it causes a **double flipping** of the Checksum bit which cancels each other out. This is a limitation of the simple parity scheme.



Error Detection and Correction in Memory Systems

Example of Error Detection with Parity

- Suppose the following 9 bit data word i.e. 8 bit data + Even Parity Bit were read from memory. [10111010 | 1].
- EX-OR'ing these 9 bits together gives a checksum = logic 0, meaning (for an even parity scheme) that the data was read correctly (or possibly that an even number of errors had occurred).
- Suppose instead that we read back the following 9 bit data word [10111011 | 1] where the data bit b_0 has been corrupted. The result of EX-OR'ing (adding with no carry) these 9 bits together to produce the checksum, would result in a logic 1, i.e. an error.
- Suppose instead that we read back the following 9 bit data word [10111001 | 1] where the data bits b_0 and b_1 have been corrupted. The result of EX-OR'ing these 9 bits together would be a logic 0, i.e. no error for an even parity scheme.
- Note that even with single bit errors, this simple 1-bit parity scheme cannot identify the *location* of the erroneous bit and thus it cannot be used to as a means of correcting the faulty bit.

Error Detection and Correction in Memory Systems

Error Detection and Correction Schemes – Hamming Codes

- Depending upon the chosen level of sophistication, Hamming Codes will enable a memory controller to detect and correct single-bit errors within each memory location.

Such a scheme may well impose significant additional storage requirements on our system leading to increased cost and complexity.

For example, by storing an additional 4 bits of hamming code, alongside each and every byte of data, i.e. making a 12 bit wide memory system instead of 8 bit wide, we would be able to detect and correct all single bit errors arising within that location, including errors in the hamming code bits themselves.

The more sophisticated the hamming code scheme, the more erroneous bits can be detected and corrected but at a tradeoff of increased memory cost and a more sophisticated memory controller (*how much is increased reliability worth?*)

Error Detection and Correction in Memory Systems

Error **Correcting** Techniques: **Hamming Codes**

- Basic Error Correcting techniques in memory systems borrow concepts from **Digital Communications theory**.
- The idea is based on **Forward Error Codes** (F.E.C.s), where **multiple Parity bits** are encoded into a data stream and transmitted as a **Code Word**. That is, a code word consisting of [**Data Bits + Multiple Parity Bits**]
- F.E.C.'s are essentially multiple **parity bits**, **carefully chosen** so that they can **identify** the **location** of any **single bit error** (**data** or **parity**) in the resultant code word. Hardware *could* use these parity bits to **identify** and **correct** bit errors.
- The table opposite highlights the number of parity bits required for a specified data width to enable single bit error **detection** and single bit error **correction** (i.e. **SED/SEC**). We can use this idea in memory systems.
- As you can see this scheme becomes more cost effective with wider data widths.

<u>Data bits</u>	<u># of Parity Bits</u>
4	3
8	4
12	5
16	5
20	5
24	5
28	6
32	6
64	7
etc	

1 extra byte wide
memory chip required to
protect 8 bytes of Data.

Error Detection and Correction in Memory Systems

- In communications, where the data is transmitted **serially**, the extra **Parity Bits** were transmitted in **bit positions** corresponding to **powers of 2**, i.e. bit positions **1, 2, 4, 8, 16, 32** etc. depending on data. →

Code Word Bit Position

7	6	5	4	3	2	1
d_4	d_3	d_2	p_3	d_1	p_2	p_1

- The significance of this, as well as counting from position **1** (*not 0*) will become apparent shortly.
- Using an **even parity scheme**, as shown below, (where \oplus means **EX-OR** or “**add with no carry**”), we generate the values for the parity bits p_1 , p_2 and p_3 (assuming just for 4 bits of data) according to these equations.
- This will ensure that during a read, **EX-OR'ing** the **data** and the **parity bits** together in the **same way** will produce three **checksums all = 0**.

$$p_1 = d_1 \oplus d_2 \oplus d_4$$

i.e. during a read

$$p_2 = d_1 \oplus d_3 \oplus d_4$$

i.e. during a read

$$p_3 = d_2 \oplus d_3 \oplus d_4$$

i.e. during a read

$$p_1 \oplus d_1 \oplus d_2 \oplus d_4 = 0$$

$$p_2 \oplus d_1 \oplus d_3 \oplus d_4 = 0$$

$$p_3 \oplus d_2 \oplus d_3 \oplus d_4 = 0$$

Each Checksum Bit = 0

Error Detection and Correction in Memory Systems

7	6	5	4	3	2	1
d_4	d_3	d_2	p_3	d_1	p_2	p_1

The data in positions 3, 6, 7 when written in binary have a '1' in bit position 1

The data in positions 3, 5, 7 when written in binary have a '1' in bit position 0

How do we determine the Parity Bit Equations?

- Write out in binary the position numbers occupied by all data bits, i.e. data bits occupying positions 3, 5, 6 and 7, i.e.

d_1 = Position 3 = 011
 d_2 = Position 5 = 101
 d_3 = Position 6 = 110
 d_4 = Position 7 = 111

- Identify all those position numbers where there is a '1' in the least significant bit position (*right most column*). These indicate the positions of the data bits that have to be **EX-OR**'ed to produce parity bit p_1 .
- Likewise parity bit p_2 will be formed by 'EX-OR' ing the data bits occupying positions where there is a 1 in the next l.s.b. of their position number.

011 = 3	} p_1	d_1
101 = 5		d_2
111 = 7		d_4

011 = 3	} p_2	d_1
110 = 6		d_3
111 = 7		d_4




101 = 5	} p_3	d_2
110 = 6		d_3
111 = 7		d_4

The data in positions 5, 6, 7 when written in binary have a '1' in bit position 2

Error Detection and Correction in Memory Systems

7	6	5	4	3	2	1
d_4	d_3	d_2	p_3	d_1	p_2	p_1

- Parity bit p_1 is thus formed from the **exclusive OR** of **data** bits occupying positions **3, 5 & 7**, i.e. the EX-OR of data bits d_1 , d_2 and d_4 .
- Parity bit p_2 is formed from the **exclusive OR** of data bits occupying positions **3, 6 & 7**. i.e. data bits d_1 , d_3 and d_4
- Parity bit p_3 is formed from the **exclusive OR** of data bits occupying positions **5, 6 & 7**. i.e. data bits d_2 , d_3 and d_4
- When we read back the **data** and the previously stored **parity** bits and **EX-OR** them all together (*like we did for the simple parity scheme*) to create a **3 bit checksum**, the checksum will be **[000]** if there were no errors.



Error Detection and Correction in Memory Systems

Problem 1

- How many Parity Bits would be required to enable correction of a **single bit error** in a **6 bit data word**? What **positions** would the parity and data bits occupy in the code word?

Solution


- First construct the code word as follows placing parity bits at positions in the code word corresponding **to powers of two**, until you have placed all data bits.

Error Detection and Correction in Memory Systems

Problem 2

- Write down the **parity equations** for a **6 bit data word**.
- First write out in binary the position numbers occupied by the Data Bits in the code word, then look for **1's** in the columns.

3	0011	d1
5	0101	d2
6	0110	d3
7	0111	d4
9	1001	d5
10	1010	d6



P1 = EXOR of data bits occupying positions 3,5,7,9, i.e. data bits d1, d2, d4, d5
P2 = EXOR of data bits occupying positions 3,6,7,10 i.e. data bits d1, d3, d4, d6
P3 = EXOR of data bits occupying positions 5,6,7 i.e. data bits d2, d3, d4
P4 = EXOR of data bits occupying positions 9,10 i.e. data bits d5, d6

Error Detection and Correction in Memory Systems

Example Problem using 4 bit Data

- Take the 4 bit data word $d_4, d_3, d_2, d_1 = [1011]$. We calculate the parity bits according to the equations developed earlier.

$$\begin{aligned} p_1 &= d_1 \oplus d_2 \oplus d_4 \\ p_2 &= d_1 \oplus d_3 \oplus d_4 \\ p_3 &= d_2 \oplus d_3 \oplus d_4 \end{aligned}$$

- We generate each parity bit thus

$$\begin{aligned} p_1 &= 1 \oplus 1 \oplus 1 & \text{i.e. } p_1 &= 1 \\ p_2 &= 1 \oplus 0 \oplus 1 & \text{i.e. } p_2 &= 0 \\ p_3 &= 1 \oplus 0 \oplus 1 & \text{i.e. } p_3 &= 0 \end{aligned}$$

- These 3 parity bits are stored such that the resultant code word (*data plus parity*) would be

$$d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [1 \ 0 \ 1 \ 0 \ 1 \ 0 \ 1]$$

Error Detection and Correction in Memory Systems

How do we Check for an Error during Reading?

- We calculate a set of check bits (*the checksum*) using the **original equations** that were used to create the parity bits in the first instance. For example, assuming the correctly generated code word given earlier (shown below).

$$d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [1 \ 0 \ 1 \ 0 \ 1 \ 0 \ 1]$$

- We calculate check bit C_1 as the EX-OR of p_1 plus the **data bits** that were originally used to generate p_1 in the **first place**, ditto for check bits C_2 , C_3 etc. Thus

$$C_1 = p_1 \oplus d_1 \oplus d_2 \oplus d_4$$

$$C_2 = p_2 \oplus d_1 \oplus d_3 \oplus d_4$$

$$C_3 = p_3 \oplus d_2 \oplus d_3 \oplus d_4$$

$$C_1 = 1 \oplus 1 \oplus 1 \oplus 1 = 0$$

$$C_2 = 0 \oplus 1 \oplus 0 \oplus 1 = 0$$

$$C_3 = 0 \oplus 1 \oplus 0 \oplus 1 = 0$$

Each Checksum bit will be 0 when there are no errors

- If the **Data** and **Parity** Bits (*i.e. the code word stored*) were read **correctly** then each of the resultant check bits will be 0, (*as we can see above, indicating no corruption of data or parity bits has taken place*).

Error Detection and Correction in Memory Systems

- Suppose an **error** did occur some time in the future and that instead of reading back the code word that was stored (*repeated again below*)

$$d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [1 \ 0 \ 1 \ 0 \ 1 \ 0 \ 1]$$

we actually read back a code word with an error in d_3 (i.e. *bit position 6 in the code word*) thus we actually read back :-

$$d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [1 \ 1 \ 1 \ 0 \ 1 \ 0 \ 1]$$

Error in bit position **6** causing a '0' to be read back as a '1'

Error Detection and Correction in Memory Systems

- To detect the **error** and **locate** its position, we calculate the set of **check bits** by EX-OR'ing the data bits we read back with the **parity bits** according to the **original encoding** equations.

$$\begin{aligned} C_1 &= p_1 \oplus d_1 \oplus d_2 \oplus d_4 \\ C_2 &= p_2 \oplus d_1 \oplus d_3 \oplus d_4 \\ C_3 &= p_3 \oplus d_2 \oplus d_3 \oplus d_4 \end{aligned}$$

- However, in our case we end up with the following **check sum** bits

$$\begin{aligned} C_1 &= 1 \oplus 1 \oplus 1 \oplus 1 = 0 \\ C_2 &= 0 \oplus 1 \oplus 1 \oplus 1 = 1 \\ C_3 &= 0 \oplus 1 \oplus 1 \oplus 1 = 1 \end{aligned}$$

Bit position **6 [110]** meaning **d_3** is wrong

- This just so happens to identify the **bit position** which was in **error** and thus it can be **flipped/inverted** in hardware '**on-the-fly**' during the **read** to **correct it**.

Error Detection and Correction in Memory Systems

Class Exercise

- The following 6 bit data is to be stored [101101]
- Calculate the Parity bits and construct the code word for this data.
- We know from the previous example that 4 parity bits are needed for 6 bits of data in accordance with the equation and code word format below

P_1 = EXOR of data bits occupying positions?

P_2 = EXOR of data bits occupying positions?

P_3 = EXOR of data bits occupying positions?

P_4 = EXOR of data bits occupying positions?

10	9	8	7	6	5	4	3	2	1
d_6	d_5	p_4	d_4	d_3	d_2	p_3	d_1	p_2	p_1

d_6	d_5	p_4	d_4	d_3	d_2	p_3	d_1	p_2	p_1
1	0	?	1	1	0	?	1	?	?

Error Detection and Correction in Memory Systems

Solution

d_6	d_5	p_4	d_4	d_3	d_2	p_3	d_1	p_2	p_1
1	0	1	1	1	0	0	1	0	0

Problem 3

- Now introduce an **error** into data bit d_4 (bit position 7) and calculate the check bits based on what is read back. Can you detect where the error is?

d_6	d_5	p_4	d_4	d_3	d_2	p_3	d_1	p_2	p_1
1	0	1	0	1	0	0	1	0	0

Error Detection and Correction in Memory Systems

How does it work?

- Take a look at the 4 bit data check sum equations shown again below

$$\begin{aligned} C_1 &= p_1 \oplus d_1 \oplus d_2 \oplus d_4 \\ C_2 &= p_2 \oplus d_1 \oplus d_3 \oplus d_4 \\ C_3 &= p_3 \oplus d_2 \oplus d_3 \oplus d_4 \end{aligned}$$

7	6	5	4	3	2	1
d_4	d_3	d_2	p_3	d_1	p_2	p_1

- From this, we see that if any single bit is incorrectly read, it will invert those check bits for which the corrupted bit is part of the equation.
- For example, if d_2 were corrupted, then check bits C_3 and C_1 would be inverted to logic 1 and C_2 would remain logic 0, therefore the check bits would be {101} identifying bit position 5 (i.e. d_2) as the faulty bit.

Error Detection and Correction in Memory Systems

- The apparent 'magic' with this scheme is in choosing the Parity equations such that the faulty bit affects the check bits in just such a way as to identify the source of the faulty bit.
- For example, if we are going to detect a fault in data bit d_4 (code word position 7), then we have to ensure that all three check bits get flipped to 1 (to give a checksum = [111] i.e. 7_{10}) to identify the error in position 7 i.e. d_4 .

This is why d_4 is part of the parity equations for all three check bits as an error in d_4 will flip those check bits to logic 1.

- Likewise an error in d_1 (code word position 3) would have to generate the checksum [011] to identify its location, that is why it is part of the parity equations for check bits C_1 and C_2 and not C_3 .

Error Detection and Correction in Memory Systems

What if the Error lies with one of the Parity Bits?

- If one of the **parity** bits is **erroneous** then it only affects **one check bit** (*see below*) and hence it is only able to produce the check sum values [001], [010] or [100] corresponding to bit positions 1, 2 or 4 which corresponds exactly to the positions occupied by the parity bits in the code word (*i.e. positions that are powers of 2*).

$$\begin{aligned} C_1 &= p_1 \oplus d_1 \oplus d_2 \oplus d_4 \\ C_2 &= p_2 \oplus d_1 \oplus d_3 \oplus d_4 \\ C_3 &= p_3 \oplus d_2 \oplus d_3 \oplus d_4 \end{aligned}$$

An error in a **parity** bit flips exactly one check bit leading to checksum values of [001], [010] or [100], positions 1, 2 or 4 where the parity bits are stored.

- This is why the data and parity bits are stored in the **position/order** they are and why the bit positions in the code word start off at position 1 (*as a check sum of 0 means no error and thus 0 cannot be used as an erroneous position number*).

Error Detection and Correction in Memory Systems

- What happens to the check bits if you get a **double bit error**?
- Obviously, the **3 check bits** can only identify **one** position number so there is no way it could be used to identify the position of **2** incorrect bits.
- However, look at the equations again, what happens if a double bit error were to occur in say **d1** and **d2** (positions **3** and **5**)?

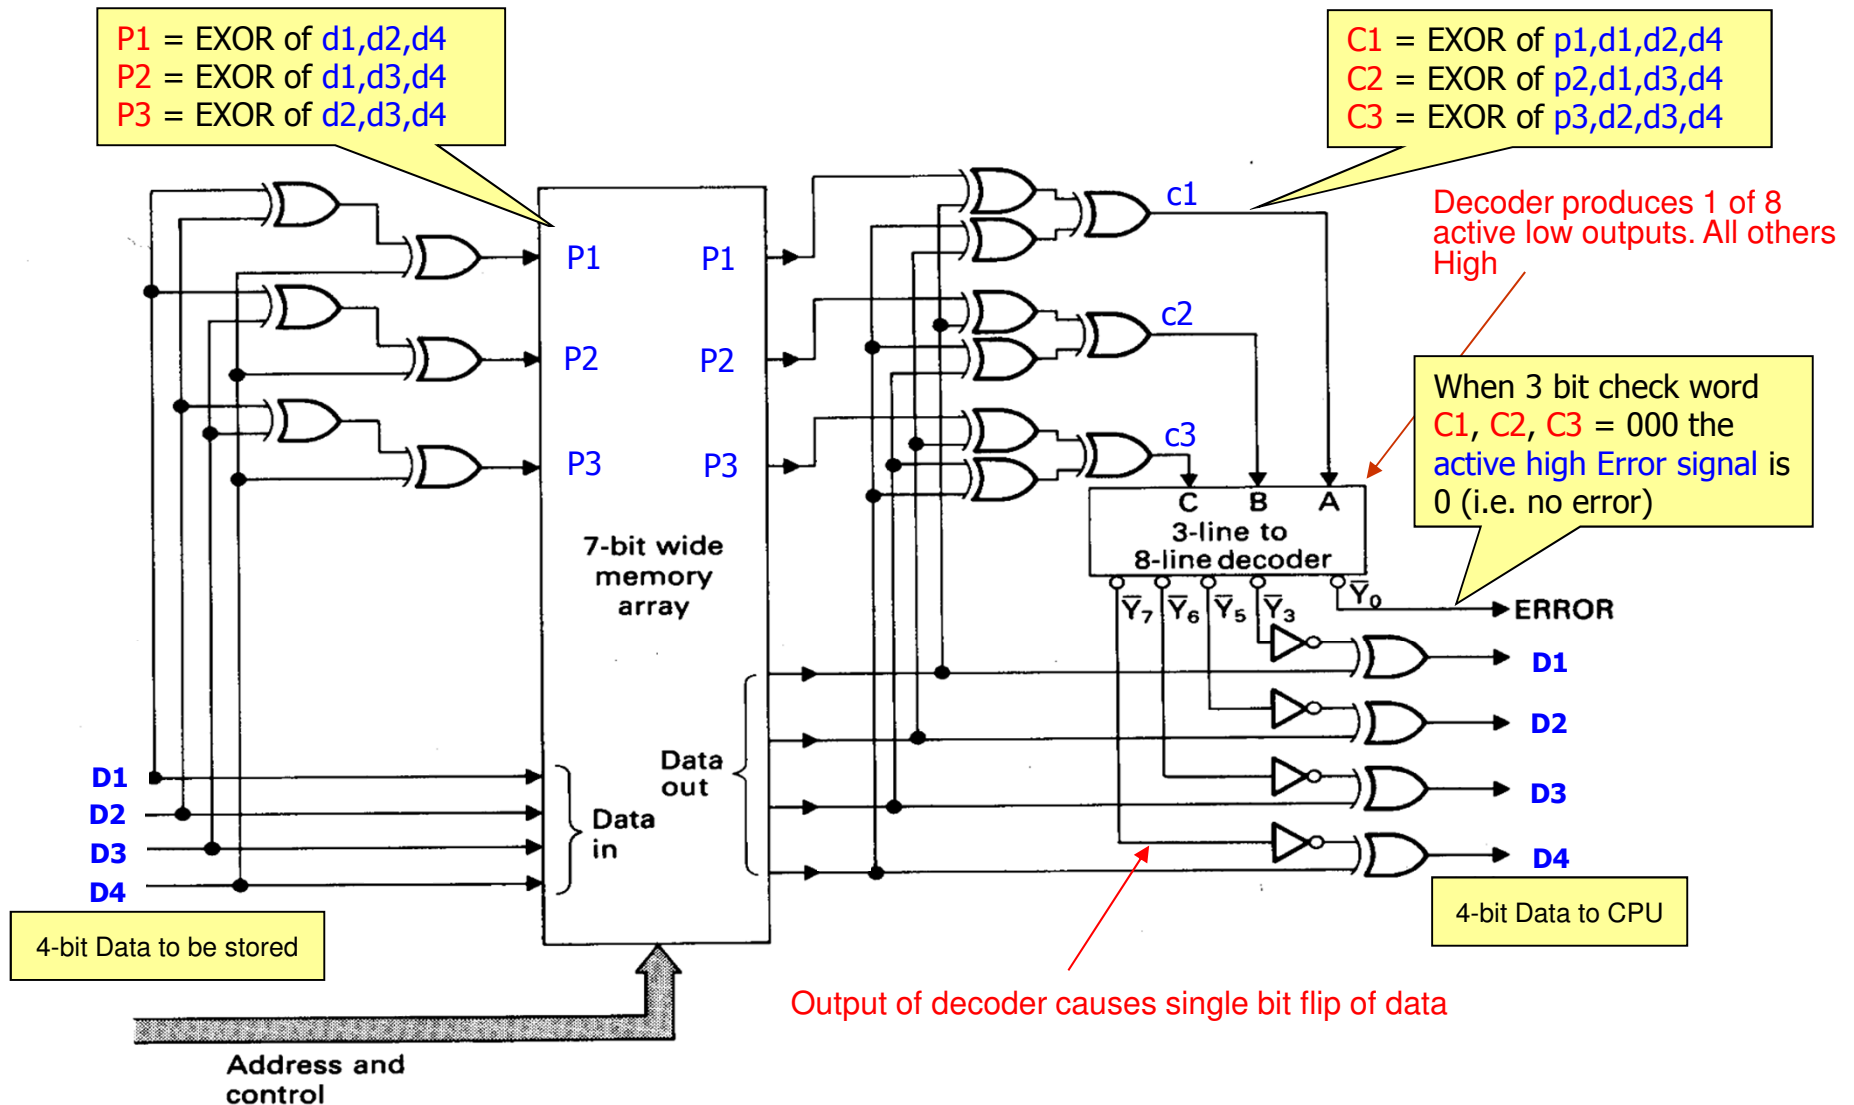
$$C_1 = p_1 \oplus d_1 \oplus d_2 \oplus d_4$$

$$C_2 = p_2 \oplus d_1 \oplus d_3 \oplus d_4$$

$$C_3 = p_3 \oplus d_2 \oplus d_3 \oplus d_4$$

Error Detection and Correction in Memory Systems

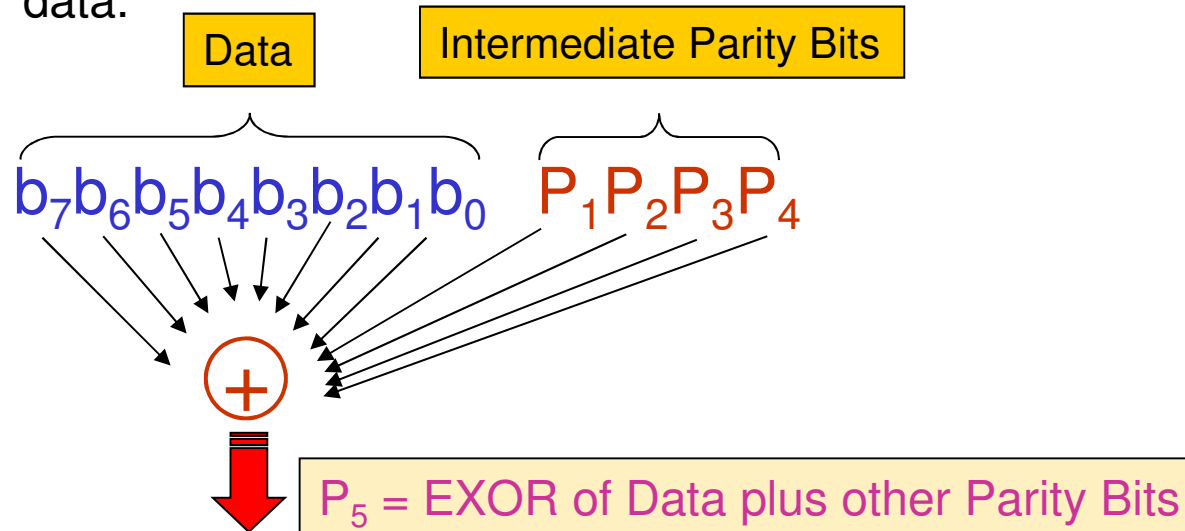
- A Simple Check bit generator plus Error Detector and Corrector Circuit for **4 bit data + 3 bit Parity** (normally this circuitry would be integrated into an ECC Dram controller)



Error Detection and Correction in Memory Systems

Enhanced Error Correction Schemes

- With the simple addition of an **extra parity bit** we could also enhance this scheme to include the **detection** of **double bit errors** (but only with the ability to correct **single bit errors** i.e. **DED/SEC**).
- The **extra Parity bit** is generated by **Exclusive OR'ing** the original **data bits** and the **intermediate parity bits** calculated as previously. We can see this below for 8 bit data.



<u>Data bits</u>	<u># of Check Bits</u>
4	3
8	4
12	5
16	5
20	5
24	5
28	6
32	6
64	7
etc	

4 bits of hamming code could detect and correct single bit errors in 8 bit wide data

Error Detection and Correction in Memory Systems

Checking and Correcting Errors

- During reading, identical hardware to that used to generate the 5 original parity bits is then used to re-create 5 new check bits (*the checksum*) based on the 8 bits of data and 5 bits of parity actually read from memory.

13	12	11	10	9	8	7	6	5	4	3	2	1	Bit Position
p ₅	d ₈	d ₇	d ₆	d ₅	p ₄	d ₄	d ₃	d ₂	p ₃	d ₁	p ₂	p ₁	Parity Bit

- That is, each calculated check bit (C_1 - C_4) is formed from EX-OR'ing each parity bit (P_1 - P_4) read from memory with the data bits that were used to generate the parity bits in the first place, exactly as before. The result will be a 4 bit checksum. If all is well, this 4 bit checksum should be [0000]
- The fifth check bit (C_5) is calculated as the EX-OR of the 8 data bits plus the 5 parity bits READ BACK FROM MEMORY it should also be 0 when no error has occurred.
- If a single bit error occurs reading back original data + parity, then the value for the 5th check bit C_5 will **always** be flipped to 1 regardless of where that error lies.
- However if a double bit error had occurred during the read, the check bit C_5 will be double flipped i.e. still 0 but the other check bits will be non-zero.

Error Detection and Correction in Memory Systems

Interpreting the results

- If check bit C_5 is 0 and the 4 bit check code C_1 - C_4 is 0000, then no errors have occurred and no correction is required.
- If check bit C_5 is 0 and the 4 bit check code C_1 - C_4 is not 0, then a **double bit, non-correctable error** has occurred (i.e. no bits can be corrected). In this case the check code cannot be used to identify the location of any of the errors (*i.e. it must be ignored*)
- If check bit C_5 is 1 then a **single bit error** has occurred in which case, if the 4 bit check code C_1 - C_4 is 0000 then the error has occurred only in Parity Bit p_5 , the data itself is correct and no correction is required.

However if the 4 bit check code C_1 - C_4 is not 0000, then a single bit error has occurred in the bit position indicated by C_1 - C_4 and it can be corrected, this includes the parity bits P_1 - P_4 .

Error Detection and Correction in Memory Systems

Example DED/SEC Parity Encoding with a 4 bit Data Word

- Take the 4 bit data word $d_4, d_3, d_2, d_1 = [1011]$. We calculate the parity bits according to the equations developed earlier.

$$\begin{aligned} p_1 &= d_1 \oplus d_2 \oplus d_4 \\ p_2 &= d_1 \oplus d_3 \oplus d_4 \\ p_3 &= d_2 \oplus d_3 \oplus d_4 \end{aligned}$$

- We generate each parity bit thus

$$\begin{aligned} p_1 &= 1 \oplus 1 \oplus 1 & \text{i.e. } p_1 &= 1 \\ p_2 &= 1 \oplus 0 \oplus 1 & \text{i.e. } p_2 &= 0 \\ p_3 &= 1 \oplus 0 \oplus 1 & \text{i.e. } p_3 &= 0 \end{aligned}$$

Extra Parity bit $p_4 = \text{EX-OR}(p_1, p_2, p_3, d_1, d_2, d_3, d_4) = 0$, Thus code word is

$$p_4 \ d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [0 \ 1 \ 0 \ 1 \ 0 \ 1 \ 0 \ 1]$$

Error Detection and Correction in Memory Systems

- Let's introduce a **single bit error** into parity bit **p4** so that we read back the code word below :-

$$\underline{p_4} \ d_4 \ d_3 \ d_2 \ p_3 \ d_1 \ p_2 \ p_1 = [\underline{1} \ 1 \ 0 \ 1 \ 0 \ 1 \ 0 \ 1]$$

- The only bit affected by a change in **p4** will be **C₄** because no other **check** bits use **p4** in their equation, thus the computed value of the checksum bit **C₄** will be

$$C_4 = \text{EX-OR}(p_1, p_2, p_3, p_4, d_1, d_2, d_3, d_4) = 1$$

The remaining checksum bits **C₁- C₃** will be **000**, indicating that the error is just in **p4** and the data was read **correctly**.

Error Detection and Correction in Memory Systems

- Let's introduce a single bit error into data bit d4.

$$p_4 \quad \underline{d_4} \quad d_3 \quad d_2 \quad p_3 \quad d_1 \quad p_2 \quad p_1 = [0 \quad \underline{0} \quad 0 \quad 1 \quad 0 \quad 1 \quad 0 \quad 1]$$

- This single bit error will cause C_4 to be flipped to logic 1 since

$$C_4 = \text{EX-OR}(p_1, p_2, p_3, p_4, d_1, d_2, d_3, d_4) = 1$$

however, d4 is part of the equations for the other 3 checksum/parity bits hence $C_1, C_2, C_3 = \{111\}$. Hence we have a single bit correctable error in position 7, i.e. d4.

Error Detection and Correction in Memory Systems

- Let's introduce a double bit error : i.e. d_4 and p_1

$$p_4 \quad \underline{d_4} \quad d_3 \quad d_2 \quad p_3 \quad d_1 \quad p_2 \quad \underline{p_1} = [0 \quad \underline{0} \quad 0 \quad 1 \quad 0 \quad 1 \quad 0 \quad \underline{0}]$$

- This double bit error will cause C_4 to be logic 0 due to 2 bit flips since

$$C_4 = \text{EX-OR}(p_1, p_2, p_3, p_4, d_1, d_2, d_3, d_4) = 0$$

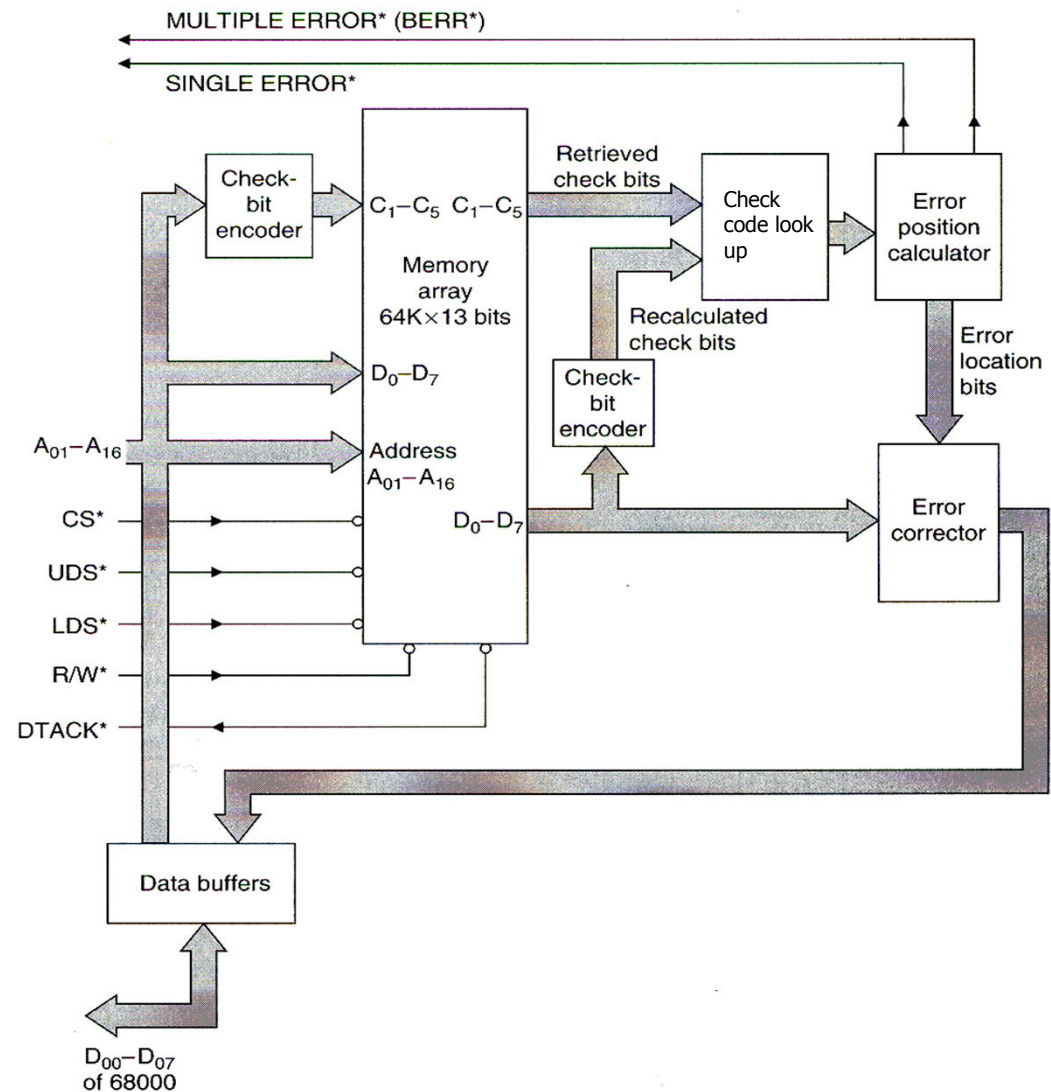
however, d_4 (position 7) is part of the equations for the other 3 checksum/parity bits, thus the other 3 checksum bits should in theory indicate {111} highlighting position 7 in the data stream as the source of the corruption, but because p_1 (the 2nd faulty bit) is also incorrect, C_1 will be double flipped so position of error is read as {110}, incorrectly identifying position 6 as the source of the error.

However this does not matter since our circuitry will detect this is a double bit non-correctable error ($C_4 = 0$, $C_3 - C_1 = \text{non-zero}$).

Error Detection and Correction in Memory Systems

Example SEC/DED ECC Circuit

- The illustration opposite shows the block diagram of a typical error detection and correction circuit for a byte wide memory array.
- This circuit (taken from Clements) shows the **Parity bits** labelled as **C₁-C₅**.
- What about protecting 16 bit data. What are the issues here?
- What is good about using 64 bit wide data as in the PC?



Error Detection and Correction in Memory Systems


A Typical Micron Memory Module with support for SED/SEC ECC: 64 data bits, 8 Check Bits

Features


- 184-pin, dual in-line memory modules (DIMM)
- Fast data transfer rates; PC1600 or PC2100
- Utilizes 200 MT/s and 266 MT/s DDR SDRAM stacked components
- ECC, 1-bit error detection and correction
- Registered inputs with one-clock delay
- Phase-lock loop (PLL) clock driver to reduce loading
- 1GB (128 Meg x 72), 2GB (256 Meg x 72), and 4GB (128 Meg x 72)
- $V_{dd} = V_{ddQ} = +2.5V$
- $V_{ddspd} = +2.3V$ to $+3.6V$
- 2.5V I/O (SSTL_2 compatible)
- Commands entered on each positive CK edge
- DQS edge-aligned with data for READs; center-aligned with data for WRITEs
- Internal, pipelined double data rate (DDR) architecture; two data accesses per clock cycle
- Bidirectional data strobe (DQS) transmitted/received with data, i.e., source-synchronous data capture
- Differential clock inputs (CK and CK#)
- Four internal device banks for concurrent operation
- Selectable burst lengths: 2, 4, or 8
- Auto Refresh and Self Refresh Modes
- 7.8125 μ s maximum average periodic refresh interval
- Serial Presence Detect (SPD) with EEPROM
- Selectable READ CAS latency
- Gold edge contacts

Figure 1: 184-Pin DIMM (MO-206)

Standard PCB (1.7in./43.18mm)



Low Profile PCB (1.2in./30.48mm)



OPTIONS

- Package

184-pin DIMM (Standard)	G
184-pin DIMM (Lead-free) ¹	Y
- Memory Clock/Speed, CAS Latency²

7.5ns (133 Mhz), 266 MT/s, CL = 2	-262
7.5ns (133 Mhz), 266 MT/s, CL = 2	-26A
7.5ns (133 Mhz), 266 MT/s, CL = 2.5	-265
10ns (100 Mhz), 200 MT/s, CL = 2	-202
- Circuit Board

Standard (1.7in./43.18mm)	See note, page 2
Low Profile (1.2in./30.48mm)	See note, page 2

MARKING

ECC used with Sram

- Due to the increased density, reduced feature size and lower operating voltages of Sram cells, Sram technology is also becoming increasingly susceptible to the effects of radiation and electrical noise.
- More and more manufacturers are offering **ECC**, at least as an option, for many of their SRAM chips such as the ISSI Part #**IS61WV51216EDALL** - 512K x 16 SRam (with ECC), as used more recently on the **DE2**.
- You can see much of what has been discussed in this lecture reflected in this block diagram taken from their data sheet. The ECC circuitry is entirely built into the chip.

